
“I’ve Seen How This Goes”: Characterizing the Diversity of LLM Generations and Human Writing via Progressive Conditional Surprise

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Measuring the diversity of creative outputs is central to evaluating post-training
2 mode collapse, comparing decoding strategies, and quantifying creative behavior
3 in both AI and human writing. We propose a new approach to measuring diver-
4 sity using in-context learning, of which the “Decan” metric, $D_{Ca_n} = C \times a_n$,
5 is the working instance we evaluate: a per-byte score read off the per-token log-
6 probabilities of a base model θ in a *single forward pass* per permutation, with no
7 embedding model, no reference corpus, and no human labels. This approach is
8 grounded in information theory, makes use of language model in-context learning
9 to detect a wide range of similarities between any number of inputs, and obviates
10 the need to train a special-purpose model. The same pipeline scores AI samples
11 and human-written response sets, with diversity treated as a property of (responses,
12 prompt, scoring model). On Tevet and Berant’s human-grounded McDiv bench-
13 mark, D_{Ca_n} reaches OCA 0.846 on the McDiv prompt_gen set where it performs
14 best, behind the strongest neural baseline reported in Tevet and Berant (SentBERT,
15 0.897). On the OLMo-2-7B post-training pipeline, D_{Ca_n} drops monotonically
16 across the base $\rightarrow$ SFT $\rightarrow$ DPO $\rightarrow$ RLVR stages, detecting the type of diversity
17 loss that creative-writing applications care about.

1 Motivation

19 The possibility for diverse outputs is necessary, but not sufficient, for creativity. Machine learning
20 researchers that study creativity routinely need to compare the diversity of outputs across generation
21 processes: post-training stages of the same base model (mode collapse), decoding strategies (what
22 does temperature trade off against?), prompting interventions (which prompts have a wide response
23 distribution?), and human-AI hybrid pipelines (would AI use cause human-AI writing to be less
24 creative?). Existing diversity metrics rely on embedding distances [Du and Black, 2019] or surface-
25 level n -gram statistics, both of which lack a human-like ability to recognize arbitrary patterns when
26 those patterns differ from their training data or surface similarities, respectively (Appendix A). A
27 complementary line of work applies information theory directly: Maximum Mutual Information
28 decoding [Li et al., 2016] and Adversarial Information Maximization [Zhang et al., 2018] optimise
29 pairwise mutual information between input and response, but as a training or decoding objective
30 rather than an evaluation-time diagnostic. We propose a new approach to measuring diversity using
31 in-context learning, of which the “Decan” metric $D_{Ca_n} = C \times a_n$ is the working instance we evaluate:
32 it scores AI samples and human-written response sets through the same pipeline of per-byte log-
33 probabilities of a base model θ , in a *single forward pass* over the concatenated responses (Figure 1).
34 There is no embedding model, no reference corpus, no human labels, no auxiliary classifier. The
35 approach uses only the per-token probabilities θ already produces.




Figure 1: The diversity-metric pipeline. Given prompt p and n responses from policy π , we format them with response labels (“Response A:”, “Response B:”, ...) and tokenize. The **conditional track** (left) feeds the concatenated context to the base model θ in a single forward pass per permutation σ , extracts per-response total surprise, divides by each response’s UTF-8 byte count, and averages over permutations to get $\bar{a}_k$. Its last point is a_n . The **unconditional track** (right) scores each r_i independently against the prompt alone, again per-byte; the geometric mean of these surprises defines coherence $C = 1/\text{PPL}_\theta(\pi, p)$, the reciprocal of the geometric-mean per-byte perplexity. The product $D_{Ca_n} = C \times a_n$ measures plausibility-weighted residual diversity (bits/byte). Formal definitions: Sections 2, 3.1, 3.2, 3.3.

36 The intuition is that if responses from a policy π are diverse, seeing one should not help θ predict the
 37 next; if they are repetitive or constrained to a few modes, conditioning on earlier responses should
 38 sharply reduce θ ’s surprise at later ones. We exploit the in-context learning capability described by
 39 Brown et al. [2020], applied here as a measurement lens rather than as a few-shot-prompting setup.
 40 The progressive conditional surprise curve a_k (the per-byte cross-entropy of the k -th response given
 41 the previous $k-1$) captures this signal directly, and its last point is a_n (Section 3.1). A separate
 42 coherence term $C = 1/\text{PPL}_\theta(\pi, p)$, the reciprocal of the geometric-mean per-byte perplexity that θ
 43 assigns to each response individually, prevents pure noise from registering as “diverse” (Section 3.2).
 44 The product $D_{Ca_n} = C \times a_n$ is the working scalar we adopt; it is plausibility-weighted residual
 45 diversity in bits per byte. We validate it against Tevet and Berant’s human-grounded diversity
 46 benchmark (Section 4) and on the OLMo-2-7B post-training pipeline (Section 5), and release a
 47 working open-source implementation alongside.¹ The metric measures diversity as θ perceives it:
 48 outputs differing only in ways θ cannot distinguish in context appear less diverse. This relativity
 49 gives the metric the potential to tighten as base models improve (see Appendix H for a preliminary
 50 scaling experiment).

51 2 Setup and Notation

52 We use the following notation:

- 53 • p be a prompt,
- 54 • π be the policy under evaluation,
- 55 • $r_1, r_2, \dots, r_n \sim \pi(\cdot \mid p)$ be n i.i.d. responses sampled from π ,
- 56 • θ be a trusted base model from which we can obtain per-token log-probabilities,
- 57 • $|r|_{\text{tok}}$ denote the number of tokens in response r ,
- 58 • $\|r\|$ denote the number of bytes in the UTF-8 encoding of response r .

¹Code and data: <https://anonymous.4open.science/r/icl-diversity-67E6/>

59 We define the **cross-entropy** (total surprise) of a response r under θ :

$$-\log_2 \theta(r \mid p) = \sum_{t=1}^{|r|_{\text{tok}}} -\log_2 \theta(r^t \mid r^{<t}, p) \quad (1)$$

60 where r^t is the t -th token and $r^{<t}$ the preceding tokens. Units: bits.² This is the total surprise of the
 61 response under θ , a function of the string and of θ 's distribution but not of θ 's tokenizer (since the
 62 chain rule yields the same total regardless of how the sequence is factored).

63 Similarly, the **conditional cross-entropy** given previous responses $r_{<k} = (r_1, \dots, r_{k-1})$ is
 64 $-\log_2 \theta(r_k \mid r_{<k}, p)$ (bits).

65 For the coherence term and diversity scores (Sections 3.2 and 3.3), we also use the **per-byte cross-**
 66 **entropy rate**:

$$h_\theta(r \mid p) = \frac{-\log_2 \theta(r \mid p)}{\|r\|} \quad (2)$$

67 Units: bits/byte. We adopt this per-byte rate because it works better than total bits in our experiments;
 68 we have not investigated why. Normalising by byte count rather than token count keeps the rate
 69 independent of θ 's tokenizer when comparing base models with different vocabularies.

70 In practice, computing $\theta(r_k \mid r_{<k}, p)$ requires feeding θ a formatted context containing the prompt
 71 and all previous responses (see Section B).

72 3 Method

73 This section defines the three quantities used in the main body: the progressive conditional surprise
 74 curve a_k , the coherence term C , and the scalar score $D_{Ca_n} = C \times a_n$. The full information-theoretic
 75 motivation, the alternative excess-entropy summary E , and the cross-mode learning analysis appear
 76 in Appendices E–H.

77 3.1 Progressive Conditional Surprise

78 Given prompt p and n responses $r_1, \dots, r_n$ from policy π , define

$$a_k = -\log_2 \theta(r_k \mid r_{<k}, p) \quad \text{for } k = 1, \dots, n, \quad (3)$$

79 the total surprise of the k -th response under base model θ given the previous $k-1$ responses. The
 80 sequence $(a_1, \dots, a_n)$ is the **progressive conditional surprise curve**. We normalize each a_k by the
 81 byte count $\|r_k\|$ to get a per-byte (bits/byte) curve that is independent of θ 's tokenizer. All quantities
 82 below are per-byte unless otherwise stated.

83 The intuition is that as θ sees more responses from π , the conditional surprise drops if the responses
 84 share patterns (θ 's in-context learning picks up on modes, stylistic regularities, topic patterns) and
 85 stays flat if they do not. For a policy with rich coherent diversity the curve declines and levels off at a
 86 positive floor; for a policy producing one repeated mode the curve drops sharply to near zero; for
 87 pure noise the curve stays roughly constant at a high value.

88 In practice we average a_k over uniformly random permutations of the response ordering (so each
 89 position averages over all responses and the curve reflects only how θ 's predictions improve with
 90 more context, not which response happened to land in slot k). We take the metric value to be the
 91 last observed point a_n . Equivalently, $a_n = H_\theta(r_n \mid p) - I_\theta(r_n; r_1, \dots, r_{n-1} \mid p)$: the response's
 92 individual surprise under θ minus how much the prior responses make it predictable. High a_n
 93 therefore requires r_n to be both individually surprising and not made predictable by the others, two
 94 properties one would want a diversity score to reward. No fitting or extrapolation step is involved.
 95 The curve a_k is itself informative: its shape distinguishes a one-mode collapse (sharp drop to a low
 96 floor) from richer diversity (gradual decline to a higher floor) in ways the endpoint alone obscures.

²Throughout, "bits" refers to self-information in $\log_2$ units, identical to the "shannon" (Sh) of IEC 80000-13. We use "bits" to match standard usage in the information-theory literature.

Table 1: Intended edge-case behavior of $D_{Ca_n} = C \times a_n$ on the five synthetic scenarios (Appendix C.2), with empirical D_{Ca_n} values under GPT-2 (124M) and Qwen2.5-3B base. The product correctly suppresses pure noise and incoherent multi-mode (via C) and one-mode sets (via a_n); multi-mode coherent is the intended winner. Empirical numbers show smaller a_n for multi-mode coherent than this story predicts: with only 3 distinct modes in a set of 10 responses, θ ’s in-context learning identifies them and a_n is no longer high. In such ambiguous cases reporting the full a_k curve is more principled than the scalar a_n : a different choice of n would reveal the diversity, and showing the curve sidesteps having to pick that n in advance. Appendix C.3 shows a_n rising with mode count. Mixed empirically scores highest (via a high a_n) rather than the predicted “mid” position; reweighting strategies are discussed in Section C.6. On Qwen2.5-3B, multi-mode incoherent also outranks multi-mode coherent; Section 3.3 discusses.

Scenario	Predicted			Empirical D_{Ca_n}	
	C	a_n	D_{Ca_n}	GPT-2	Qwen2.5-3B
Pure noise	low	high	low (via C)	0.04	0.05
Multi-mode incoherent	low	high	low (via C)	0.19	0.29
Multi-mode coherent	high	high	high	0.26	0.19
One-mode	high	low	low (via a_n)	0.10	0.09
Mixed	mid	high	mid	0.52	0.41

97 3.2 The Coherence Term

98 The curve alone is fooled by pure noise: θ can never predict random tokens from one another, so a_n
99 stays at the high unconditional surprise level and noise would dominate the metric. The distinguishing
100 signal lies in how plausible θ finds each individual response, independent of the others.

101 Let $h_\theta(r | p) = -\frac{1}{\|r\|} \log_2 \theta(r | p)$ be the per-byte cross-entropy of response r under θ conditioned
102 on the prompt alone. Define **coherence** as the geometric mean of the per-byte probabilities:

$$C = 2^{-\frac{1}{n} \sum_{i=1}^n h_\theta(r_i | p)} = \frac{1}{\text{PPL}_\theta(\pi, p)}, \quad (4)$$

103 the reciprocal of the geometric-mean per-byte perplexity³ that θ assigns to the responses individually.⁴
104 Perplexity is a standard metric of incoherence, and the geometric form is intended to suppress sets
105 containing incoherent responses: a single sample with high per-byte cross-entropy drives C toward
106 zero, limiting the rescue effect of any single fluent response on an otherwise incoherent set.

107 3.3 The Diversity Score

108 The scalar score is the product:

$$D_{Ca_n} = C \times a_n \quad (\text{bits/byte}). \quad (5)$$

109 It can be read as “how many bits of surprise per byte remain after θ has learned what it can from the
110 other responses, weighted by output plausibility.” The intended edge-case behaviors are summarized
111 in Table 1, alongside empirical D_{Ca_n} on synthetic instances of each scenario under two base models
112 (full per-metric breakdown in Appendix C.2, Table 5).

113 The score depends on p , π ’s responses, and θ . Diversity is therefore a property of (responses, prompt,
114 scoring model), not of the policy in isolation: the same response set can score differently under
115 a stronger or weaker θ . This is by design: θ ’s in-context learning capability is the lens through
116 which diversity is measured, giving the score the potential to improve as base models improve
117 (see Appendix H for a preliminary scaling experiment). For example, Table 1 shows Qwen2.5-3B
118 scoring multi-mode incoherent (0.29) above multi-mode coherent (0.19), reversing our predicted
119 ranking; Figure 3 shows Qwen’s $\bar{a}_k$ curve dropping over k on the multi-mode incoherent scenario,

³In our Tevet evaluation with Qwen2.5-3B, the central 90% (5th to 95th percentile) of per-response bits-per-byte values across 48705 responses span $[0.65, 2.20]$, giving $C \in [0.22, 0.64]$; the mean per-set C across 9741 response sets is 0.42.

⁴For scale, the bits/byte score is the average per-byte cross-entropy $\bar{\ell}$, related to coherence by $C = 2^{-\bar{\ell}}$. GPT-3 davinci on The Pile [Gao et al., 2020] scored ~ 0.72 bits/byte and GPT-2 XL ~ 1.05 bits/byte, giving $C \approx 0.61$ and $C \approx 0.49$ respectively; stronger base models reach lower bits/byte and correspondingly higher C .

Table 2: ConTest results: Spearman ρ and OCA between a set’s diversity class and each metric score. Baselines from Tevet’s pre-computed metrics; $C \times a_n$ and a_n are ours (Qwen2.5-3B, 50 permutations, per-byte). Best result per task in bold.

Metric	prompt_gen		resp_gen		story_gen	
	ρ	OCA	ρ	OCA	ρ	OCA
<i>ConTest (200, with_hds)</i>						
$C \times a_n$ (ours)	+0.584	0.785	+0.391	0.668	+0.686	0.828
a_n (ours)	+0.444	0.715	+0.274	0.641	+0.387	0.684
C (ours)	+0.214	0.625	−0.001	0.555	+0.247	0.632
SentBERT	+0.682	0.815	+0.591	0.791	+0.770	0.896
BERTsts	+0.646	0.820	+0.463	0.714	+0.601	0.780
distinct- n	+0.333	0.675	+0.346	0.677	+0.573	0.772
<i>McDiv_nuggets (~1K, no_hds)</i>						
$C \times a_n$ (ours)	+0.636	0.785	+0.345	0.649	+0.317	0.634
a_n (ours)	+0.487	0.705	+0.225	0.619	+0.124	0.557
C (ours)	+0.138	0.567	+0.082	0.545	+0.251	0.643
SentBERT	+0.728	0.850	+0.532	0.758	+0.633	0.803
BERTsts	+0.683	0.830	+0.393	0.676	+0.344	0.638
distinct- n	−0.003	0.514	−0.002	0.507	−0.002	0.510
<i>McDiv (full, no_hds, ~2K)</i>						
$C \times a_n$ (ours)	+0.729	0.846	+0.500	0.724	+0.523	0.717
a_n (ours)	+0.617	0.781	+0.432	0.698	+0.402	0.668
C (ours)	+0.138	0.565	−0.007	0.512	+0.171	0.594
SentBERT	+0.796	0.897	+0.678	0.830	+0.753	0.867
BERTsts	+0.780	0.893	+0.614	0.781	+0.571	0.740
distinct- n	+0.476	0.746	+0.517	0.738	+0.535	0.744

indicating that its in-context learning recognises the shared template structure as a pattern despite the within-response scrambling. GPT-2’s weaker in-context learning fits the predicted ordering on this row better (0.19 for incoherent vs. 0.26 for coherent), but we still recommend stronger base models when possible – sophisticated patterns are exactly what an ICL-based metric should be able to detect. When comparing policies across a prompt suite, D_{Ca_n} can be averaged over prompts or differenced; see Appendix G.

4 Tevet–Berant Diversity-Eval: Human-Grounded Validation

A diversity metric should be tested against human judgements on a standardized eval; Tevet and Berant [2021] provide both. Their McDiv and ConTest datasets pair human-grounded diversity labels with a fixed comparison protocol (Spearman ρ and OCA: *optimal classification accuracy*, the best accuracy achievable by a one-dimensional threshold separating low- and high-diversity sets), the standard methodology for ranking diversity metrics against human judgements. The response sets are written by Mechanical Turk workers, and for both McDiv_nuggets and ConTest the high/low label is fixed by the construction protocol itself: a worker writes five different continuations (the high-diversity set), then self-selects one of those continuations and paraphrases it five times preserving content but varying form (the low-diversity set; see Appendix F for the full protocol and a construction-confound caveat). We run D_{Ca_n} through this evaluation pipeline.

Setup. We use Qwen2.5-3B (base) with 50 permutations in completion format (Appendix B.1) on the released splits: 6002 McDiv sets (no_hds), 3069 McDiv_nuggets sets (no_hds), and 670 ConTest sets (with_hds), each with 5 worker-written responses per set and a binary high/low diversity label. We follow Tevet and Berant in reporting Spearman ρ and OCA, and additionally report ROC AUC (the natural metric for binary classification).

Headline result. On the binary tasks D_{Ca_n} clearly tracks diversity, with OCA between 3.7% and 21.0% behind SentBERT (Table 2). On McDiv prompt_gen, where D_{Ca_n} performs best, it reaches $\rho = +0.729$ and OCA = 0.846, against SentBERT’s $\rho = +0.796$ and OCA = 0.897; the

Table 3: DecTest results: Spearman ρ between sampling temperature and each metric score (1000 samples, no_hds). Reported for parity with Tevet and Berant; the labels here are sampling temperatures rather than human judgements, so DecTest is not part of the framing for $D_{C_{a_n}}$. That a_n tracks temperature is mechanically unsurprising: a_n is a conditional-entropy quantity, and sampling temperature directly scales the entropy of the policy, so a positive ρ here is a sanity check rather than evidence about creative diversity.

Metric	prompt_gen	resp_gen	story_gen
$C \times a_n$ (ours)	+0.847	+0.771	+0.763
a_n (ours)	+0.932	+0.924	+0.779
C (ours)	-0.727	-0.813	-0.547
distinct- n	+0.917	+0.894	+0.758
BERTScore	+0.878	+0.874	+0.694
SentBERT	+0.747	+0.801	+0.645
cos-sim	+0.873	+0.895	+0.712

corresponding ROC AUC is 0.921. The headline takeaway is that an ICL-based diversity metric reaches close to a sentence-embedding baseline on a human-grounded eval, with no embedding model, no reference corpus, and no human labels in the metric itself.

DecTest, included for completeness. Tevet and Berant also release DecTest (Table 3), a diagnostic where the “label” is the sampling temperature used to generate each response set rather than a human judgement. We include it for parity with the released benchmark only: we do not believe that temperature labels map onto creative/meaningful output diversity. The raw a_n achieves $\rho = +0.932$ on prompt_gen and $\rho = +0.924$ on resp_gen, matching or exceeding all baselines.

Caveats. The McDiv_nuggets construction protocol introduces a confound: low-diversity sets are paraphrases of specific, dramatic endings that are intrinsically more surprising to the base model than typical high-diversity continuations. Appendix F documents the mechanism, the per-byte a_1 gap, and an evidence-by-length-bin breakdown showing the effect is content-driven not length-driven. The C weighting in $D_{C_{a_n}}$ helps on this benchmark partly for that confound-related reason; on settings without such a confound we expect a_n to carry most of the signal.

5 OLMo-2-7B Post-Training: AI-Side Validation

The Tevet evaluation in Section 4 validates $D_{C_{a_n}}$ against human-grounded labels. This section evaluates $D_{C_{a_n}}$ on real policies: we sample from four post-training stages of the same base model and ask whether the metric detects the diversity loss widely attributed to RLHF [Kirk et al., 2023, Zhang et al., 2025a, Padmakumar and He, 2023]. The same scoring pipeline applies to both human-written response sets (Section 4) and the policy-sampled sets here, so θ is held fixed across the two evaluations.

Setup. We sample from the four released stages of the OLMo-2-1124-7B pipeline [OLMo et al., 2024]: the pretrained base model, its SFT checkpoint, the DPO checkpoint trained on preference data, and the final RLVR-tuned Instruct checkpoint. On two prompt sets, 200 AlpacaFarm evaluation prompts [Dubois et al., 2023] (seed=42 subsample of the 805-prompt set later adopted by AlpacaEval) and 100 curated NoveltyBench prompts [Zhang et al., 2025c], we draw $K=10$ responses per prompt per stage (temperature 1.0, top- p 1.0, max_new_tokens=100). Base runs raw; SFT / DPO / Instruct receive the prompt through their own chat template. Because per-byte cross-entropy in a causal LM decreases with within-response context, we truncate every (stage, prompt) tuple’s responses to a common per-prompt UTF-8 byte length before scoring,⁵ and drop prompts where that common length falls below 50 bytes (150/200 AlpacaEval and 39/100 NB-curated prompts retained; see Section 6).

⁵Truncation is applied to the response *string* via `text.encode("utf-8")[:N].decode("utf-8", errors="ignore")` and the truncated string is then re-tokenised for the forward pass. We deliberately do not truncate at token boundaries: a_n is per-byte, so we need a fixed byte denominator across stages, but tokens-per-byte is tokenizer-dependent and varies across responses with the same byte budget. That variation is harmless because both the bits-numerator and the byte-denominator are computed on the same truncated string.

We score each (stage, prompt) group with $D_{C_{a_n}}$ using Qwen2.5-3B as θ and 25 permutations. For comparison we also compute Kirk et al.’s EAD and distinct- n (averaged over $n=1 \dots 5$) and a SentBERT-embedding [Reimers and Gurevych, 2019] similarity-to-diversity reduction.

Hypotheses. We pre-register three one-sided paired Wilcoxon signed-rank tests [Dror et al., 2018] with family-wise Bonferroni correction across the three contrasts [Dror et al., 2017] ($\alpha = 0.05/3$):

- **H1a** $D_{\text{base}} > D_{\text{SFT}}$: SFT narrows toward the helpful-assistant style.
- **H1b** $D_{\text{SFT}} > D_{\text{DPO}}$: DPO is trained on preference data and inherits its typicality bias [Zhang et al., 2025a].
- **H1c** $D_{\text{base}} > D_{\text{RLVR}}$: cumulative effect of the full post-training pipeline.

The DPO-vs-RLVR comparison is reported as an exploratory two-sided contrast (H1’): we have no directional prediction for which post-training stage loses more diversity.

Table 4: Per-prompt $D_{C_{a_n}}$ by OLMo-2-7B stage on the length-matched subset, with pre-registered paired tests. Reported p -values for H1a–H1c are Bonferroni-corrected ($\times 3$); the H1’ row (marked *) is an uncorrected, two-sided exploratory comparison whose direction we had no prior expectation for, and is not a load-bearing claim.

AlpacaEval			
Stage	$D_{C_{a_n}}$ mean	std	n
Base	0.481	0.038	150
SFT	0.329	0.084	150
DPO	0.286	0.072	150
Instruct (RLVR)	0.281	0.071	150
Contrast	Δ	d_z	p_{Bonf}
Base>SFT	0.151	1.615	3.4×10^{-24}
SFT>DPO	0.044	0.675	1.7×10^{-13}
Base>Instruct (RLVR)	0.200	2.425	5.2×10^{-26}
DPO $\neq$ Instruct (H1’)	0.005	0.227	0.004*
NoveltyBench curated			
Stage	$D_{C_{a_n}}$ mean	std	n
Base	0.481	0.030	39
SFT	0.369	0.085	39
DPO	0.312	0.082	39
Instruct (RLVR)	0.303	0.086	39
Contrast	Δ	d_z	p_{Bonf}
Base>SFT	0.112	1.212	4.1×10^{-8}
SFT>DPO	0.057	0.957	2.8×10^{-6}
Base>Instruct (RLVR)	0.179	1.889	2.3×10^{-10}
DPO $\neq$ Instruct (H1’)	0.010	0.375	0.028*

Results. Table 4 reports per-stage means and the pre-registered Wilcoxon contrasts on both prompt sets. On AlpacaEval and NoveltyBench-curated alike, $D_{C_{a_n}}$ decreases monotonically across base $\rightarrow$ SFT $\rightarrow$ DPO $\rightarrow$ Instruct (RLVR), and all three pre-registered contrasts (H1a–H1c) are significant after Bonferroni correction. The underlying $\bar{a}_k$ curves and per-prompt $D_{C_{a_n}}$ distributions are visualised in Appendix I (Figure 17): each later stage’s curve lies below the base curve at every $k \geq 2$, and the per-prompt distribution shifts toward lower $D_{C_{a_n}}$ as the pipeline advances.

Discussion. The monotone drop across all three pre-registered contrasts is consistent with the post-training-reduces-diversity findings of Kirk et al. [2023], Padmakumar and He [2023], now measured by an ICL-based diversity metric that needs neither an embedding model nor a sentence similarity function. Per-prompt $D_{C_{a_n}}$ also agrees with the lexical (EAD) and semantic (SentBERT) baselines on the same length-matched subset, confirming the metrics see the same diversity-loss signal (Appendix I, Figure 18).

Released artefacts. The $200 + 100$ prompts $\times 4$ stages $\times K=10$ generations are released as a public dataset under `results/rlhf_experiment/` in the project repository. To our knowledge this

is the first public release of $K \geq 10$ sampled responses across a paired SFT/DPO/RL pipeline, filling a gap that blocked direct diversity replication of Kirk et al. [2023] (whose $K=16$ generations were not publicly released).

6 Limitations

We collect the main limitations of the framework and of our $D = C \times a_n$ choice in one place so readers can calibrate how strongly to update on the results above. Several points below recap caveats already made in their natural context and cross-reference those for detail.

Measurement is relative to θ 's perception. The metric evaluates diversity *as perceived by the base model θ* : if π 's outputs differ only in ways θ cannot distinguish from context, the metric underestimates diversity (Section 1, "Important caveat"). Concretely, the off-diagonal sign of the pairwise surprise-reduction matrix changes with θ 's scale (Section C.4), so a_k curves for the same response set are not directly comparable across different choices of θ .

Metric selection saw the full Tevet benchmark. We chose $C \times a_n$ (rather than alternative summaries such as $C \times E$; Appendix E) after observing its performance on the full Tevet diversity-eval benchmark, without holding out a validation split. The Tevet numbers in Section 4 may therefore be optimistically biased: the scalar form was selected with visibility into those results. The OLMo-2-7B post-training experiment (Section 5) was scored after the metric form was fixed and provides independent validation; the synthetic mode-count experiments (Section C.3) informally influenced the choice (had they failed, we would likely have revised the metric) and so are not strictly independent.

External benchmark performance trails the strongest baseline. On Tevet and Berant's diversity-eval benchmark, $C \times a_n$ trails SentBERT by 3.7%–21.0% OCA across the nine binary tasks (Section 4); scaling the base model to Qwen3-30B-A3B-Base [Yang et al., 2025] did not close the gap (Appendix H). We position ICL-based diversity metrics (of which $D_{C a_n}$ is one instance we found to work well) as a principled complement to embedding-based metrics, not a replacement.

Length-matching drops short-response prompts. Reporting in bits/byte (rather than total bits) makes per-byte cross-entropy not invariant to response length, so the OLMo-2-7B experiment scores a length-matched subset and drops 111/300 prompts whose common per-prompt byte length falls below 50 bytes; full mechanics and the un-truncated robustness check are in Section C.6.

$D = C \times a_n$ is a pragmatic choice, not a theoretically derived one. The product form was selected because it captures the properties we wanted a diversity score to have (residual surprise remains after the base model has seen previous responses, low-coherence outputs are penalised, and the result stays on a bits-per-byte scale), not because it is derived from an axiomatic information-theoretic setup. Other scalars respecting the same constraints (weighted integrals of $a_k - a_\infty$, slope-based scores, or coherence applied at a different point on the curve) could equally be defended (see "The framework admits other metrics" in Section C.6). We adopt $D_{C a_n}$ because it is interpretable and works empirically across our validation regimes, not because it is uniquely correct. We invite other researchers to build on this work to develop other formulas that characterize the a_k curve in meaningful ways.

7 Conclusion

We have presented a new approach to measuring diversity using in-context learning, of which $D_{C a_n} = C \times a_n$ is the working instance we evaluate: a base model's in-context learning detects similarities between arbitrary numbers of responses in a single forward pass. The approach requires no embedding model, no reference corpus, no human labels, and no special-purpose training. The same pipeline scores AI samples and human-written sets. On Tevet and Berant's human-grounded benchmark it is near the level of the strongest sentence-embedding baseline. On the OLMo-2-7B post-training pipeline it drops across the base $\rightarrow$ SFT $\rightarrow$ DPO/RLVR stages, tracking the diversity loss attributed to RLHF-style training.

Impact Statement

This paper presents work which may be used in evaluation pipelines that select for more capable AI systems. We view the prospect of increasingly capable future AI as one of unclear positive or negative net impact. This work may also contribute to the automation of human creative roles. We publish it in the hope that better understanding of machine-generated diversity will contribute to better evaluation and oversight, and that these benefits outweigh the risks.

References

- Lukas Berglund, Meg Tong, Max Kaufmann, Mikita Balesni, Asa Cooper Stickland, Tomasz Korbak, and Owain Evans. The reversal curse: LLMs trained on "a is b" fail to learn "b is a", 2023. URL <http://arxiv.org/abs/2309.12288v4>.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners, 2020. URL <http://arxiv.org/abs/2005.14165v4>.
- James P. Crutchfield and David P. Feldman. Regularities unseen, randomness observed: Levels of entropy convergence. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 13(1):25–54, March 2003. ISSN 1089-7682. doi: 10.1063/1.1530990. URL <http://dx.doi.org/10.1063/1.1530990>.
- Rotem Dror, Gili Baumer, Marina Bogomolov, and Roi Reichart. Replicability analysis for natural language processing: Testing significance with multiple datasets. *Transactions of the Association for Computational Linguistics*, 5:471–486, 2017. doi: 10.1162/tacl_a_00074. URL <https://aclanthology.org/Q17-1033/>.
- Rotem Dror, Gili Baumer, Segev Shlomov, and Roi Reichart. The hitchhiker’s guide to testing statistical significance in natural language processing. In Iryna Gurevych and Yusuke Miyao, editors, *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1383–1392, Melbourne, Australia, July 2018. Association for Computational Linguistics. doi: 10.18653/v1/P18-1128. URL <https://aclanthology.org/P18-1128/>.
- Wenchao Du and Alan W Black. Boosting dialog response generation. In Anna Korhonen, David Traum, and Lluís Màrquez, editors, *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 38–43, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1005. URL <https://aclanthology.org/P19-1005/>.
- Yann Dubois, Xuechen Li, Rohan Taori, Tianyi Zhang, Ishaan Gulrajani, Jimmy Ba, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. AlpacaFarm: A simulation framework for methods that learn from human feedback, 2023. URL <http://arxiv.org/abs/2305.14387v4>.
- Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, Shawn Presser, and Connor Leahy. The pile: An 800gb dataset of diverse text for language modeling, 2020. URL <http://arxiv.org/abs/2101.00027v1>.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego

298 Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova,
 299 Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang, Gabriel
 300 Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire Mialon,
 301 Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan
 302 Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet,
 303 Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer van der Linde,
 304 Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie
 305 Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua
 306 Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani, Kate Plawiak,
 307 Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz Malik, Kuenley
 308 Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yeary, Laurens van der Maaten, Lawrence
 309 Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo, Lukas Blecher, Lukas
 310 Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat Singh, Manohar Paluri,
 311 Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya Pavlova, Melanie
 312 Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman Goyal, Narjes
 313 Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang, Olivier Duchenne,
 314 Onur Çelebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic, Peter Weng, Prajjwal
 315 Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu, Qing He, Qingxiao Dong,
 316 Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira Cabral, Robert Stojnic,
 317 Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain Sauvestre, Ronnie
 318 Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar Hosseini, Sahana
 319 Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov, Shao-liang Nie,
 320 Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale, Shun Zhang, Simon
 321 Vandenheide, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane Collot, Suchin Gururangan,
 322 Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha, Thomas Georgiou, Thomas
 323 Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal Karn, Vedanuj Goswami,
 324 Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gonguet, Virginie Do, Vish Vogeti,
 325 Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenyin Fu, Whitney Meers, Xavier
 326 Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide Xia, Xinfeng Xie, Xuchao
 327 Jia, Xuewei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei, Yi Wen, Yiwen Song,
 328 Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpierre Coudert, Zheng Yan, Zhengxing Chen, Zoe
 329 Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey, Adam Shajnfeld, Adithya
 330 Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma, Alex Boesenberg, Alexei
 331 Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo, Anam Yunus, Andrei Lupu,
 332 Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew Poulton, Andrew Ryan, Ankit
 333 Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita Saraf, Arkabandhu Chowdhury,
 334 Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh Yazdan, Beau James, Ben Maurer,
 335 Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola, Bhargavi Paranjape, Bing Liu,
 336 Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence, Brani Stojkovic, Brian Gamido,
 337 Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu, Changhan Wang, Changkyu
 338 Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris Tindal, Christoph Feichtenhofer,
 339 Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel Li, David Adkins, David Xu,
 340 Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich, Didem Foss, Dingkang Wang, Duc
 341 Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine Montgomery, Eleonora Presani, Emily
 342 Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban Arcaute, Evan Dunbar, Evan Smothers,
 343 Fei Sun, Felix Kreuk, Feng Tian, Filippas Kokkinos, Firat Ozgenel, Francesco Caggioni, Frank
 344 Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella Schwarz, Gada Badeer, Georgia Swee,
 345 Gil Halpern, Grant Herman, Grigory Sizov, Guangyi, Zhang, Guna Lakshminarayanan, Hakan Inan,
 346 Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha, Haroun Habeeb, Harrison Rudolph,
 347 Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan Zhan, Ibrahim Damlaj, Igor Molybog,
 348 Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai Gat, Jake Weissman, James Geboski, James
 349 Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya, Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny
 350 Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica Zhong, Jian Jin, Jingyi Yang, Joe Cummings,
 351 Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan Torres, Josh Ginsburg, Junjie Wang, Kai
 352 Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal, Katayoun Zand, Kathy Matosich, Kaushik
 353 Veeraraghavan, Kelly Michelena, Keqian Li, Kiran Jagadeesh, Kun Huang, Kunal Chawla, Kyle
 354 Huang, Lailin Chen, Lakshya Garg, Lavender A, Leandro Silva, Lee Bell, Lei Zhang, Liangpeng
 355 Guo, Licheng Yu, Liron Moshkovich, Luca Wehrstedt, Madian Khabsa, Manav Avalani, Manish
 356 Bhatt, Martynas Mankus, Matan Hasson, Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim

Naumov, Maya Lathi, Meghan Keneally, Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov, Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat, Mohammad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White, Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager, Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang, Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra, Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh, Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Sheng Feng, Shenghao Lin, Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang, Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe, Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou, Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov, Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocheng Tang, Xiaojuan Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi, Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu, Wang, Yu Zhao, Yuchen Hao, Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu Yang, Zhiwei Zhao, and Zhiyu Ma. The llama 3 herd of models, 2024. URL <http://arxiv.org/abs/2407.21783v3>.

Yanzhu Guo, Guokan Shang, and Chloé Clavel. Benchmarking linguistic diversity of large language models, 2024. URL <http://arxiv.org/abs/2412.10271v2>.

Robert Kirk, Ishita Mediratta, Christoforos Nalmpantis, Jelena Luketina, Eric Hambro, Edward Grefenstette, and Roberta Raileanu. Understanding the effects of rlhf on llm generalisation and diversity, 2023. URL <http://arxiv.org/abs/2310.06452v3>.

Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. A diversity-promoting objective function for neural conversation models. In Kevin Knight, Ani Nenkova, and Owen Rambow, editors, *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 110–119, San Diego, California, June 2016. Association for Computational Linguistics. doi: 10.18653/v1/N16-1014. URL <https://aclanthology.org/N16-1014/>.

Meta AI. Llama 3.2: Revolutionizing edge AI and vision with open, customizable models. Meta AI Blog, September 2024. URL <https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/>.

Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, Akshita Bhagia, Yuling Gu, Shengyi Huang, Matt Jordan, Nathan Lambert, Dustin Schwenk, Oyvind Tafjord, Taira Anderson, David Atkinson, Faeze Brahman, Christopher Clark, Pradeep Dasigi, Nouha Dziri, Allyson Ettinger, Michal Guerquin, David Heineman, Hamish Ivison, Pang Wei Koh, Jiacheng Liu, Saumya Malik, William Merrill, Lester James V. Miranda, Jacob Morrison, Tyler Murray, Crystal Nam, Jake Poznanski, Valentina Pyatkin, Aman Rangapur, Michael Schmitz, Sam Skjonsberg, David Wadden, Christopher Wilhelm, Michael Wilson, Luke Zettlemoyer, Ali Farhadi, Noah A. Smith, and Hannaneh Hajishirzi. 2 olmo 2 furious, 2024. URL <http://arxiv.org/abs/2501.00656v3>.

Vishakh Padmakumar and He He. Does writing with language models reduce content diversity?, 2023. URL <http://arxiv.org/abs/2309.05196v3>.

Tianyi Qiu, Ahmed Hani Ismail, Zhonghao He, and Shi Feng. Self-improvement as coherence optimization: A theoretical account, 2026. URL <http://arxiv.org/abs/2601.13566v1>.

Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI,

2019. URL https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.

Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks, 2019. URL <http://arxiv.org/abs/1908.10084v1>.

Guy Tevet and Jonathan Berant. Evaluating the evaluation of diversity in natural language generation. In Paola Merlo, Jorg Tiedemann, and Reut Tsarfay, editors, *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 326–346, Online, April 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.eacl-main.25. URL <https://aclanthology.org/2021.eacl-main.25/>.

Zihan Wang, Kangrui Wang, Qineng Wang, Pingyue Zhang, Linjie Li, Zhengyuan Yang, Xing Jin, Kefan Yu, Minh Nhat Nguyen, Licheng Liu, Eli Gottlieb, Yiping Lu, Kyunghyun Cho, Jiajun Wu, Li Fei-Fei, Lijuan Wang, Yejin Choi, and Manling Li. Ragen: Understanding self-evolution in llm agents via multi-turn reinforcement learning, 2025. URL <http://arxiv.org/abs/2504.20073v2>.

Jiaxin Wen, Zachary Ankner, Arushi Somani, Peter Hase, Samuel Marks, Jacob Goldman-Wetzler, Linda Petrini, Henry Sleight, Collin Burns, He He, Shi Feng, Ethan Perez, and Jan Leike. Unsupervised elicitation of language models, 2025. URL <http://arxiv.org/abs/2506.10139v2>.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report, 2024. URL <http://arxiv.org/abs/2412.15115v2>.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jing Zhou, Jingren Zhou, Junyang Lin, Kai Dang, Keqin Bao, Kexin Yang, Le Yu, Lianghao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang Ren, Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yinger Zhang, Yu Wan, Yuqiong Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan Qiu. Qwen3 technical report, 2025. URL <http://arxiv.org/abs/2505.09388v1>.

Jiayi Zhang, Simon Yu, Derek Chong, Anthony Sicilia, Michael R. Tomz, Christopher D. Manning, and Weiyang Shi. Verbalized sampling: How to mitigate mode collapse and unlock llm diversity, 2025a. URL <http://arxiv.org/abs/2510.01171v3>.

Tianhui Zhang, Bei Peng, and Danushka Bollegala. Evaluating the evaluation of diversity in commonsense generation, 2025b. URL <http://arxiv.org/abs/2506.00514v1>.

Yiming Zhang, Harshita Diddee, Susan Holm, Hanchen Liu, Xinyue Liu, Vinay Samuel, Barry Wang, and Daphne Ippolito. Noveltybench: Evaluating language models for humanlike diversity, 2025c. URL <http://arxiv.org/abs/2504.05228v4>.

Yizhe Zhang, Michel Galley, Jianfeng Gao, Zhe Gan, Xiuju Li, Chris Brockett, and Bill Dolan. Generating informative and diverse conversational responses via adversarial information maximization, 2018. URL <http://arxiv.org/abs/1809.05972v5>.

Yaoming Zhu, Sidi Lu, Lei Zheng, Jiaxian Guo, Weinan Zhang, Jun Wang, and Yong Yu. Taxygen: A benchmarking platform for text generation models, 2018. URL <http://arxiv.org/abs/1802.01886v1>.

A Related Work

Sampling diversity metrics. Work on decoding-time diversity typically operates at the surface level: n -gram overlap [Li et al., 2016] and self-BLEU [Zhu et al., 2018]. Our approach operates at the distributional level and can distinguish policies that produce lexically varied but semantically redundant outputs.

Diversity evaluation benchmarks. Tevet and Berant [2021] introduced a systematic framework for evaluating diversity metrics, with their McDiv benchmark providing human-labeled response sets at two diversity levels. We use McDiv for validation (Section 4), though we identify a construction confound in how its low-diversity sets are produced (Appendix F). More recent work has revisited the meta-evaluation problem. NoveltyBench [Zhang et al., 2025c] defines diversity through pairwise functional equivalence rather than binary labels, training a DeBERTa classifier to group generations into equivalence classes and computing a $Distinct_k$ score (the number of meaningfully different outputs in k samples), conceptually close to our mode count. However, their ground truth is itself a trained classifier (79% accuracy), making it unclear whether correlation with their labels validates a metric or merely measures agreement between two model-based scores. Zhang et al. [2025b] conduct a meta-evaluation of diversity metrics for constrained commonsense generation, using GPT-4o as an annotator in place of crowd workers. Guo et al. [2024] benchmark linguistic diversity of LLMs, building on Tevet and Berant’s framework with a broader set of generation tasks.

Perplexity. The coherence term $C = 1/\text{PPL}$ (Section 3.2) connects our framework to the standard language modeling evaluation metric. The primary diversity score $D_{Ca_n} = C \times a_n$ operates in bits/byte, weighting the residual surprise floor by output plausibility.

Coherence as a signal. Our coherence term C uses the base model’s predictive distribution as an unsupervised quality signal. This connects to a broader line of work using model-internal coherence without external supervision: Qiu et al. [2026] show theoretically that feedback-free self-improvement methods work by optimizing coherence (compressibility of context-to-behavior mappings), and Wen et al. [2025] use internal coherence maximization to elicit capabilities from language models. Our framework leverages a related insight: the base model’s ability to compress responses in context provides a meaningful signal about their diversity structure.

Embedding-based diversity. Du and Black [2019] introduced the embedding-based approach, clustering sentence embeddings of generated responses with k -means and reporting cluster inertia as the diversity score; subsequent work has explored alternative geometric statistics over embedded responses. These approaches are complementary to ours: they capture semantic distance in a learned representation space, while our metric captures statistical independence under a generative model. The approaches may disagree when θ ’s in-context learning captures structure invisible to the embedding model, or vice versa.

Output collapse in RL training. Wang et al. [2025] document an “Echo Trap” pattern during multi-turn agent RL: early-stage agents produce varied symbolic reasoning, then collapse to deterministic templates as training proceeds. They detect this collapse using within-prompt reward standard deviation as an early-warning signal, a proxy that relies on the reward distinguishing degenerate outputs from diverse ones. The a_k framework provides a text-level alternative: collapse should appear as a drop in a_n on the agent’s own outputs, independent of reward structure. A downside of using the a_k framework is that it is unconnected from a downstream task. Responses can both be “diverse” as measured by a_k -based metrics and fail uniformly.

Excess entropy. The a_k curve also admits an excess-entropy summary related to the computational mechanics literature [Crutchfield and Feldman, 2003]; see Appendix E. We found it empirically inferior to D_{Ca_n} on Tevet and Berant’s diversity-eval benchmark [Tevet and Berant, 2021].

B Practical Considerations

B.1 Formatting the Conditioning Context

Computing $\theta(r_k \mid r_{<k}, p)$ requires feeding θ a context containing the prompt and previous responses. The formatting choice affects results, and we use two formats depending on whether the responses are best read as parallel answers to the prompt or as continuations of it. Both are implemented as `format_conditioning_context` in `src/icl_diversity/core.py` and selected via a `format_mode` argument throughout the pipeline.

Instruct format. The default, used for the synthetic scenarios (Appendix C.2), the mode-count experiments (Appendix C.3), and the OLMo-2-7B post-training case study (Section 5):

```
[prompt p]
Response A: [r_1]
Response B: [r_2]
Response C: [r_3]
```

The prompt appears once and each response is introduced by a labelled header (“Response A:”, “Response B:”, ..., with labels rolling over to “AA”, “AB” past “Z”), encouraging the base model’s in-context learning to treat the responses as parallel answers to the same prompt.

Completion format. Used for the Tevet and Berant evaluation (Section 4), where each response set is a continuation of a shared narrative prompt rather than an instruction-style answer:

```
1. [prompt p] [r_1]
2. [prompt p] [r_2]
3. [prompt p] [r_3]
```

The prompt is repeated immediately before each response so θ scores r_k in the same prompt-as-context that produced it. Per-token log-probabilities are extracted only over the response tokens; the bits accumulated over the repeated prompt prefix are excluded from a_k , keeping the bits-numerator and the per-byte denominator both attributable to r_k alone.

Base vs. instruction-tuned models. We use a base (non-instruction-tuned) model as θ for two reasons. First, an instruction-tuned model has had its output distribution shaped by RLHF or similar procedures, so using one as θ reintroduces the kind of distributional bias we seek to avoid. Second, an instruction-tuned θ may assign systematically different probabilities to text that follows or violates its alignment training, introducing a confound between coherence-as-fluency and coherence-as-alignment that the metric should not conflate. Modern base models already exhibit strong in-context learning from pretraining alone, so this choice does not sacrifice ICL capability. Optimizing the prompt format to maximally elicit in-context learning from θ is an interesting direction but out of scope for this paper; we use the neutral “Response A/B/C” template throughout.

B.2 Computational Cost

Because θ is a causal language model, the full a_k curve for one response ordering can be computed from a **single forward pass**. The prompt and all responses are concatenated into one sequence (Figure 2). The forward pass produces the log-probability of every token conditioned on all preceding tokens.

The tokens belonging to r_k are conditioned on exactly $p, r_1, \dots, r_{k-1}$ (plus formatting), which is the conditioning required for a_k . Partitioning the output log-probabilities by response boundary and summing within each partition yields all n values of a_k simultaneously for that ordering, with FLOPs scaling as $O((|p| + n\bar{L}_{\text{tok}})^2)$ from causal attention, where $\bar{L}_{\text{tok}}$ is the average response length

[prompt p] Response A: [r_1] Response B: [r_2] ... Response N: [r_n]

Figure 2: Single-pass input layout for computing the full a_k curve. The prompt is followed by all n responses with “Response A/B/C/...” labels; one forward pass over this sequence yields every a_k value simultaneously, since causal attention conditions the tokens of r_k on exactly $p, r_1, \dots, r_{k-1}$ (plus formatting).

in tokens. Because we report $\bar{a}_k$ averaged over $|\Sigma|$ random permutations of the response ordering (Section B.3), the total a_k -curve cost is $|\Sigma|$ such passes; the permutations are independent and we batch them on a single GPU.

The coherence term C additionally requires the *unconditional* per-byte cross-entropies $h_\theta(r_i | p)$ for each response. These cannot be extracted from the concatenated pass, since in that pass r_i for $i > 1$ is conditioned on all prior responses. Computing them requires n independent forward passes, each over the short context (p, r_i) . These passes are embarrassingly parallel and batchable, and they do not depend on the ordering, so C is computed once per response set regardless of $|\Sigma|$.

In summary, scoring one (prompt, response set) tuple takes:

- a_k curve: $|\Sigma|$ forward passes (long context, one per permutation).
- C : n forward passes (short contexts, batchable; not multiplied by $|\Sigma|$).

With $n = 10$ responses and $|\Sigma| = 50$ permutations (our default for larger-scale experiments), this is $50 + 10 = 60$ forward passes per (prompt, response set) tuple; the long-context $|\Sigma|$ passes dominate wall-clock time.

B.3 Dependence on Sample Ordering

The a_k values depend on the ordering of $\{r_i\}$. Individual responses differ in how surprising they are to θ given the responses that precede them, so a_k depends on which response sits at position k and which others precede it, making a single ordering’s curve jagged. Averaging over random permutations removes this dependence: each position averages over many choices of response and preceding context, so the curve reflects only how θ ’s predictions improve with more context. We therefore average over $|\Sigma|$ random permutations; Section C.5 compares low- and high-permutation runs on the validation scenarios and finds that the low setting misranks adjacent scenarios. We have not swept intermediate values and cannot pinpoint a cheap-and-reliable minimum, so we default to $|\Sigma| = 100$ for scenario-level experiments and $|\Sigma| = 50$ elsewhere.

Per-response normalization by byte count must be performed *before* averaging across permutations, since the byte count depends on which response lands at each position. Concretely, the per-byte curve is a mean of per-permutation rates,

$$\bar{a}_k = \frac{1}{|\Sigma|} \sum_{\sigma \in \Sigma} \frac{a_k^\sigma}{\|r_{\sigma(k)}\|}, \quad (6)$$

not a ratio of permutation-averaged bits to permutation-averaged byte counts (which would degenerate, after enough permutations, into the total-bits curve rescaled by the mean response length).

B.4 Choice of n

The diversity score depends on n as a measurement parameter, not through any estimate of an asymptotic floor. Given enough in-context examples, even genuinely diverse policies eventually become predictable to θ , so a_k keeps decreasing rather than converging to a meaningful irreducible value: there is no a_∞ for a_n to approximate. Too small an n means θ has not yet exploited the inter-response structure, and D_{Ca_n} overstates diversity. Too large an n means θ has accumulated enough in-context examples to reduce its surprise even within genuine diversity; for any policy with finitely many distinct modes, a_n continues to decrease toward zero as θ learns to predict within-mode variation from prior examples. For comparisons to be meaningful, n must be held fixed across all policies under evaluation. Our experiments use $n = 5$ on Tevet (Section 4), $n = 10$ on OLMo-2-7B and scenario validation, and $n = 20$ on the mode-count experiment (Appendix C.3). This is one reason to report the full a_k curve rather than D_{Ca_n} alone. D is a lossy summary, but the curve shows directly how response surprise evolves at every context size.

Table 5: Scenario validation metrics (mean across 5 prompts, 100 permutations, $n = 10$ responses, per-byte). We report several candidate scalars: unconditional surprise a_1 , the curve’s last point a_n , coherence C , our recommended score $D_{Ca_n} = C \times a_n$ (**bold**), the excess entropy E (Appendix E), $C \times E$, and the coherence spread σ_ℓ (the standard deviation of the per-byte cross-entropies $\{h_\theta(r_i | p)\}_{i=1}^n$). D_{Ca_n} correctly ranks multi-mode above one-mode and suppresses incoherent scenarios via C . Both E and $C \times E$ go negative for the mixed scenario on GPT-2. Mixed empirically scores the highest D_{Ca_n} on both models, exceeding multi-mode coherent rather than landing at the “mid” position Table 1 predicts; this is not the intended ranking. The coherence spread σ_ℓ flags this case in its intended diagnostic role (the largest σ_ℓ across scenarios on both models), reflecting the within-set heterogeneity between coherent and incoherent responses; reweighting variants such as $C^\alpha \times a_n$ are discussed in Section C.6.

Scenario	GPT-2 (124M)								Qwen2.5-3B							
	a_1	a_n	C	D_{Ca_n}	E	$C \times E$	σ_ℓ		a_1	a_n	C	D_{Ca_n}	E	$C \times E$	σ_ℓ	
Pure noise	7.52	6.94	0.005	0.04	0.68	0.004	0.20		7.18	6.68	0.007	0.05	0.75	0.005	0.11	
Multi-incoher.	3.26	1.85	0.10	0.19	5.94	0.61	0.48		2.36	1.40	0.21	0.29	3.12	0.64	0.69	
Multi-mode	1.06	0.53	0.48	0.26	1.61	0.76	0.08		0.78	0.31	0.59	0.19	1.17	0.68	0.11	
One-mode	1.06	0.22	0.48	0.10	1.31	0.63	0.06		0.80	0.16	0.58	0.09	0.92	0.53	0.08	
Mixed	1.90	2.02	0.26	0.52	−1.29	−0.31	1.17		1.13	0.91	0.46	0.41	0.52	0.23	0.73	

C Experiments

C.1 Experimental Setup

We evaluate the metric primarily using Qwen2.5-3B (3B parameters, 32K-token context window) [Yang et al., 2024], with GPT-2 (124M parameters, 1024-token context window) [Radford et al., 2019] as a smaller-model comparison point for scenario validation. Qwen2.5-3B is used for mode-count scaling, cross-mode learning, and external validation (Sections C.3, C.4, 4). The cross-model scaling study (Section C.6) additionally uses the full Llama 3 family for comparison.

All a_k curves are computed using single-pass forward computation (Section B.2). Unless otherwise noted, we use 100 permutations for scenario validation and 50 permutations for larger-scale experiments. Responses are formatted as described in Section B.2. All code and data, including the Olmo-2 RLHF generations under `results/rlhf_experiment/`, are publicly available.⁶

C.2 Scenario Validation

We construct five synthetic scenarios with known diversity structure:

1. **Pure noise:** Random ASCII characters (letters, digits, punctuation, spaces) with no learnable structure.
2. **Multi-incoherent:** Responses from 5 distinct modes, each internally incoherent (scrambled words within a template).
3. **Multi-mode:** Responses from 5 distinct modes, each internally coherent (e.g., recipe, poem, code).
4. **One-mode:** All responses from a single coherent mode (paraphrases of the same content).
5. **Mixed:** A mixture of coherent and incoherent responses.

Each scenario contains 10 responses per prompt, 5 prompts per scenario, with 100 permutations and seed 42.

Results. Table 5 summarizes the metrics. Figure 3 shows the a_k curves for all scenarios.

The coherence term C correctly separates coherent from incoherent scenarios: $C \approx 0.5$ – 0.6 for multi-mode and one-mode, $C \approx 0.1$ – 0.3 for multi-incoherent, and $C \approx 0.01$ for pure noise. The coherence spread σ_ℓ discriminates mixed from pure scenarios: $\sigma_\ell > 1.0$ for mixed (GPT-2), reflecting the within-set heterogeneity between coherent and incoherent responses.

⁶<https://anonymous.4open.science/r/icl-diversity-67E6/>




Figure 3: Progressive conditional surprise curves (a_k , per-byte) for all five scenarios, comparing GPT-2 (top row) and Qwen2.5-3B (bottom row). In each panel, faint colored curves are individual permutations (100 per prompt), medium colored curves are per-prompt averages, and the bold black curve is the mean across prompts. Pure noise curves are flat (no learnable structure); multi-mode curves show progressive decline; one-mode curves decline steeply then flatten.

622 C consistently improves with model strength (Qwen2.5-3B assigns higher coherence to coherent text
 623 than GPT-2). D_{Ca_n} correctly ranks multi-mode above one-mode on both models, and suppresses
 624 pure noise and multi-incoherent via the C factor.

625 C.3 Mode Count Scaling

626 To test whether the metric reflects the number of distinct modes, we construct response sets with
 627 $m \in \{1, \dots, 10\}$ modes drawn from a pool of 50 format-based generators (haiku, code, recipe, legal
 628 disclaimer, etc.), all responding to the same prompt (“Write a short piece about rain”). For Qwen2.5-
 629 3B: $n = 20$ responses, 1000 random draws of mode assignments, averaged over permutations.




Figure 4: Mode count scaling on Qwen2.5-3B ($n = 20$, 1000 draws). The a_k curves fan out with increasing m : higher floors, slower convergence. All curves are exponential (no sigmoidal plateau), even at $m = 10$, due to cross-mode learning (Section C.4). Shaded bands around each curve are ± 1 standard error of the mean across the 1000 random draws of mode assignments (i.e., the across-draw standard deviation divided by $\sqrt{1000}$); they are narrow because averaging over 1000 draws sharply reduces the across-draw spread.

Key findings. Figure 4 shows the a_k curves. a_n increases monotonically with m (9.3 bits at $m = 1$ to 77.8 bits at $m = 10$; see Table 6), reflecting that with more modes there are fewer same-mode repetitions within n responses, so each response remains more surprising even after θ has seen the others. This is exactly the behavior $D_{C_{a_n}} = C \times a_n$ needs: a_n tracks mode count directly. All curves are purely exponential (no sigmoidal plateau), even at $m = 10$; Section C.4 investigates why. For a detailed comparison of the excess-entropy-based scores on this data, including the sigmoid fit parameters and the non-monotonicity of $\hat{E}_n$, see Appendix E.

C.4 Cross-Mode Learning and Curve Shape

The authors predicted that at high m , the a_k curve should be sigmoidal: an initial plateau (early responses come from different modes and are uninformative about each other), followed by decline once mode repetitions accumulate. GPT-2 shows hints of this pattern at some m values (e.g. $m = 10$). Qwen2.5-3B does not: it shows immediate exponential decay at all m . This section investigates the discrepancy through systematic pairwise analysis.

Pairwise cross-mode surprise matrix. For 15 modes, we compute a 15×15 matrix M_{ij} : the surprise reduction (in bits) when a response from mode j is in the context and mode i is the target. Diagonal entries measure same-mode learning; off-diagonal entries measure cross-mode information transfer. Each entry averages over 5 context–target samples (using different samples for context and target to avoid self-prediction inflation).

Results. (Cell standard deviations are the mean within-cell sampling std across the 15 diagonal or 210 off-diagonal matrix cells, each estimated from 5 context-target sample pairs.)

- **Qwen2.5-3B:** diagonal mean = 60.5 ± 14.4 bits, off-diagonal mean = $+1.9 \pm 2.5$ bits, 64% of off-diagonal entries positive.
- **GPT-2:** diagonal mean = 83.0 ± 20.0 bits, off-diagonal mean = -3.7 bits, only 30% of off-diagonal entries positive.

The sign of the off-diagonal determines the curve shape (Figure 5):

- **Positive off-diagonal (Qwen):** Every response in context lowers expected surprise on every other response. The a_k curve drops from position 1, producing exponential decay.
- **Negative off-diagonal (GPT-2):** Cross-mode responses actively raise surprise on subsequent responses. At $m = 10$, the expected first-step net drop under the additive independent-mode model is only +4.9 bits (the 8.3-bit diagonal gain from the $1/m$ same-mode chance is largely offset by the -3.4 -bit cross-mode damage from the $(m-1)/m$ different-mode chance), producing a flat plateau until same-mode repetitions accumulate, the predicted sigmoid.

GPT-2 has a larger diagonal and a negative off-diagonal: it gains more from same-mode context but is actively penalised by cross-mode context. The two effects are linked, since probability mass that conditioning concentrates on same-mode continuations is mass taken away from other modes; a sharper conditional on the seen mode is necessarily worse on the unseen ones. Qwen does the opposite: a smaller diagonal but a positive off-diagonal, distributing its update across modes.

Pairwise asymmetry. The mean absolute asymmetry $|M_{ij} - M_{ji}|$ is 5.6 bits ($3.0 \times$ the off-diagonal mean), falsifying the hypothesis that pairwise surprise reduction is symmetric (Figure 6). Information-theoretically $I(X; Y) = I(Y; X)$, so this asymmetry cannot reflect asymmetric information structure between modes; it reflects θ ’s imperfection as an in-context reasoner. The same content is more useful as context than as target depending on its surface form, mirroring the *reversal curse* observed in autoregressive LLMs trained on “A is B” but failing on “B is A” [Berglund et al., 2023].

Connections to information-theoretic properties. Two information-theoretic properties are relevant to interpreting the pairwise matrix:

1. **Non-negative conditioning** ($H(X) \geq H(X | Y)$): Under the true distribution, conditioning can never increase entropy on average. For cross-entropies this is not guaranteed (the KL term can increase), but a well-calibrated θ should show mostly non-negative off-diagonal entries.

680 2. **Symmetry of mutual information** ($I(X; Y) = I(Y; X)$): Under the true joint, row means
681 (how informative mode i is as context) should correlate with column means (how much
682 mode i benefits from context).

683 Figure 7 compares the two models. Qwen shows a tighter row–column relationship (slope = 0.84,
684 $R^2 = 0.34$) with only 1 of 15 modes having a negative column mean. GPT-2 shows widespread
685 violations of non-negative conditioning (6 modes with negative column means) and a barely-existent
686 row–column relationship (slope = 0.35, $R^2 = 0.08$). Across the four Llama models [Grattafiori et al.,
687 2024, Meta AI, 2024] we additionally tested (Section C.6, Figure 9), the fraction of positive off-
688 diagonal entries rises monotonically with size (25% at 1B to 63% at 70B), and the off-diagonal mean
689 follows the same trend (−4.9 bits at 1B to +2.1 bits at 70B): the *magnitude* of cross-mode information
690 transfer increases with size. The row–column relationship does not: R^2 peaks at Llama-3B (0.67)
691 and degrades at the larger Llamas (0.09 at 70B), so the *symmetry* of the pairwise matrix does not
692 correspondingly improve.

693 **Token-level attribution.** Figure 8 shows per-token surprise reduction (in bits) for one pair from
694 each of three strata: same-mode (letter → letter), cross-mode where conditioning helps
695 (math_stats → json_data), and cross-mode where conditioning hurts (scientific_fact →
696 haiku). Pairs were selected by stratified-median total $|\Delta|$.

697 **Discussion.** The paper’s sigmoid prediction assumes modes are approximately independent. For
698 capable models, this fails: cross-mode learning creates positive information transfer between distinct
699 modes. The Llama 1B/3B/8B/70B series confirms this scales with model size (Section C.6): the
700 total cross-mode information θ extracts grows monotonically with size, and with it non-negative
701 conditioning is better satisfied. The matrix nevertheless remains asymmetric (row–column R^2 peaks
702 at Llama-3B and degrades at the larger Llamas), which under $I(X; Y) = I(Y; X)$ signals that even
703 Llama-70B is an imperfect in-context reasoner.

704 C.5 Practical Findings

705 **Permutation sensitivity.** The per-byte a_k curve depends on the response ordering. Comparing
706 the $D_{C_{a_n}}$ ranking of the 5 validation scenarios from Section C.2 between 3- and 100-permutation
707 runs: on GPT-2, 2 of 5 scenarios shift rank between the two settings; on Qwen2.5-3B, 2 of 5. Where
708 scenarios do shift rank, it is between adjacent scenarios with similar $D_{C_{a_n}}$ values rather than dramatic
709 reorderings. We use 100 permutations throughout for scenario-level experiments; we have not swept
710 intermediate values, so we cannot pinpoint a minimum that is both cheap and reliable. The coherence
711 term C is permutation-invariant by construction and unaffected.

712 **Boundary handling.** When Qwen merges a response’s final . with the following \n\n into a single
713 .\n\n token, our boundary detector (which uses a character-span overlap rule) attributes that token to
714 the response, not the separator. The response’s cross-entropy therefore includes a small contribution
715 from predicting the upcoming separator alongside its own characters; the byte-count denominator
716 remains the response’s literal byte length. We have not measured the magnitude of this bias on our
717 reported numbers. Boundary correctness is verified by tests/test_response_boundaries.py.

718 C.6 Discussion

719 **σ_ℓ is a diagnostic, not a diversity signal.** σ_ℓ increased monotonically with m in our synthetic
720 experiments because our modes have very different formats (code vs. haiku vs. recipe → different
721 per-byte cross-entropies). But σ_ℓ measures coherence *heterogeneity*, not diversity: modes with
722 similar fluency would have low σ_ℓ despite high diversity, and a single mode with high quality variance
723 could have high σ_ℓ with zero diversity. Its role is diagnostic (detecting mixed coherence), not as a
724 standalone diversity score.

725 **Cross-mode information scales with model quality.** We test how cross-mode information transfer
726 varies with θ ’s scale by running the pairwise matrix experiment on 4 dense Llama models (1B,
727 3B, 8B, 70B) using the same 15 modes and 5 samples per mode. Figure 9 shows the results. The
728 off-diagonal mean increases monotonically with model size: −4.9 (1B), −1.4 (3B), −1.2 (8B), +2.1
729 (70B) bits, transitioning from negative (cross-mode context hurts) to positive (cross-mode context

helps). The fraction of positive off-diagonal entries follows the same monotone trend: 25%, 41%, 44%, 63%. This hypothesis was pre-registered before running any Llama models, and the probability of accidental monotonicity across 4 models under the null is $1/4! \approx 4.2\%$.

The framework admits other metrics. The a_k curve is the primary object of our framework, and $D_{Ca_n} = C \times a_n$ is one natural summary of it among many. Reporting the curve directly preserves information that any single scalar (including ours) collapses. Alternative scalars could be derived from the same curve to capture different desiderata: weighted sums of $a_k - a_\infty$ (emphasizing early or late positions), the slope or curvature at a specific k (measuring how quickly θ learns), curve-shape descriptors, or coherence-weighted variants that combine C with quantities other than a_n . Different applications may warrant different choices; for instance, a policy-comparison setting might prefer a metric sensitive to *changes* in the curve shape, while a ranking setting needs a single well-behaved scalar. We have not explored this space systematically; we focus on D_{Ca_n} because it empirically works well across regimes and has a clean information-theoretic interpretation, but we expect the framework to support a family of related metrics with complementary properties. One untested family is coherence-reweighting: a practitioner with a clear preference for sharper Mixed penalisation could replace $C \times a_n$ with $C^\alpha \times a_n$ for $\alpha > 1$ to drive low-coherence sets toward zero faster. We have not evaluated such variants on the Tevet or OLMo experiments and so make no empirical claim about them.

Bits/byte vs. total bits, and the length-matching consequence. A separate choice within the framework is whether to report total bits (length-dependent but tokenizer-agnostic) or per-byte bits/byte (also tokenizer-agnostic, with length scaled out). We default to bits/byte for the headline diversity score, but per-byte quantities are not invariant to response length: in a causal LM, longer responses give θ more within-response context to predict later tokens, lowering per-byte surprise. For the OLMo-2-7B experiment (Section 5) raw response length is not monotone across stages, so we score the length-matched subset as the headline analysis: every response in a (stage, prompt) tuple is truncated to a common per-prompt UTF-8 byte length (the minimum across all 40 responses) before scoring. 111/300 prompts had a per-prompt common length below 50 bytes – driven by the base model on AlpacaEval (no instruction-following prior) and by SFT on NB-curated short-answer prompts – and were dropped because length-matching them would compare essentially-empty fragments. The headline numbers therefore condition on the 150/200 AlpacaEval and 39/100 NB-curated prompts where length-matching is well-defined; the dropped prompts are themselves a stage-specific behavioural signal we are not measuring here. We also verified that the monotone D drop holds on the un-truncated 200 + 100 prompt sets, with all three pre-registered H1 contrasts remaining Bonferroni-significant at $p < 10^{-13}$ on both prompt sets. Full numbers and reproduction commands are in `investigations/length_matched_rlh/VERDICT.md`.

D Future Work

Several directions remain unexplored.

Prompt format optimization. We use a fixed prompt template throughout and do not attempt to optimize it for ICL elicitation. Prompt engineering could improve the metric’s discriminative power. Dimension-specific framings (e.g., “Here are several stories in different genres”) are one such variant: priming θ to expect variation along a dimension removes the small portion of the surprise attributable to the mere existence of that variation, but it does not isolate per-dimension diversity, since the rest of the cross-response surprise reduction remains.

Ensembling base models. A single base model θ may produce non-monotone a_k curves when pushed out of distribution by long contexts. Ensembling multiple base models at the token level (averaging softmax probabilities) could stabilize the estimates while preserving the autoregressive structure. Our codebase supports token-level ensembling, but we have not yet validated it experimentally.

Disagreement cases with other metrics. We report aggregate correlations with embedding-based and reference metrics on Tevet’s tasks (Section 4) but do not look at *which* response sets the metrics rank differently. Identifying prompts or response sets where $C \times a_n$ disagrees with SentBERT, distinct- n , or McDiv is the most direct way to characterise what each metric is actually measuring; the

disagreements are where the choice of metric matters. We expect two regimes to drive disagreement: response sets that share arbitrary rare patterns a base model can pick up on but that fall outside what a fixed sentence embedder represents (where $C \times a_n$ should detect repetition that embeddings miss), and response sets where one response is a mixture of two others (where progressive conditioning credits the mixture for being predictable from its components, while pairwise embedding-distance averages do not).

Total-bits variants of D . We report $D_{Ca_n} = C \times a_n$ in bits/byte for length control, normalising by bytes rather than tokens so that scores remain tokenizer-agnostic. The implementation also exposes a_n in total bits, which is likewise tokenizer-agnostic but not length-controlled. A hybrid scalar $C \times a_n^{\text{bits}}$ rewards length (longer responses have more positions at which to be surprising), which some practitioners may want; characterising when each variant is the appropriate measurement target is left to future work.

Optimizing generators against D . We use $C \times a_n$ as an evaluation-time diagnostic, but it could in principle serve as a training or decoding objective. Future work could fine-tune or sample from a generator π with the goal of maximizing $C \times a_n$ under a fixed base model θ to push π toward response sets that are both coherent and diverse.

E Excess Entropy and the $C \times E$ Score

This appendix develops the excess entropy E , an alternative summary of the a_k curve that is theoretically interesting (connecting to the excess entropy of computational mechanics [Crutchfield and Feldman, 2003] and to the total correlation) but empirically inferior to D_{Ca_n} on the external benchmarks of Section 4. We tried both the last-point estimator $\hat{E}_n$ and a sigmoid-extrapolated projected floor for a_∞ (Section E.1); we stopped using E when it became clear it cannot distinguish diverse from non-diverse response sets in the few-draws regime where modes do not repeat within n samples. We include it here for theoretical completeness, as a null result, and to explain the empirical comparisons that motivate our preference for D_{Ca_n} .

E.1 The Excess Entropy

The primary metric $D_{Ca_n} = C \times a_n$ uses the curve’s endpoint a_n : as the surprise after conditioning on $n - 1$ peer responses, a_n reflects both the curve’s level and the drop from a_1 in a single scalar. Section 4 shows this combination is what discriminates diverse from non-diverse response sets in the few-draws regime, while a_1 alone (the level) and E (the integrated drop, defined below) each underperform.

We thought that E ’s learnable structure would be interesting and track diversity. Following the concept of excess entropy from computational mechanics [Crutchfield and Feldman, 2003], define the **excess entropy**:

$$E = \sum_{k=1}^{\infty} (a_k - a_\infty) \quad (7)$$

Units: bits. This is the total learnable structural information in the progressive conditioning process, above the irreducible residual noise. It converges whenever θ ’s surprise at new responses eventually stabilizes (since $a_k \rightarrow a_\infty$ and the excess $e_k = a_k - a_\infty$ decays to zero). The “structure” captured by E is not limited to discrete mode identity: it includes any inter-response regularity to which θ assigns lower conditional surprise.

In practice, a_∞ is unknown. We estimate it as a_n (the last observed value) and compute:

$$\hat{E}_n = \sum_{k=1}^n (a_k - a_n) \quad (8)$$

This underestimates E (since $a_n \geq a_\infty$), but the bias decreases with n .

822 **Parametric extrapolation of a_∞ .** Rather than using a_n directly, one can fit a parametric model
 823 to the observed curve and extrapolate a_∞ . The a_k curve is theoretically expected to be sigmoidal
 824 (with concave-up/exponential decay as the degenerate case when the inflection point $k_0 < 1$). This
 825 motivates fitting a four-parameter sigmoid:

$$a_k = a_\infty + \frac{\alpha}{1 + e^{\beta(k-k_0)}} \quad (9)$$

826 where a_∞ is the asymptotic floor, α is the total drop from a_1 to a_∞ , $\beta > 0$ controls the steepness
 827 of the transition, and k_0 is the inflection point. The fit yields a_∞ without requiring n to be large
 828 enough for convergence. Section C.3 provides empirical evidence that the sigmoid-extrapolated E_{fit}
 829 recovers the expected monotonic relationship between mode count and excess entropy, while the raw
 830 $\hat{E}_n$ estimator does not.

831 **Relationship to total correlation.** The excess entropy and total correlation are complementary
 832 decompositions of the same underlying structure. Working in expectation over i.i.d. draws from π , let
 833 $\bar{a}_1 = \mathbb{E}[-\log_2 \theta(r \mid p)]$ denote the expected unconditional cross-entropy, $\bar{a}_k = \mathbb{E}[a_k]$ the expected
 834 conditional cross-entropy at position k , and $e_k = \bar{a}_k - a_\infty$ the excess at step k . The per-step mutual
 835 information is $I_k = \bar{a}_1 - \bar{a}_k = (\bar{a}_1 - a_\infty) - e_k$. Summing:

$$\text{TC}_n = \sum_{k=1}^n I_k = n(\bar{a}_1 - a_\infty) - E_n \quad (10)$$

836 For large n where $E_n \rightarrow E$, the total correlation grows linearly with slope $(\bar{a}_1 - a_\infty)$: $\text{TC}_n \approx$
 837 $n(\bar{a}_1 - a_\infty) - E$. Figure 10 illustrates. The three quantities have clean interpretations, all in bits:
 838 $(\bar{a}_1 - a_\infty)$ is the per-response redundancy once θ has fully learned the available structure; E is
 839 the total finite structural information, consumed during θ 's transient learning phase; TC_n is the
 840 cumulative redundancy, which grows without bound.

841 E.2 Per-Byte Excess Entropy Rate

842 The excess entropy $E = \sum_k (a_k - a_\infty)$ is in bits. For a length-normalized variant, we define the
 843 **per-byte excess entropy rate** by normalizing each response's surprise by its byte count *before*
 844 averaging across permutations. For a single permutation σ , the per-byte conditional surprise at
 845 position k is $a_k^\sigma / \|r_{\sigma(k)}\|$ (bits/byte). Averaging over permutations gives

$$\hat{a}_k^{\text{rate}} = \mathbb{E}_\sigma \left[\frac{a_k^\sigma}{\|r_{\sigma(k)}\|} \right] \quad (11)$$

846 and the per-byte floor is estimated analogously from the last position: $\hat{a}_\infty^{\text{rate}} = \hat{a}_n^{\text{rate}}$. The per-byte
 847 excess entropy rate is

$$E_{\text{rate}} = \sum_{k=1}^n (\hat{a}_k^{\text{rate}} - \hat{a}_\infty^{\text{rate}}). \quad (12)$$

848 This treats each response equally regardless of length. Note that $E_{\text{rate}} \neq E/\bar{B}$ in general: because
 849 total surprise a_k and byte count $\|r_k\|$ are positively correlated, $E/\bar{B}$ overweights long responses,
 850 while the per-response normalization in E_{rate} avoids this.

851 E.3 The $C \times E$ Score

852 Combining E with the coherence term C (Section 3.2) yields scalar scores

$$C \times E \quad (\text{bits}) \quad \text{or} \quad C \times E_{\text{rate}} \quad (\text{bits/byte}). \quad (13)$$

853 The first measures total structural information, and the second normalizes per byte so long and short
 854 responses are treated equally.

855 E.4 Why Not Weight Excess Entropy Inside the Sum?

856 A natural alternative to reporting E and C separately is to weight the terms within E by coherence,
 857 computing $E_w = \sum_k w_k \cdot (a_k - a_\infty)$ where $w_k = 2^{-h_\theta(r_k|p)}$. This would suppress contributions

from incoherent modes directly at the point of measurement. We considered and rejected this approach because it breaks the information-theoretic interpretation. $E = \sum_k e_k$ has a clean meaning as the total structural information above the noise floor, connected to Crutchfield and Feldman’s excess entropy, and inserting weights makes E_w a hybrid quantity that is not the excess entropy of any well-defined process.

E.5 Limitations of $C \times E$

$C \times E$ inherits two limitations from E . First, E measures *recurrence*: surprise reduction as multiple responses accumulate, rather than continuous spread. A policy that draws from a broad, roughly continuous distribution of coherent outputs has $a_k \approx a_1$ for all k , yielding $E \approx 0$ despite maximal diversity. Second, E is near-zero in the few-draws regime: with only 5–10 responses and no repeated modes, θ finds little learnable structure regardless of diversity, so $C \times E$ carries almost no signal. These limitations are the reason $D_{C a_n}$ is preferred as the primary score: a_∞ is the *level* of the floor, which differs meaningfully between diverse and non-diverse response sets even when the curve stays nearly flat.

Empirical inferiority. On Tevet’s diversity-eval benchmark with only 5 responses per set (Section 4), $C \times a_n$ achieves ROC AUC of 0.867 on McDiv_nuggets prompt_gen while the $D_{\text{fit}} = C \times E_{\text{fit}}$ variant cannot be computed at all (the four-parameter sigmoid fails to converge on only five points) and the discretized D_{disc} variant is anti-correlated with diversity (AUC = 0.303, well below chance). Figure 11 shows the ROC curves.

Why E fails in this regime. With only 5 responses and no repeated modes, θ finds little learnable structure regardless of diversity. The a_k curve stays approximately flat, yielding $E \approx 0$ for both diverse and non-diverse sets. The E -based scores thus carry almost no signal. The asymptotic floor a_∞ , by contrast, differs substantially between the two conditions: diverse responses remain surprising after conditioning (a_∞ high), while paraphrases become predictable (a_∞ low). This is the fundamental reason $D_{C a_n}$ is preferred.

E.6 Mode Count Scaling: Excess-Entropy Metrics

In the synthetic mode count experiments (Section C.3), the sigmoid-extrapolated E_{fit} is monotonic in mode count for Qwen2.5-3B but has wide confidence intervals at high m , while the raw $\hat{E}_n$ estimator is non-monotonic. Table 6 and Figure 12 show the details.

Table 6: Mode count scaling metrics for Qwen2.5-3B ($n = 20$, 1000 draws). E_{fit} is the sigmoid-extrapolated excess entropy; a_n is the mean floor at $k = 20$; σ_ℓ is per-byte coherence spread; k_0 is the sigmoid inflection point. Confidence intervals are 95% bootstrap. The high- m rows ($m \geq 8$) have E_{fit} CIs spanning more than half the mean and should be read as preliminary; the trend direction is robust but the point values are imprecise.

m	E_{fit} (bits)	a_n (bits)	σ_ℓ (bits/byte)	k_0
1	296 ± 11	9.3	0.1	−10.0
2	568 ± 24	16.9	0.2	−10.0
3	704 ± 40	26.7	0.2	−10.0
4	799 ± 67	37.4	0.3	−10.0
5	855 ± 113	48.0	0.3	−10.0
6	940 ± 127	56.1	0.3	−10.0
7	1069 ± 194	61.5	0.3	−10.0
8	1138 ± 253	66.9	0.3	−10.0
9	1563 ± 705	72.8	0.3	−10.0
10	1807 ± 1194	77.8	0.3	−10.0

F Dataset Construction Confounds in McDiv

When validating against the McDiv_nuggets benchmark [Tevet and Berant, 2021], we observed that the mean $\hat{a}_1$ curve (per-byte) for the “low diversity” group starts *above* the “high diversity” group in

the story_gen task (Figure 14). Since a_1 measures the surprise of the first response conditioned only on the prompt (before any other responses appear in context), this gap cannot reflect diversity. It is a confound in the dataset construction.

F.1 Mechanism

The McDiv protocol (Tevet and Berant §6.4), from which McDiv_nuggets is sampled (Tevet and Berant Appendix C.2 specifies McDiv_nuggets as the 3K subset of McDiv on which distinct- n correlates at zero), pairs a low-diversity set with each high-diversity set as follows. Workers first write five *different* continuations (the high-diversity set). The same worker is then asked to *self-select one* of their own five responses and paraphrase it five times, preserving content while varying form (the low-diversity set). Tevet and Berant do not characterize the distribution of which responses workers self-select.

Empirically (see Table 7 and Figure 14, and the illustrative examples from the McDiv_nuggets story_gen dataset in Section F.3 below), we observe that the self-selected endings for the low-diversity sets tend to be more specific or dramatic (e.g., “He scored winner” or “they quit”) than the typical high-diversity continuations (e.g., “Joel fired the cook when things went too far downhill”). This means that individual low-diversity responses are intrinsically more surprising to the base model, not because of diversity, but because of the specificity of the endings workers self-selected for paraphrasing.

F.2 Evidence

Table 7 summarizes the gap. The per-byte a_1 difference is substantial (a content effect) while the total-bits difference is small because high-diversity responses are on average several bytes longer.

Table 7: Unconditional surprise (a_1) by diversity label, McDiv_nuggets story_gen.

	High diversity	Low diversity	Gap
a_1 (bits/byte)	0.977	1.154	+0.177
a_1 (total bits)	49.1	50.5	+1.4
Mean response bytes	52.7	46.5	−6.3
n	103	97	

The gap persists after binning by response length (Table 8), confirming it is primarily a content effect rather than a length artifact: the high-minus-low per-byte gap remains positive across the bulk of the length range.

Table 8: Unconditional per-byte surprise by response length bin, story_gen.

Byte range	High div. mean	Low div. mean	Gap
[20, 40)	1.172	1.340	+0.167
[40, 60)	0.878	1.019	+0.141
[60, 80)	0.853	1.029	+0.176
[80, 120)	0.905	0.908	+0.003

F.3 Illustrative Examples

Figures 15 and 16 show per-token surprise (in bits) for the first response of a high-diversity and low-diversity sample, respectively. Samples are drawn by ranking each label group by per-byte a_1 (ascending for high-diversity, descending for low-diversity) and taking the sample at the $\lfloor N/4 \rfloor$ position, chosen to be representative of the expected confound direction rather than an extreme outlier; the specific pinned IDs (sa_00698, sa-b_00279) are fixed in scripts/dataset_confound_analysis.py for reproducibility. Red bars are measured response tokens; grey bars are masked context tokens.

High diversity (Figure 15). Context: “Joel owned a restaurant. He hired a cook that didn’t care about his work. The cook didn’t clean after himself. The kitchen was a mess.” Response: “Joel fired

923 *the cook when things went too far downhill.* This is a natural, predictable continuation ($a_1 = 0.783$
924 bits/byte).

925 **Low diversity** (Figure 16). Context: *“Harry and his basketball team was losing the game. The*
926 *coach called for a timeout. Harry boosted morale by talking to his team. The team caught up and the*
927 *game was tied.”* Response: *“He scored winner.”* This is a specific, somewhat ungrammatical ending
928 ($a_1 = 1.777$ bits/byte). All five responses in this set are paraphrases of the same idea (“Harry scored
929 the winning point”).

930 F.4 Implications

931 The construction confound (low-diversity sets are paraphrases of specific dramatic endings rather
932 than a random subsample of low-diversity content) lifts the entire $\bar{a}_k$ curve on low-diversity sets,
933 since their content is intrinsically surprising under θ at every k , and pushes C down because the
934 same content is less coherent. a_1 ’s upward shift opposes the diversity label; C ’s downward shift
935 happens to align with it (a confound artifact, not coherence genuinely tracking diversity). At a_n ,
936 the diversity signal we are measuring overrides the confound’s upward pull: paraphrases collapse to
937 predictable once θ has seen the others while genuinely diverse responses do not, so a_n ends up lower
938 on low-diversity sets and $C \times a_n$ tracks the label despite the confound (Section 4).

939 G Aggregation Across Prompts

940 This appendix concerns aggregation across prompts. Aggregation *within* a single prompt, across
941 permutations of the response ordering, is described in Section B.3 and given by Eq. (6): the per-
942 byte $\bar{a}_k$ is a mean of per-permutation per-byte rates, not a ratio of permutation-averaged bits to
943 permutation-averaged byte counts.⁷

944 The diversity score D_{Ca_n} is defined relative to a single prompt p . To summarize a policy’s diversity
945 across a prompt suite $P = \{p_1, \dots, p_M\}$, one can average:

$$\bar{D}_{Ca_n}(\pi) = \frac{1}{M} \sum_{j=1}^M D_{Ca_n}(\pi, p_j) \quad (14)$$

946 To compare two policies π_1 and π_2 , the natural scalar is the difference:

$$\Delta D_{Ca_n} = \bar{D}_{Ca_n}(\pi_1) - \bar{D}_{Ca_n}(\pi_2) \quad (15)$$

947 ΔD_{Ca_n} answers “how much more plausibility-weighted residual diversity does π_1 have than π_2 , on
948 average across prompts?” No division is involved, so ΔD_{Ca_n} is stable even when either policy’s
949 score is near zero.

950 H Scaling the Base Model: Qwen2.5-3B vs Qwen3-30B-A3B-Base

951 To test whether a stronger base model improves the metric on Tevet’s benchmarks, we re-ran the full
952 evaluation using Qwen3-30B-A3B-Base [Yang et al., 2025] (a 30B-parameter mixture-of-experts
953 model with $\sim 3B$ active parameters per token) via the Tinker API,⁸ with the same setup as in Section 4:
954 completion-format prompting, 50 permutations, no fine-tuning. Table 9 reports per-task $C \times a_n$
955 (per-byte) for both base models.

956 **Result: scaling did not help on the binary tasks.** On the 12 binary classification tasks (McDiv,
957 McDiv_nuggets, ConTest), Qwen2.5-3B wins on 11 of 12, with ConTest prompt_gen the only
958 marginal win for Qwen3-30B-A3B (+0.005 AUC). The mean degradation from scaling is small but

⁷An earlier revision of our implementation computed the ratio-of-means form, in which the bits and byte counts at each position are each averaged across permutations before dividing. After enough permutations, the byte counts at each position approach the mean response length, so that estimator collapses into the total-bits curve rescaled by a single constant, losing the per-permutation per-byte signal that Eq. (6) preserves. Switching the implementation to the form in Eq. (6) shifts the headline numbers in this paper by 0–17% (in mean-of-ratios’ favor), with no qualitative change to any ranking we report.

⁸<https://thinkingmachines.ai/tinker>

Table 9: Per-task comparison of $C \times a_n$ (per-byte) for Qwen2.5-3B vs Qwen3-30B-A3B-Base on Tevet diversity-eval. ΔAUC is positive when the larger model wins. ROC AUC is reported for the binary tasks (McDiv, McDiv_nuggets, ConTest); only Spearman ρ is meaningful for DecTest (continuous temperature labels).


Section	Task	Qwen2.5-3B		Qwen3-30B-A3B		ΔAUC
		ρ	AUC	ρ	AUC	
McDiv	prompt_gen (no_hds)	0.729	0.921	0.718	0.915	-0.006
	resp_gen (no_hds)	0.500	0.788	0.483	0.779	-0.009
	story_gen (no_hds)	0.523	0.802	0.493	0.785	-0.017
McDiv_nuggets	prompt_gen (no_hds)	0.636	0.867	0.615	0.855	-0.012
	prompt_gen (with_hds)	0.683	0.895	0.655	0.878	-0.017
	resp_gen (no_hds)	0.345	0.699	0.325	0.687	-0.012
	resp_gen (with_hds)	0.437	0.753	0.425	0.746	-0.007
	story_gen (no_hds)	0.317	0.683	0.293	0.669	-0.014
	story_gen (with_hds)	0.405	0.734	0.376	0.717	-0.017
ConTest	prompt_gen (with_hds)	0.584	0.837	0.592	0.842	+0.005
	resp_gen (with_hds)	0.391	0.726	0.333	0.692	-0.034
	story_gen (with_hds)	0.686	0.896	0.653	0.877	-0.019
DecTest	prompt_gen (no_hds)	0.842	—	0.877	—	—
	prompt_gen (with_hds)	0.845	—	0.875	—	—
	resp_gen (no_hds)	0.771	—	0.804	—	—
	resp_gen (with_hds)	0.760	—	0.777	—	—
	story_gen (no_hds)	0.763	—	0.771	—	—
	story_gen (with_hds)	0.785	—	0.785	—	—

consistent (mean $\Delta\text{AUC} = -0.013$ across the 12 binary tasks). On the 6 DecTest tasks (continuous temperature labels), Qwen3-30B-A3B is mildly better on 5 of 6 (mean $+0.021$ in Spearman ρ , with one tie).


Interpretation. The binary result is a clear negative: scaling from 3B to 30B active parameters does not improve discrimination on McDiv or ConTest. We do not have a confident explanation. The Qwen3-30B-A3B run was not repeated after a bug fix applied to the primary experiments, so the comparison should be treated as preliminary. We leave the question of whether larger base models improve the metric on binary tasks to future work.

I Cross-Metric Agreement on the OLMo-2-7B RLHF Experiment

The RLHF case study (Section 5) reports the headline monotone D_{Ca_n} drop in Table 4. Figure 17 below visualises the underlying $\bar{a}_k$ curves and per-prompt D_{Ca_n} distributions for the same length-matched subset, and Figure 18 plots D_{Ca_n} against the lexical (EAD) and semantic (SentBERT) baselines, confirming that the three metrics see the same diversity-loss signal.



(a) Qwen2.5-3B: positive off-diagonal.



(b) GPT-2: negative off-diagonal.

Figure 5: Pairwise cross-mode surprise reduction matrices. Each cell (i, j) shows how many bits of surprise reduction mode i (target) receives from seeing mode j (context). Qwen shows diagonal dominance with pervasive positive off-diagonal (+1.9 bits mean); GPT-2 shows diagonal dominance with negative off-diagonal (-3.7 bits mean).






Figure 6: Pairwise symmetry scatter for Qwen2.5-3B: each point is one (i, j) pair, plotting M_{ij} vs. M_{ji} . Under symmetric mutual information, points would lie on $y = x$. The best-fit slope is 0.46 ($R^2 = 0.24$), and many pairs have opposite signs: mode i can reduce surprise for mode j while j *increases* surprise for i .



(a) Qwen2.5-3B: slope = 0.84, $R^2 = 0.34$.



(b) GPT-2: slope = 0.35, $R^2 = 0.08$.

Figure 7: Row mean (how informative as context) vs. column mean (how much benefit from context) for each mode. Qwen shows tighter correlation closer to the identity, consistent with approximate symmetry of mutual information. GPT-2 shows widespread violations of non-negative conditioning (modes in the negative-row-mean region).




Figure 8: Per-token surprise reduction (bits) for three representative context–target pairs (Qwen2.5-3B). Bars are mean reductions across 5 context samples per pair; error bars are ± 1 standard deviation. Green bars indicate the conditioning context lowered the per-token surprise; red bars indicate it raised it. The three pairs are stratified-median selections (by total $|\Delta|$) from the same-mode, positive cross-mode, and negative cross-mode strata of a 9-pair pool. Note how in scientific fact to haiku, the early newline triggers the largest surprise. Scientific facts may discuss "still water", but they almost never have a line with three words in it.




Figure 9: Cross-mode information transfer scales with model size. **Left:** Mean off-diagonal surprise reduction (with bootstrap 95% CI) transitions from negative to positive as models grow, indicating that larger models extract more information from cross-mode context. **Right:** The fraction of mode pairs showing positive cross-mode transfer increases monotonically across the 4 Llama models (blue; $p = 4.2\%$ under monotone null). GPT-2 and Qwen2.5-3B (gray) are included for context but differ in architecture.




Figure 10: Decomposition of the gap between unconditional surprise $\bar{a}_1$ and the asymptotic floor a_∞ at each step k . The per-step gap $\bar{a}_1 - a_\infty$ splits into mutual information $I_k = \bar{a}_1 - \bar{a}_k$ (red, above the curve) and excess $e_k = \bar{a}_k - a_\infty$ (blue, below the curve). Summing across k : the red area is the total correlation TC_n ; the blue area is the excess entropy E . As k grows, $e_k \rightarrow 0$ and each step contributes the full $\bar{a}_1 - a_\infty$ to TC_n , so TC_n grows linearly while E converges. All quantities are in bits. This figure is theoretical: the sigmoidal $\bar{a}_k$ shape reflects our initial expectation, but we subsequently found that the empirical curves are exponential decay without an initial plateau (see Section C.4).




Figure 11: ROC curves for McDiv_nuggets binary classification (Qwen2.5-3B, 50 permutations). Five per-byte metrics are overlaid per panel: $C \times a_n$ (red, solid), a_n alone (orange, dashed), $D_{\text{fit}} = C \times E_{\text{fit}}$ (blue, solid), $D_{\text{disc}} = C \times \hat{E}_n$ (cyan, dashed), and a_1 (gray, dotted). Line style is redundant with color so the encoding remains legible under common color-vision deficiencies. $C \times a_n$ dominates the other four scores on every task; the E -based scores ($D_{\text{fit}}, D_{\text{disc}}$) hug or fall below the chance diagonal in this 5-response regime.




Figure 12: Summary metrics vs. mode count m (Qwen2.5-3B, $n = 20$, 1000 draws). E_{fit} (sigmoid-extrapolated) increases monotonically, while the raw $\hat{E}_n$ does not.

Sigmoid Fit Parameters vs Mode Count — Qwen/Qwen2.5-3B




Figure 13: Sigmoid fit parameters vs. mode count m (Qwen2.5-3B). The inflection point k_0 remains at the lower bound (-10) across all m , indicating pure exponential decay without an initial plateau (see Section C.4).




Figure 14: Distribution of a_1 (unconditional surprise of the first response) for high- vs. low-diversity story_gen samples. Left: per-byte. Right: total bits. The per-byte gap (see Table 7) is a dataset construction artifact.




Figure 15: Per-token surprise for a high-diversity sample’s first response. The continuation (“Joel fired the cook...”) is predictable, yielding low per-token surprise across response tokens.






Figure 16: Per-token surprise for a low-diversity sample’s first response. The specific ending (“He scored winner”) is inherently more surprising to the base model, despite being labeled “low diversity.”




(a) AlpacaEval $\bar{a}_k$ curves (length-matched).

(b) Per-prompt D_{Ca_n} by stage (length-matched).

Figure 17: Progressive conditional surprise curves and per-prompt diversity distributions across the four OLMo-2-7B stages on the length-matched AlpacaEval subset. Each later stage’s $\bar{a}_k$ curve lies below the base curve at every $k \geq 2$; the per-prompt D_{Ca_n} distribution shifts toward lower values as the pipeline advances.



(a) AlpacaEval (length-matched, $n = 150$ prompts).



(b) NoveltyBench curated (length-matched, $n = 39$ prompts).

Figure 18: Per-prompt D_{Ca_n} versus EAD (left subpanel) and SentBERT-similarity diversity (right subpanel), coloured by stage, on the length-matched subset of prompts. Length-matching truncates each (stage, prompt) tuple’s responses to a common per-prompt byte budget so the per-byte conditional surprise that defines a_n is not depressed by response length. Pearson r and Spearman ρ (two-sided p) are reported in each panel. The rank correlations are positive across both prompt sets and both baselines: D_{Ca_n} tracks the same diversity-loss signal EAD and SentBERT detect, and the agreement is not an artefact of response length.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Section 1 state that we (i) define the per-byte a_k progressive-surprise curve and the scalar $D_{Ca_n} = C \times a_n$, (ii) evaluate it on Tevet & Berant’s McDiv human-grounded benchmark, and (iii) apply it to the OLMo-2-7B post-training pipeline; Section 4 and Section 5 report exactly these results, with limitations (Section 6) flagging where the metric trails embedding baselines.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 is dedicated to limitations: measurement is relative to the base model θ , cross-mode behaviour is model-dependent, the metric form $D = C \times a_n$ was selected with full visibility into the Tevet benchmark, performance trails the strongest embedding baseline, the prompt format is fixed, and the per-byte denominator forces length-matching that drops short-response prompts.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best

1024 judgment and recognize that individual actions in favor of transparency play an impor-
1025 tant role in developing norms that preserve the integrity of the community. Reviewers
1026 will be specifically instructed to not penalize honesty concerning limitations.

1027 3. Theory assumptions and proofs

1028 Question: For each theoretical result, does the paper provide the full set of assumptions and
1029 a complete (and correct) proof?

1030 Answer: [N/A]

1031 Justification: The contribution is empirical: an estimator and a benchmark study, not a
1032 theorem. The information-theoretic identities we draw on (excess entropy, entropy rate) are
1033 standard and cited; informal derivations connecting the a_k curve to those quantities are in
1034 Appendix E.

1035 Guidelines:

- 1036 • The answer [N/A] means that the paper does not include theoretical results.
- 1037 • All the theorems, formulas, and proofs in the paper should be numbered and cross-
1038 referenced.
- 1039 • All assumptions should be clearly stated or referenced in the statement of any theorems.
- 1040 • The proofs can either appear in the main paper or the supplemental material, but if
1041 they appear in the supplemental material, the authors are encouraged to provide a short
1042 proof sketch to provide intuition.
- 1043 • Inversely, any informal proof provided in the core of the paper should be complemented
1044 by formal proofs provided in appendix or supplemental material.
- 1045 • Theorems and Lemmas that the proof relies upon should be properly referenced.

1046 4. Experimental result reproducibility

1047 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
1048 perimental results of the paper to the extent that it affects the main claims and/or conclusions
1049 of the paper (regardless of whether the code and data are provided or not)?

1050 Answer: [Yes]

1051 Justification: Section 3 fully specifies the metric (prompt template, permutation-averaged
1052 single-pass forward computation, per-byte cross-entropy with explicit byte-count denomi-
1053 nators); Section 4 and Section 5 state the base model (θ), permutation count, response set
1054 sizes, and length-matching protocol used in each experiment; the anonymous code mir-
1055 ror at <https://anonymous.4open.science/r/icl-diversity-67E6/> contains the
1056 exact run scripts (scripts/run_scenarios.py, scripts/analyze_c_ainf.py) that
1057 produced every reported number.

1058 Guidelines:

- 1059 • The answer [N/A] means that the paper does not include experiments.
- 1060 • If the paper includes experiments, a [No] answer to this question will not be perceived
1061 well by the reviewers: Making the paper reproducible is important, regardless of
1062 whether the code and data are provided or not.
- 1063 • If the contribution is a dataset and/or model, the authors should describe the steps taken
1064 to make their results reproducible or verifiable.
- 1065 • Depending on the contribution, reproducibility can be accomplished in various ways.
1066 For example, if the contribution is a novel architecture, describing the architecture fully
1067 might suffice, or if the contribution is a specific model and empirical evaluation, it may
1068 be necessary to either make it possible for others to replicate the model with the same
1069 dataset, or provide access to the model. In general, releasing code and data is often
1070 one good way to accomplish this, but reproducibility can also be provided via detailed
1071 instructions for how to replicate the results, access to a hosted model (e.g., in the case
1072 of a large language model), releasing of a model checkpoint, or other means that are
1073 appropriate to the research performed.
- 1074 • While NeurIPS does not require releasing code, the conference does require all submis-
1075 sions to provide some reasonable avenue for reproducibility, which may depend on the
1076 nature of the contribution. For example

- (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
- (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Anonymous code and data mirror at <https://anonymous.4open.science/r/icl-diversity-67E6/> contains the full implementation (src/icl_diversity/), all experiment scripts, the released OLMo-2-7B generations under results/r1hf_experiment/, and command-line invocations in CLAUDE.md that reproduce every reported scalar and figure.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: We do not train models; the metric is a forward-pass-only estimator over a base θ . Section 3 specifies the response-conditioning template, permutation count ($|\Sigma| = 50$ at scale, $|\Sigma| = 20$ for synthetic studies), choice of θ (Qwen2.5-3B base for headline; ablations in Appendix H), and the length-matching protocol used in the OLMo-2-7B experiment.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: a_k curves are reported with permutation-averaged means and standard-deviation bands across response orderings (Section 3). The OLMo-2-7B comparison reports per-prompt $D_{C_{a_n}}$ means with standard deviations and pre-registered Bonferroni-corrected paired tests for the H1a–H1c stage contrasts, with the H1' comparison explicitly marked as exploratory (Section 5, Table 4).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The metric requires only forward passes through a base language model. Section B.2 reports the cost as $|\Sigma| + n$ forward passes per (prompt, response set). Local experiments (Qwen2.5-3B base, GPT-2) ran on a workstation with $2 \times$ Quadro RTX 8000 GPUs (48 GB VRAM each); each individual experiment completed in at most a few hours. The Qwen3-30B-A3B-Base ablation (Appendix H) was run through the Tinker API rather than on local hardware; total Tinker API spend across all 30B runs (including those reported in the paper) was under \$17. No model training was performed in this work.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 1182 • The paper should disclose whether the full research project required more compute
1183 than the experiments reported in the paper (e.g., preliminary or failed experiments that
1184 didn't make it into the paper).

1185 **9. Code of ethics**

1186 Question: Does the research conducted in the paper conform, in every respect, with the
1187 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

1188 Answer: [Yes]

1189 Justification: We use only publicly released models (GPT-2, Qwen2.5-3B, Qwen3-30B-
1190 A3B-Base, OLMo-2-7B) and the publicly released Tevet & Berant McDiv benchmark; no
1191 human subjects were enrolled, no data was scraped, and no PII is processed.

1192 Guidelines:

- 1193 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
1194 Ethics.
1195 • If the authors answer [No], they should explain the special circumstances that require a
1196 deviation from the Code of Ethics.
1197 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
1198 eration due to laws or regulations in their jurisdiction).

1199 **10. Broader impacts**

1200 Question: Does the paper discuss both potential positive societal impacts and negative
1201 societal impacts of the work performed?

1202 Answer: [Yes]

1203 Justification: The Impact Statement (immediately preceding the references) discusses both
1204 the positive impact (better measurement of mode collapse and creative-writing diversity
1205 supports evaluation and oversight of generative AI) and the potential negative impacts (use
1206 in pipelines that select for more capable AI; possible automation of human creative roles).

1207 Guidelines:

- 1208 • The answer [N/A] means that there is no societal impact of the work performed.
1209 • If the authors answer [N/A] or [No], they should explain why their work has no societal
1210 impact or why the paper does not address societal impact.
1211 • Examples of negative societal impacts include potential malicious or unintended uses
1212 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
1213 (e.g., deployment of technologies that could make decisions that unfairly impact specific
1214 groups), privacy considerations, and security considerations.
1215 • The conference expects that many papers will be foundational research and not tied
1216 to particular applications, let alone deployments. However, if there is a direct path to
1217 any negative applications, the authors should point it out. For example, it is legitimate
1218 to point out that an improvement in the quality of generative models could be used to
1219 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
1220 that a generic algorithm for optimizing neural networks could enable people to train
1221 models that generate Deepfakes faster.
1222 • The authors should consider possible harms that could arise when the technology is
1223 being used as intended and functioning correctly, harms that could arise when the
1224 technology is being used as intended but gives incorrect results, and harms following
1225 from (intentional or unintentional) misuse of the technology.
1226 • If there are negative societal impacts, the authors could also discuss possible mitigation
1227 strategies (e.g., gated release of models, providing defenses in addition to attacks,
1228 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
1229 feedback over time, improving the efficiency and accessibility of ML).

1230 **11. Safeguards**

1231 Question: Does the paper describe safeguards that have been put in place for responsible
1232 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
1233 image generators, or scraped datasets)?

1234 Answer: [N/A]

1235 Justification: We release neither a high-risk pretrained model nor scraped data. The released
1236 artefact is the set of OLMo-2-7B generations on AlpacaEval and NoveltyBench-curated
1237 prompts (a transcript of standard sampled outputs from a public model), plus the metric
1238 implementation; none of these introduce misuse risks beyond what the underlying public
1239 OLMo-2 release already poses.

1240 Guidelines:

- 1241 • The answer [N/A] means that the paper poses no such risks.
- 1242 • Released models that have a high risk for misuse or dual-use should be released with
1243 necessary safeguards to allow for controlled use of the model, for example by requiring
1244 that users adhere to usage guidelines or restrictions to access the model or implementing
1245 safety filters.
- 1246 • Datasets that have been scraped from the Internet could pose safety risks. The authors
1247 should describe how they avoided releasing unsafe images.
- 1248 • We recognize that providing effective safeguards is challenging, and many papers do
1249 not require this, but we encourage authors to take this into account and make a best
1250 faith effort.

1251 12. Licenses for existing assets

1252 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
1253 the paper, properly credited and are the license and terms of use explicitly mentioned and
1254 properly respected?

1255 Answer: [Yes]

1256 Justification: All models and datasets used in the paper are cited by their canonical references.
1257 Models: GPT-2 [Radford et al., 2019], Qwen2.5-3B [Yang et al., 2024], Qwen3-30B-A3B-
1258 Base [Yang et al., 2025], OLMo-2-7B [OLMo et al., 2024], and the Llama 3 family at 8B/70B
1259 [Grattafiori et al., 2024] together with the Llama 3.2 lightweight 1B/3B variants [Meta AI,
1260 2024]; SentBERT-style sentence embeddings [Reimers and Gurevych, 2019] are used as
1261 a baseline. Datasets: Tevet & Berant’s diversity-eval / McDiv [Tevet and Berant, 2021],
1262 AlpacaFarm/AlpacaEval prompts [Dubois et al., 2023], and NoveltyBench-curated prompts
1263 [Zhang et al., 2025c]. Most assets are released under permissive licenses (Apache-2.0 for the
1264 Qwen and OLMo families, MIT for GPT-2 and the diversity-eval repo, research-use licenses
1265 for the benchmark datasets); Llama 3 and Llama 3.2 are released under the bespoke Meta
1266 Llama Community License, which is more restrictive than Apache-2.0/MIT but permits
1267 research use. All assets are used within their stated terms.

1268 Guidelines:

- 1269 • The answer [N/A] means that the paper does not use existing assets.
- 1270 • The authors should cite the original paper that produced the code package or dataset.
- 1271 • The authors should state which version of the asset is used and, if possible, include a
1272 URL.
- 1273 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1274 • For scraped data from a particular source (e.g., website), the copyright and terms of
1275 service of that source should be provided.
- 1276 • If assets are released, the license, copyright information, and terms of use in the
1277 package should be provided. For popular datasets, paperswithcode.com/datasets
1278 has curated licenses for some datasets. Their licensing guide can help determine the
1279 license of a dataset.
- 1280 • For existing datasets that are re-packaged, both the original license and the license of
1281 the derived asset (if it has changed) should be provided.
- 1282 • If this information is not available online, the authors are encouraged to reach out to
1283 the asset’s creators.

1284 13. New assets

1285 Question: Are new assets introduced in the paper well documented and is the documentation
1286 provided alongside the assets?

1287 Answer: [Yes]

1288 Justification: New assets released: (i) the metric implementation (open-source code at
1289 <https://anonymous.4open.science/r/icl-diversity-67E6/> with README, unit
1290 tests, and a CLI documented in CLAUDE.md); (ii) the OLMo-2-7B generations dataset (200
1291 AlpacaEval + 100 NoveltyBench-curated prompts $\times$ 4 stages $\times$ $K = 10$ generations)
1292 released under `results/rlhf_experiment/` with a sidecar README describing schema,
1293 sampling parameters, and license. No new model weights are released.

1294 Guidelines:

- 1295 • The answer [N/A] means that the paper does not release new assets.
- 1296 • Researchers should communicate the details of the dataset/code/model as part of their
1297 submissions via structured templates. This includes details about training, license,
1298 limitations, etc.
- 1299 • The paper should discuss whether and how consent was obtained from people whose
1300 asset is used.
- 1301 • At submission time, remember to anonymize your assets (if applicable). You can either
1302 create an anonymized URL or include an anonymized zip file.

1303 14. Crowdsourcing and research with human subjects

1304 Question: For crowdsourcing experiments and research with human subjects, does the paper
1305 include the full text of instructions given to participants and screenshots, if applicable, as
1306 well as details about compensation (if any)?

1307 Answer: [N/A]

1308 Justification: We did not run crowdsourcing or human-subjects experiments. The human
1309 ratings used in the Tevet evaluation are from the previously published McDiv benchmark,
1310 which is cited as the source.

1311 Guidelines:

- 1312 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1313 with human subjects.
- 1314 • Including this information in the supplemental material is fine, but if the main contribu-
1315 tion of the paper involves human subjects, then as much detail as possible should be
1316 included in the main paper.
- 1317 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
1318 or other labor should be paid at least the minimum wage in the country of the data
1319 collector.

1320 15. Institutional review board (IRB) approvals or equivalent for research with human 1321 subjects

1322 Question: Does the paper describe potential risks incurred by study participants, whether
1323 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
1324 approvals (or an equivalent approval/review based on the requirements of your country or
1325 institution) were obtained?

1326 Answer: [N/A]

1327 Justification: No human-subjects research was performed in this work; we use only previ-
1328 ously released human ratings from the Tevet & Berant McDiv benchmark.

1329 Guidelines:

- 1330 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1331 with human subjects.
- 1332 • Depending on the country in which research is conducted, IRB approval (or equivalent)
1333 may be required for any human subjects research. If you obtained IRB approval, you
1334 should clearly state this in the paper.
- 1335 • We recognize that the procedures for this may vary significantly between institutions
1336 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
1337 guidelines for their institution.
- 1338 • For initial submissions, do not include any information that would break anonymity (if
1339 applicable), such as the institution conducting the review.

1340 16. Declaration of LLM usage

1341 Question: Does the paper describe the usage of LLMs if it is an important, original, or
1342 non-standard component of the core methods in this research? Note that if the LLM is used
1343 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
1344 scientific rigor, or originality of the research, declaration is not required.

1345 Answer: [\[Yes\]](#)

1346 Justification: An LLM is the central object of the proposed method: a base language model θ
1347 (e.g. Qwen2.5-3B) is used at inference time as the scoring model that produces the a_k curve
1348 and the per-byte cross-entropy C (Section 3). The choice of θ is documented per-experiment,
1349 and ablations across model scales are reported in Appendix H.

1350 Guidelines:

- 1351 • The answer [\[N/A\]](#) means that the core method development in this research does not
- 1352 involve LLMs as any important, original, or non-standard components.
- 1353 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
- 1354 be described.