

---

# “I’ve Seen How This Goes”: Characterizing the Diversity of LLM Generations and Human Writing via Progressive Conditional Surprise

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Measuring the diversity of creative outputs is central to evaluating post-training  
2 mode collapse, comparing decoding strategies, and quantifying creative behavior  
3 in both AI and human writing. We propose a new approach to measuring diver-  
4 sity using in-context learning, of which the “Decan” metric,  $D_{Ca_n} = C \times a_n$ ,  
5 is the working instance we evaluate: a per-byte score read off the per-token log-  
6 probabilities of a base model  $\theta$  in a *single forward pass* per permutation, with no  
7 embedding model, no reference corpus, and no human labels. This approach is  
8 grounded in information theory, makes use of language model in-context learning  
9 to detect a wide range of similarities between any number of inputs, and obviates  
10 the need to train a special-purpose model. The same pipeline scores AI samples  
11 and human-written response sets, with diversity treated as a property of (responses,  
12 prompt, scoring model). On Tevet and Berant’s human-grounded McDiv bench-  
13 mark,  $D_{Ca_n}$  reaches OCA 0.846 on the McDiv prompt\_gen set where it performs  
14 best, behind the strongest neural baseline reported in Tevet and Berant (SentBERT,  
15 0.897). On the OLMo-2-7B post-training pipeline,  $D_{Ca_n}$  drops monotonically  
16 across the base  $\rightarrow$  SFT  $\rightarrow$  DPO  $\rightarrow$  RLVR stages, detecting the type of diversity  
17 loss that creative-writing applications care about.

## 1 Motivation

19 The possibility for diverse outputs is necessary, but not sufficient, for creativity. Machine learning  
20 researchers that study creativity routinely need to compare the diversity of outputs across generation  
21 processes: post-training stages of the same base model (mode collapse), decoding strategies (what  
22 does temperature trade off against?), prompting interventions (which prompts have a wide response  
23 distribution?), and human-AI hybrid pipelines (would AI use cause human-AI writing to be less  
24 creative?). Existing diversity metrics rely on embedding distances [Du and Black, 2019] or surface-  
25 level  $n$ -gram statistics, both of which lack a human-like ability to recognize arbitrary patterns when  
26 those patterns differ from their training data or surface similarities, respectively (Appendix A). A  
27 complementary line of work applies information theory directly: Maximum Mutual Information  
28 decoding [Li et al., 2016] and Adversarial Information Maximization [Zhang et al., 2018] optimise  
29 pairwise mutual information between input and response, but as a training or decoding objective  
30 rather than an evaluation-time diagnostic. We propose a new approach to measuring diversity using  
31 in-context learning, of which the “Decan” metric  $D_{Ca_n} = C \times a_n$  is the working instance we evaluate:  
32 it scores AI samples and human-written response sets through the same pipeline of per-byte log-  
33 probabilities of a base model  $\theta$ , in a *single forward pass* over the concatenated responses (Figure 1).  
34 There is no embedding model, no reference corpus, no human labels, no auxiliary classifier. The  
35 approach uses only the per-token probabilities  $\theta$  already produces.




Figure 1: The diversity-metric pipeline. Given prompt  $p$  and  $n$  responses from policy  $\pi$ , we format them with response labels (“Response A:”, “Response B:”, ...) and tokenize. The **conditional track** (left) feeds the concatenated context to the base model  $\theta$  in a single forward pass per permutation  $\sigma$ , extracts per-response total surprise, divides by each response’s UTF-8 byte count, and averages over permutations to get  $\bar{a}_k$ . Its last point is  $a_n$ . The **unconditional track** (right) scores each  $r_i$  independently against the prompt alone, again per-byte; the geometric mean of these surprises defines coherence  $C = 1/\text{PPL}_\theta(\pi, p)$ , the reciprocal of the geometric-mean per-byte perplexity. The product  $D_{Ca_n} = C \times a_n$  measures plausibility-weighted residual diversity (bits/byte). Formal definitions: Sections 2, 3.1, 3.2, 3.3.

36 The intuition is that if responses from a policy  $\pi$  are diverse, seeing one should not help  $\theta$  predict the  
 37 next; if they are repetitive or constrained to a few modes, conditioning on earlier responses should  
 38 sharply reduce  $\theta$ ’s surprise at later ones. We exploit the in-context learning capability described by  
 39 Brown et al. [2020], applied here as a measurement lens rather than as a few-shot-prompting setup.  
 40 The progressive conditional surprise curve  $a_k$  (the per-byte cross-entropy of the  $k$ -th response given  
 41 the previous  $k-1$ ) captures this signal directly, and its last point is  $a_n$  (Section 3.1). A separate  
 42 coherence term  $C = 1/\text{PPL}_\theta(\pi, p)$ , the reciprocal of the geometric-mean per-byte perplexity that  $\theta$   
 43 assigns to each response individually, prevents pure noise from registering as “diverse” (Section 3.2).  
 44 The product  $D_{Ca_n} = C \times a_n$  is the working scalar we adopt; it is plausibility-weighted residual  
 45 diversity in bits per byte. We validate it against Tevet and Berant’s human-grounded diversity  
 46 benchmark (Section 4) and on the OLMo-2-7B post-training pipeline (Section 5), and release a  
 47 working open-source implementation alongside.<sup>1</sup> The metric measures diversity as  $\theta$  perceives it:  
 48 outputs differing only in ways  $\theta$  cannot distinguish in context appear less diverse. This relativity  
 49 gives the metric the potential to tighten as base models improve (see Appendix H for a preliminary  
 50 scaling experiment).

## 51 2 Setup and Notation

52 We use the following notation:

- 53 •  $p$  be a prompt,
- 54 •  $\pi$  be the policy under evaluation,
- 55 •  $r_1, r_2, \dots, r_n \sim \pi(\cdot \mid p)$  be  $n$  i.i.d. responses sampled from  $\pi$ ,
- 56 •  $\theta$  be a trusted base model from which we can obtain per-token log-probabilities,
- 57 •  $|r|_{\text{tok}}$  denote the number of tokens in response  $r$ ,
- 58 •  $\|r\|$  denote the number of bytes in the UTF-8 encoding of response  $r$ .

<sup>1</sup>Code and data: <https://anonymous.4open.science/r/icl-diversity-67E6/>

59 We define the **cross-entropy** (total surprise) of a response  $r$  under  $\theta$ :

$$-\log_2 \theta(r \mid p) = \sum_{t=1}^{|r|_{\text{tok}}} -\log_2 \theta(r^t \mid r^{<t}, p) \quad (1)$$

60 where  $r^t$  is the  $t$ -th token and  $r^{<t}$  the preceding tokens. Units: bits.<sup>2</sup> This is the total surprise of the  
 61 response under  $\theta$ , a function of the string and of  $\theta$ 's distribution but not of  $\theta$ 's tokenizer (since the  
 62 chain rule yields the same total regardless of how the sequence is factored).

63 Similarly, the **conditional cross-entropy** given previous responses  $r_{<k} = (r_1, \dots, r_{k-1})$  is  
 64  $-\log_2 \theta(r_k \mid r_{<k}, p)$  (bits).

65 For the coherence term and diversity scores (Sections 3.2 and 3.3), we also use the **per-byte cross-**  
 66 **entropy rate**:

$$h_\theta(r \mid p) = \frac{-\log_2 \theta(r \mid p)}{\|r\|} \quad (2)$$

67 Units: bits/byte. We adopt this per-byte rate because it works better than total bits in our experiments;  
 68 we have not investigated why. Normalising by byte count rather than token count keeps the rate  
 69 independent of  $\theta$ 's tokenizer when comparing base models with different vocabularies.

70 In practice, computing  $\theta(r_k \mid r_{<k}, p)$  requires feeding  $\theta$  a formatted context containing the prompt  
 71 and all previous responses (see Section B).

## 72 3 Method

73 This section defines the three quantities used in the main body: the progressive conditional surprise  
 74 curve  $a_k$ , the coherence term  $C$ , and the scalar score  $D_{Ca_n} = C \times a_n$ . The full information-theoretic  
 75 motivation, the alternative excess-entropy summary  $E$ , and the cross-mode learning analysis appear  
 76 in Appendices E–H.

### 77 3.1 Progressive Conditional Surprise

78 Given prompt  $p$  and  $n$  responses  $r_1, \dots, r_n$  from policy  $\pi$ , define

$$a_k = -\log_2 \theta(r_k \mid r_{<k}, p) \quad \text{for } k = 1, \dots, n, \quad (3)$$

79 the total surprise of the  $k$ -th response under base model  $\theta$  given the previous  $k-1$  responses. The  
 80 sequence  $(a_1, \dots, a_n)$  is the **progressive conditional surprise curve**. We normalize each  $a_k$  by the  
 81 byte count  $\|r_k\|$  to get a per-byte (bits/byte) curve that is independent of  $\theta$ 's tokenizer. All quantities  
 82 below are per-byte unless otherwise stated.

83 The intuition is that as  $\theta$  sees more responses from  $\pi$ , the conditional surprise drops if the responses  
 84 share patterns ( $\theta$ 's in-context learning picks up on modes, stylistic regularities, topic patterns) and  
 85 stays flat if they do not. For a policy with rich coherent diversity the curve declines and levels off at a  
 86 positive floor; for a policy producing one repeated mode the curve drops sharply to near zero; for  
 87 pure noise the curve stays roughly constant at a high value.

88 In practice we average  $a_k$  over uniformly random permutations of the response ordering (so each  
 89 position averages over all responses and the curve reflects only how  $\theta$ 's predictions improve with  
 90 more context, not which response happened to land in slot  $k$ ). We take the metric value to be the  
 91 last observed point  $a_n$ . Equivalently,  $a_n = H_\theta(r_n \mid p) - I_\theta(r_n; r_1, \dots, r_{n-1} \mid p)$ : the response's  
 92 individual surprise under  $\theta$  minus how much the prior responses make it predictable. High  $a_n$   
 93 therefore requires  $r_n$  to be both individually surprising and not made predictable by the others, two  
 94 properties one would want a diversity score to reward. No fitting or extrapolation step is involved.  
 95 The curve  $a_k$  is itself informative: its shape distinguishes a one-mode collapse (sharp drop to a low  
 96 floor) from richer diversity (gradual decline to a higher floor) in ways the endpoint alone obscures.

<sup>2</sup>Throughout, "bits" refers to self-information in  $\log_2$  units, identical to the "shannon" (Sh) of IEC 80000-13. We use "bits" to match standard usage in the information-theory literature.

Table 1: Intended edge-case behavior of  $D_{Ca_n} = C \times a_n$  on the five synthetic scenarios (Appendix C.2), with empirical  $D_{Ca_n}$  values under GPT-2 (124M) and Qwen2.5-3B base. The product correctly suppresses pure noise and incoherent multi-mode (via  $C$ ) and one-mode sets (via  $a_n$ ); multi-mode coherent is the intended winner. Empirical numbers show smaller  $a_n$  for multi-mode coherent than this story predicts: with only 3 distinct modes in a set of 10 responses,  $\theta$ ’s in-context learning identifies them and  $a_n$  is no longer high. In such ambiguous cases reporting the full  $a_k$  curve is more principled than the scalar  $a_n$ : a different choice of  $n$  would reveal the diversity, and showing the curve sidesteps having to pick that  $n$  in advance. Appendix C.3 shows  $a_n$  rising with mode count. Mixed empirically scores highest (via a high  $a_n$ ) rather than the predicted “mid” position; reweighting strategies are discussed in Section C.6. On Qwen2.5-3B, multi-mode incoherent also outranks multi-mode coherent; Section 3.3 discusses.

| Scenario              | Predicted |       |                  | Empirical $D_{Ca_n}$ |            |
|-----------------------|-----------|-------|------------------|----------------------|------------|
|                       | $C$       | $a_n$ | $D_{Ca_n}$       | GPT-2                | Qwen2.5-3B |
| Pure noise            | low       | high  | low (via $C$ )   | 0.04                 | 0.05       |
| Multi-mode incoherent | low       | high  | low (via $C$ )   | 0.19                 | 0.29       |
| Multi-mode coherent   | high      | high  | <b>high</b>      | 0.26                 | 0.19       |
| One-mode              | high      | low   | low (via $a_n$ ) | 0.10                 | 0.09       |
| Mixed                 | mid       | high  | mid              | 0.52                 | 0.41       |

### 3.2 The Coherence Term

The curve alone is fooled by pure noise:  $\theta$  can never predict random tokens from one another, so  $a_n$  stays at the high unconditional surprise level and noise would dominate the metric. The distinguishing signal lies in how plausible  $\theta$  finds each individual response, independent of the others.

Let  $h_\theta(r | p) = -\frac{1}{\|r\|} \log_2 \theta(r | p)$  be the per-byte cross-entropy of response  $r$  under  $\theta$  conditioned on the prompt alone. Define **coherence** as the geometric mean of the per-byte probabilities:

$$C = 2^{-\frac{1}{n} \sum_{i=1}^n h_\theta(r_i | p)} = \frac{1}{\text{PPL}_\theta(\pi, p)}, \quad (4)$$

the reciprocal of the geometric-mean per-byte perplexity<sup>3</sup> that  $\theta$  assigns to the responses individually.<sup>4</sup> Perplexity is a standard metric of incoherence, and the geometric form is intended to suppress sets containing incoherent responses: a single sample with high per-byte cross-entropy drives  $C$  toward zero, limiting the rescue effect of any single fluent response on an otherwise incoherent set.

### 3.3 The Diversity Score

The scalar score is the product:

$$D_{Ca_n} = C \times a_n \quad (\text{bits/byte}). \quad (5)$$

It can be read as “how many bits of surprise per byte remain after  $\theta$  has learned what it can from the other responses, weighted by output plausibility.” The intended edge-case behaviors are summarized in Table 1, alongside empirical  $D_{Ca_n}$  on synthetic instances of each scenario under two base models (full per-metric breakdown in Appendix C.2, Table 5).

The score depends on  $p$ ,  $\pi$ ’s responses, and  $\theta$ . Diversity is therefore a property of (responses, prompt, scoring model), not of the policy in isolation: the same response set can score differently under a stronger or weaker  $\theta$ . This is by design:  $\theta$ ’s in-context learning capability is the lens through which diversity is measured, giving the score the potential to improve as base models improve (see Appendix H for a preliminary scaling experiment). For example, Table 1 shows Qwen2.5-3B scoring multi-mode incoherent (0.29) above multi-mode coherent (0.19), reversing our predicted ranking; Figure 3 shows Qwen’s  $\bar{a}_k$  curve dropping over  $k$  on the multi-mode incoherent scenario,

<sup>3</sup>In our Tevet evaluation with Qwen2.5-3B, the central 90% (5th to 95th percentile) of per-response bits-per-byte values across 48705 responses span [0.65, 2.20], giving  $C \in [0.22, 0.64]$ ; the mean per-set  $C$  across 9741 response sets is 0.42.

<sup>4</sup>For scale, the bits/byte score is the average per-byte cross-entropy  $\bar{\ell}$ , related to coherence by  $C = 2^{-\bar{\ell}}$ . GPT-3 davinci on The Pile [Gao et al., 2020] scored  $\sim 0.72$  bits/byte and GPT-2 XL  $\sim 1.05$  bits/byte, giving  $C \approx 0.61$  and  $C \approx 0.49$  respectively; stronger base models reach lower bits/byte and correspondingly higher  $C$ .

Table 2: ConTest results: Spearman  $\rho$  and OCA between a set’s diversity class and each metric score. Baselines from Tevet’s pre-computed metrics;  $C \times a_n$  and  $a_n$  are ours (Qwen2.5-3B, 50 permutations, per-byte). Best result per task in bold.

| Metric                             | prompt_gen    |              | resp_gen      |              | story_gen     |              |
|------------------------------------|---------------|--------------|---------------|--------------|---------------|--------------|
|                                    | $\rho$        | OCA          | $\rho$        | OCA          | $\rho$        | OCA          |
| <i>ConTest (200, with_hds)</i>     |               |              |               |              |               |              |
| $C \times a_n$ (ours)              | +0.584        | 0.785        | +0.391        | 0.668        | +0.686        | 0.828        |
| $a_n$ (ours)                       | +0.444        | 0.715        | +0.274        | 0.641        | +0.387        | 0.684        |
| $C$ (ours)                         | +0.214        | 0.625        | −0.001        | 0.555        | +0.247        | 0.632        |
| SentBERT                           | <b>+0.682</b> | 0.815        | <b>+0.591</b> | <b>0.791</b> | <b>+0.770</b> | <b>0.896</b> |
| BERTsts                            | +0.646        | <b>0.820</b> | +0.463        | 0.714        | +0.601        | 0.780        |
| distinct- $n$                      | +0.333        | 0.675        | +0.346        | 0.677        | +0.573        | 0.772        |
| <i>McDiv_nuggets (~1K, no_hds)</i> |               |              |               |              |               |              |
| $C \times a_n$ (ours)              | +0.636        | 0.785        | +0.345        | 0.649        | +0.317        | 0.634        |
| $a_n$ (ours)                       | +0.487        | 0.705        | +0.225        | 0.619        | +0.124        | 0.557        |
| $C$ (ours)                         | +0.138        | 0.567        | +0.082        | 0.545        | +0.251        | 0.643        |
| SentBERT                           | <b>+0.728</b> | <b>0.850</b> | <b>+0.532</b> | <b>0.758</b> | <b>+0.633</b> | <b>0.803</b> |
| BERTsts                            | +0.683        | 0.830        | +0.393        | 0.676        | +0.344        | 0.638        |
| distinct- $n$                      | −0.003        | 0.514        | −0.002        | 0.507        | −0.002        | 0.510        |
| <i>McDiv (full, no_hds, ~2K)</i>   |               |              |               |              |               |              |
| $C \times a_n$ (ours)              | +0.729        | 0.846        | +0.500        | 0.724        | +0.523        | 0.717        |
| $a_n$ (ours)                       | +0.617        | 0.781        | +0.432        | 0.698        | +0.402        | 0.668        |
| $C$ (ours)                         | +0.138        | 0.565        | −0.007        | 0.512        | +0.171        | 0.594        |
| SentBERT                           | <b>+0.796</b> | <b>0.897</b> | <b>+0.678</b> | <b>0.830</b> | <b>+0.753</b> | <b>0.867</b> |
| BERTsts                            | +0.780        | 0.893        | +0.614        | 0.781        | +0.571        | 0.740        |
| distinct- $n$                      | +0.476        | 0.746        | +0.517        | 0.738        | +0.535        | 0.744        |

indicating that its in-context learning recognises the shared template structure as a pattern despite the within-response scrambling. GPT-2’s weaker in-context learning fits the predicted ordering on this row better (0.19 for incoherent vs. 0.26 for coherent), but we still recommend stronger base models when possible – sophisticated patterns are exactly what an ICL-based metric should be able to detect. When comparing policies across a prompt suite,  $D_{Ca_n}$  can be averaged over prompts or differenced; see Appendix G.

## 4 Tevet–Berant Diversity-Eval: Human-Grounded Validation

A diversity metric should be tested against human judgements on a standardized eval; Tevet and Berant [2021] provide both. Their McDiv and ConTest datasets pair human-grounded diversity labels with a fixed comparison protocol (Spearman  $\rho$  and OCA: *optimal classification accuracy*, the best accuracy achievable by a one-dimensional threshold separating low- and high-diversity sets), the standard methodology for ranking diversity metrics against human judgements. The response sets are written by Mechanical Turk workers, and for both McDiv\_nuggets and ConTest the high/low label is fixed by the construction protocol itself: a worker writes five different continuations (the high-diversity set), then self-selects one of those continuations and paraphrases it five times preserving content but varying form (the low-diversity set; see Appendix F for the full protocol and a construction-confound caveat). We run  $D_{Ca_n}$  through this evaluation pipeline.

**Setup.** We use Qwen2.5-3B (base) with 50 permutations in completion format (Appendix B.1) on the released splits: 6002 McDiv sets (no\_hds), 3069 McDiv\_nuggets sets (no\_hds), and 670 ConTest sets (with\_hds), each with 5 worker-written responses per set and a binary high/low diversity label. We follow Tevet and Berant in reporting Spearman  $\rho$  and OCA, and additionally report ROC AUC (the natural metric for binary classification).

**Headline result.** On the binary tasks  $D_{Ca_n}$  clearly tracks diversity, with OCA between 3.7% and 21.0% behind SentBERT (Table 2). On McDiv prompt\_gen, where  $D_{Ca_n}$  performs best, it reaches  $\rho = +0.729$  and OCA = 0.846, against SentBERT’s  $\rho = +0.796$  and OCA = 0.897; the

Table 3: DecTest results: Spearman  $\rho$  between sampling temperature and each metric score (1000 samples, no\_hds). Reported for parity with Tevet and Berant; the labels here are sampling temperatures rather than human judgements, so DecTest is not part of the framing for  $D_{C a_n}$ . That  $a_n$  tracks temperature is mechanically unsurprising:  $a_n$  is a conditional-entropy quantity, and sampling temperature directly scales the entropy of the policy, so a positive  $\rho$  here is a sanity check rather than evidence about creative diversity; Peeperkorn et al. [2024] likewise find only a weak relationship between temperature and creativity in LLM outputs.

| Metric                | prompt_gen    | resp_gen      | story_gen     |
|-----------------------|---------------|---------------|---------------|
| $C \times a_n$ (ours) | +0.847        | +0.771        | +0.763        |
| $a_n$ (ours)          | <b>+0.932</b> | <b>+0.924</b> | <b>+0.779</b> |
| $C$ (ours)            | -0.727        | -0.813        | -0.547        |
| distinct- $n$         | +0.917        | +0.894        | +0.758        |
| BERTScore             | +0.878        | +0.874        | +0.694        |
| SentBERT              | +0.747        | +0.801        | +0.645        |
| cos-sim               | +0.873        | +0.895        | +0.712        |

corresponding ROC AUC is 0.921. The headline takeaway is that an ICL-based diversity metric reaches close to a sentence-embedding baseline on a human-grounded eval, with no embedding model, no reference corpus, and no human labels in the metric itself.

**DecTest, included for completeness.** Tevet and Berant also release DecTest (Table 3), a diagnostic where the “label” is the sampling temperature used to generate each response set rather than a human judgement. We include it for parity with the released benchmark only: we do not believe that temperature labels map onto creative/meaningful output diversity, consistent with Peeperkorn et al. [2024]’s finding that the relationship between sampling temperature and creativity is more nuanced and weaker than the “creativity parameter” framing suggests. The raw  $a_n$  achieves  $\rho = +0.932$  on prompt\_gen and  $\rho = +0.924$  on resp\_gen, matching or exceeding all baselines.

**Caveats.** The McDiv\_nuggets construction protocol introduces a confound: low-diversity sets are paraphrases of specific, dramatic endings that are intrinsically more surprising to the base model than typical high-diversity continuations. Appendix F documents the mechanism, the per-byte  $a_1$  gap, and an evidence-by-length-bin breakdown showing the effect is content-driven not length-driven. The  $C$  weighting in  $D_{C a_n}$  helps on this benchmark partly for that confound-related reason; on settings without such a confound we expect  $a_n$  to carry most of the signal.

## 5 OLMo-2-7B Post-Training: AI-Side Validation

The Tevet evaluation in Section 4 validates  $D_{C a_n}$  against human-grounded labels. This section evaluates  $D_{C a_n}$  on real policies: we sample from four post-training stages of the same base model and ask whether the metric detects the diversity loss widely attributed to RLHF [Kirk et al., 2023, Zhang et al., 2025a, Padmakumar and He, 2023]. The same scoring pipeline applies to both human-written response sets (Section 4) and the policy-sampled sets here, so  $\theta$  is held fixed across the two evaluations.

**Setup.** We sample from the four released stages of the OLMo-2-1124-7B pipeline [OLMo et al., 2024]: the pretrained base model, its SFT checkpoint, the DPO checkpoint trained on preference data, and the final RLVR-tuned Instruct checkpoint. On two prompt sets, 200 AlpacaFarm evaluation prompts [Dubois et al., 2023] (seed=42 subsample of the 805-prompt set later adopted by AlpacaEval) and 100 curated NoveltyBench prompts [Zhang et al., 2025c], we draw  $K=10$  responses per prompt per stage (temperature 1.0, top- $p$  1.0, max\_new\_tokens=100). Base runs raw; SFT / DPO / Instruct receive the prompt through their own chat template. Because per-byte cross-entropy in a causal LM decreases with within-response context, we truncate every (stage, prompt) tuple’s responses to a common per-prompt UTF-8 byte length before scoring,<sup>5</sup> and drop prompts where that common length

<sup>5</sup>Truncation is applied to the response *string* via `text.encode("utf-8")[:N].decode("utf-8", errors="ignore")` and the truncated string is then re-tokenised for the forward pass. We deliberately do not truncate at token boundaries:  $a_n$  is per-byte, so we need a fixed byte denominator across stages, but tokens-per-byte is

falls below 50 bytes (150/200 AlpacaEval and 39/100 NB-curated prompts retained; see Section 6). We score each (stage, prompt) group with  $D_{C_{a_n}}$  using Qwen2.5-3B as  $\theta$  and 25 permutations. For comparison we also compute Kirk et al.’s EAD and distinct- $n$  (averaged over  $n=1 \dots 5$ ) and a SentBERT-embedding [Reimers and Gurevych, 2019] similarity-to-diversity reduction.

**Hypotheses.** We pre-register three one-sided paired Wilcoxon signed-rank tests [Dror et al., 2018] with family-wise Bonferroni correction across the three contrasts [Dror et al., 2017] ( $\alpha = 0.05/3$ ):

- **H1a**  $D_{\text{base}} > D_{\text{SFT}}$ : SFT narrows toward the helpful-assistant style.
- **H1b**  $D_{\text{SFT}} > D_{\text{DPO}}$ : DPO is trained on preference data and inherits its typicality bias [Zhang et al., 2025a].
- **H1c**  $D_{\text{base}} > D_{\text{RLVR}}$ : cumulative effect of the full post-training pipeline.

The DPO-vs-RLVR comparison is reported as an exploratory two-sided contrast (H1’): we have no directional prediction for which post-training stage loses more diversity.

Table 4: Per-prompt  $D_{C_{a_n}}$  by OLMo-2-7B stage on the length-matched subset, with pre-registered paired tests. Reported  $p$ -values for H1a–H1c are Bonferroni-corrected ( $\times 3$ ); the H1’ row (marked \*) is an uncorrected, two-sided exploratory comparison whose direction we had no prior expectation for, and is not a load-bearing claim.

| AlpacaEval                |                    |       |                       |
|---------------------------|--------------------|-------|-----------------------|
| Stage                     | $D_{C_{a_n}}$ mean | std   | $n$                   |
| Base                      | 0.481              | 0.038 | 150                   |
| SFT                       | 0.329              | 0.084 | 150                   |
| DPO                       | 0.286              | 0.072 | 150                   |
| Instruct (RLVR)           | 0.281              | 0.071 | 150                   |
| Contrast                  | $\Delta$           | $d_z$ | $p_{\text{Bonf}}$     |
| Base>SFT                  | 0.151              | 1.615 | $3.4 \times 10^{-24}$ |
| SFT>DPO                   | 0.044              | 0.675 | $1.7 \times 10^{-13}$ |
| Base>Instruct (RLVR)      | 0.200              | 2.425 | $5.2 \times 10^{-26}$ |
| DPO $\neq$ Instruct (H1’) | 0.005              | 0.227 | 0.004*                |
| NoveltyBench curated      |                    |       |                       |
| Stage                     | $D_{C_{a_n}}$ mean | std   | $n$                   |
| Base                      | 0.481              | 0.030 | 39                    |
| SFT                       | 0.369              | 0.085 | 39                    |
| DPO                       | 0.312              | 0.082 | 39                    |
| Instruct (RLVR)           | 0.303              | 0.086 | 39                    |
| Contrast                  | $\Delta$           | $d_z$ | $p_{\text{Bonf}}$     |
| Base>SFT                  | 0.112              | 1.212 | $4.1 \times 10^{-8}$  |
| SFT>DPO                   | 0.057              | 0.957 | $2.8 \times 10^{-6}$  |
| Base>Instruct (RLVR)      | 0.179              | 1.889 | $2.3 \times 10^{-10}$ |
| DPO $\neq$ Instruct (H1’) | 0.010              | 0.375 | 0.028*                |

**Results.** Table 4 reports per-stage means and the pre-registered Wilcoxon contrasts on both prompt sets. On AlpacaEval and NoveltyBench-curated alike,  $D_{C_{a_n}}$  decreases monotonically across base  $\rightarrow$  SFT  $\rightarrow$  DPO  $\rightarrow$  Instruct (RLVR), and all three pre-registered contrasts (H1a–H1c) are significant after Bonferroni correction. The underlying  $\bar{a}_k$  curves and per-prompt  $D_{C_{a_n}}$  distributions are visualised in Appendix I (Figure 17): each later stage’s curve lies below the base curve at every  $k \geq 2$ , and the per-prompt distribution shifts toward lower  $D_{C_{a_n}}$  as the pipeline advances.

**Discussion.** The monotone drop across all three pre-registered contrasts is consistent with the post-training-reduces-diversity findings of Kirk et al. [2023], Padmakumar and He [2023], now measured by an ICL-based diversity metric that needs neither an embedding model nor a sentence similarity function. Per-prompt  $D_{C_{a_n}}$  also agrees with the lexical (EAD) and semantic (SentBERT)

tokenizer-dependent and varies across responses with the same byte budget. That variation is harmless because both the bits-numerator and the byte-denominator are computed on the same truncated string.

199 baselines on the same length-matched subset, confirming the metrics see the same diversity-loss  
200 signal (Appendix I, Figure 18).

201 **Released artefacts.** The  $200 + 100$  prompts  $\times$  4 stages  $\times$   $K=10$  generations are released as a  
202 public dataset under `results/rlhf_experiment/` in the project repository. To our knowledge this  
203 is the first public release of  $K \geq 10$  sampled responses across a paired SFT/DPO/RL pipeline, filling  
204 a gap that blocked direct diversity replication of Kirk et al. [2023] (whose  $K=16$  generations were  
205 not publicly released).

## 206 6 Limitations

207 We collect the main limitations of the framework and of our  $D = C \times a_n$  choice in one place so  
208 readers can calibrate how strongly to update on the results above. Several points below recap caveats  
209 already made in their natural context and cross-reference those for detail.

210 **Measurement is relative to  $\theta$ 's perception.** The metric evaluates diversity *as perceived by the*  
211 *base model  $\theta$* : if  $\pi$ 's outputs differ only in ways  $\theta$  cannot distinguish from context, the metric  
212 underestimates diversity (Section 1, "Important caveat"). Concretely, the off-diagonal sign of the  
213 pairwise surprise-reduction matrix changes with  $\theta$ 's scale (Section C.4), so  $a_k$  curves for the same  
214 response set are not directly comparable across different choices of  $\theta$ .

215 **Metric selection saw the full Tevet benchmark.** We chose  $C \times a_n$  (rather than alternative  
216 summaries such as  $C \times E$ ; Appendix E) after observing its performance on the full Tevet diversity-  
217 eval benchmark, without holding out a validation split. The Tevet numbers in Section 4 may therefore  
218 be optimistically biased: the scalar form was selected with visibility into those results. The OLMo-  
219 2-7B post-training experiment (Section 5) was scored after the metric form was fixed and provides  
220 independent validation; the synthetic mode-count experiments (Section C.3) informally influenced the  
221 choice (had they failed, we would likely have revised the metric) and so are not strictly independent.

222 **External benchmark performance trails the strongest baseline.** On Tevet and Berant's diversity-  
223 eval benchmark,  $C \times a_n$  trails SentBERT by 3.7%–21.0% OCA across the nine binary tasks (Sec-  
224 tion 4); scaling the base model to Qwen3-30B-A3B-Base [Yang et al., 2025] did not close the gap  
225 (Appendix H). We position ICL-based diversity metrics (of which  $D_{Ca_n}$  is one instance we found to  
226 work well) as a principled complement to embedding-based metrics, not a replacement.

227 **Length-matching drops short-response prompts.** Reporting in bits/byte (rather than total bits)  
228 makes per-byte cross-entropy not invariant to response length, so the OLMo-2-7B experiment scores  
229 a length-matched subset and drops 111/300 prompts whose common per-prompt byte length falls  
230 below 50 bytes; full mechanics and the un-truncated robustness check are in Section C.6.

231  **$D = C \times a_n$  is a pragmatic choice, not a theoretically derived one.** The product form was  
232 selected because it captures the properties we wanted a diversity score to have (residual surprise  
233 remains after the base model has seen previous responses, low-coherence outputs are penalised, and  
234 the result stays on a bits-per-byte scale), not because it is derived from an axiomatic information-  
235 theoretic setup. Other scalars respecting the same constraints (weighted integrals of  $a_k - a_\infty$ ,  
236 slope-based scores, or coherence applied at a different point on the curve) could equally be defended  
237 (see "The framework admits other metrics" in Section C.6). We adopt  $D_{Ca_n}$  because it is interpretable  
238 and works empirically across our validation regimes, not because it is uniquely correct. We invite  
239 other researchers to build on this work to develop other formulas that characterize the  $a_k$  curve in  
240 meaningful ways.

## 241 7 Conclusion

242 We have presented a new approach to measuring diversity using in-context learning, of which  
243  $D_{Ca_n} = C \times a_n$  is the working instance we evaluate: a base model's in-context learning detects  
244 similarities between arbitrary numbers of responses in a single forward pass. The approach requires  
245 no embedding model, no reference corpus, no human labels, and no special-purpose training. The



246 same pipeline scores AI samples and human-written sets. On Tevet and Berant’s human-grounded  
 247 benchmark it is near the level of the strongest sentence-embedding baseline. On the OLMo-2-7B  
 248 post-training pipeline it drops across the base → SFT → DPO/RLVR stages, tracking the diversity  
 249 loss attributed to RLHF-style training.

## 250 Impact Statement

251 This paper presents work which may be used in evaluation pipelines that select for more capable  
 252 AI systems. We view the prospect of increasingly capable future AI as one of unclear positive or  
 253 negative net impact. This work may also contribute to the automation of human creative roles. We  
 254 publish it in the hope that better understanding of machine-generated diversity will contribute to  
 255 better evaluation and oversight, and that these benefits outweigh the risks.

## 256 References

- 257 Lukas Berglund, Meg Tong, Max Kaufmann, Mikita Balesni, Asa Cooper Stickland, Tomasz Korbak,  
 258 and Owain Evans. The reversal curse: LLMs trained on "a is b" fail to learn "b is a", 2023. URL  
 259 <http://arxiv.org/abs/2309.12288v4>.
- 260 Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal,  
 261 Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel  
 262 Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler,  
 263 Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott  
 264 Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya  
 265 Sutskever, and Dario Amodei. Language models are few-shot learners, 2020. URL <http://arxiv.org/abs/2005.14165v4>.
- 267 James P. Crutchfield and David P. Feldman. Regularities unseen, randomness observed: Levels of  
 268 entropy convergence. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 13(1):25–54,  
 269 March 2003. ISSN 1089-7682. doi: 10.1063/1.1530990. URL [http://dx.doi.org/10.1063/](http://dx.doi.org/10.1063/1.1530990)  
 270 [1.1530990](http://dx.doi.org/10.1063/1.1530990).
- 271 Rotem Dror, Gili Baumer, Marina Bogomolov, and Roi Reichart. Replicability analysis for natural  
 272 language processing: Testing significance with multiple datasets. *Transactions of the Association*  
 273 *for Computational Linguistics*, 5:471–486, 2017. doi: 10.1162/tacl\_a\_00074. URL [https:](https://aclanthology.org/Q17-1033/)  
 274 [//aclanthology.org/Q17-1033/](https://aclanthology.org/Q17-1033/).
- 275 Rotem Dror, Gili Baumer, Segev Shlomov, and Roi Reichart. The hitchhiker’s guide to testing  
 276 statistical significance in natural language processing. In Iryna Gurevych and Yusuke Miyao,  
 277 editors, *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics*  
 278 *(Volume 1: Long Papers)*, pages 1383–1392, Melbourne, Australia, July 2018. Association for  
 279 Computational Linguistics. doi: 10.18653/v1/P18-1128. URL [https://aclanthology.org/](https://aclanthology.org/P18-1128/)  
 280 [P18-1128/](https://aclanthology.org/P18-1128/).
- 281 Wenchao Du and Alan W Black. Boosting dialog response generation. In Anna Korhonen, David  
 282 Traum, and Lluís Màrquez, editors, *Proceedings of the 57th Annual Meeting of the Association for*  
 283 *Computational Linguistics*, pages 38–43, Florence, Italy, July 2019. Association for Computational  
 284 Linguistics. doi: 10.18653/v1/P19-1005. URL <https://aclanthology.org/P19-1005/>.
- 285 Yann Dubois, Xuechen Li, Rohan Taori, Tianyi Zhang, Ishaan Gulrajani, Jimmy Ba, Carlos Guestrin,  
 286 Percy Liang, and Tatsunori B. Hashimoto. AlpacaFarm: A simulation framework for methods that  
 287 learn from human feedback, 2023. URL <http://arxiv.org/abs/2305.14387v4>.
- 288 Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang,  
 289 Horace He, Anish Thite, Noa Nabeshima, Shawn Presser, and Connor Leahy. The pile: An  
 290 800gb dataset of diverse text for language modeling, 2020. URL [http://arxiv.org/abs/2101.](http://arxiv.org/abs/2101.00027v1)  
 291 [00027v1](http://arxiv.org/abs/2101.00027v1).
- 292 Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad  
 293 Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan,  
 294 Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev,

295 Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru,  
 296 Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak,  
 297 Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu,  
 298 Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle  
 299 Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego  
 300 Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova,  
 301 Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang, Gabriel  
 302 Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire Mialon,  
 303 Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan  
 304 Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet,  
 305 Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelder van der Linde,  
 306 Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie  
 307 Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua  
 308 Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani, Kate Plawiak,  
 309 Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz Malik, Kuenley  
 310 Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yeary, Laurens van der Maaten, Lawrence  
 311 Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo, Lukas Blecher, Lukas  
 312 Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat Singh, Manohar Paluri,  
 313 Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya Pavlova, Melanie  
 314 Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman Goyal, Narjes  
 315 Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang, Olivier Duchenne,  
 316 Onur Çelebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic, Peter Weng, Prajjwal  
 317 Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu, Qing He, Qingxiao Dong,  
 318 Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira Cabral, Robert Stojnic,  
 319 Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain Sauvestre, Ronnie  
 320 Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar Hosseini, Sahana  
 321 Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov, Shaoliang Nie,  
 322 Sharan Narang, Sharath Rapparthi, Sheng Shen, Shengye Wan, Shruti Bhosale, Shun Zhang, Simon  
 323 Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane Collot, Suchin Gururangan,  
 324 Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha, Thomas Georgiou, Thomas  
 325 Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal Karn, Vedanuj Goswami,  
 326 Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gonguet, Virginie Do, Vish Voleti,  
 327 Vítor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenyin Fu, Whitney Meers, Xavier  
 328 Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide Xia, Xinfeng Xie, Xuchao  
 329 Jia, Xuwei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei, Yi Wen, Yiwen Song,  
 330 Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpiere Coudert, Zheng Yan, Zhengxing Chen, Zoe  
 331 Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey, Adam Shajnfeld, Adithya  
 332 Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma, Alex Boesenberg, Alexei  
 333 Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo, Anam Yunus, Andrei Lupu,  
 334 Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew Poulton, Andrew Ryan, Ankit  
 335 Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita Saraf, Arkabandhu Chowdhury,  
 336 Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh Yazdan, Beau James, Ben Maurer,  
 337 Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola, Bhargavi Paranjape, Bing Liu,  
 338 Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence, Brani Stojkovic, Brian Gamido,  
 339 Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu, Changan Wang, Changkyu  
 340 Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris Tindal, Christoph Feichtenhofer,  
 341 Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel Li, David Adkins, David Xu,  
 342 Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich, Didem Foss, Dingkan Wang, Duc  
 343 Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine Montgomery, Eleonora Presani, Emily  
 344 Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban Arcaute, Evan Dunbar, Evan Smothers,  
 345 Fei Sun, Felix Kreuk, Feng Tian, Filippas Kokkinos, Firat Ozgenel, Francesco Caggioni, Frank  
 346 Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella Schwarz, Gada Badeer, Georgia Sweet,  
 347 Gil Halpern, Grant Herman, Grigory Sizov, Guangyi, Zhang, Guna Lakshminarayanan, Hakan Inan,  
 348 Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha, Haroun Habeeb, Harrison Rudolph,  
 349 Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan Zhan, Ibrahim Damlaj, Igor Molybog,  
 350 Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai Gat, Jake Weissman, James Geboski, James  
 351 Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya, Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny  
 352 Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica Zhong, Jian Jin, Jingyi Yang, Joe Cummings,  
 353 Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan Torres, Josh Ginsburg, Junjie Wang, Kai

354 Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal, Katayoun Zand, Kathy Matosich, Kaushik  
355 Veeraraghavan, Kelly Michelena, Keqian Li, Kiran Jagadeesh, Kun Huang, Kunal Chawla, Kyle  
356 Huang, Lailin Chen, Lakshya Garg, Lavender A, Leandro Silva, Lee Bell, Lei Zhang, Liangpeng  
357 Guo, Licheng Yu, Liron Moshkovich, Luca Wehrstedt, Madian Khabsa, Manav Avalani, Manish  
358 Bhatt, Martynas Mankus, Matan Hasson, Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim  
359 Naumov, Maya Lathi, Meghan Keneally, Miao Liu, Michael L. Seltzer, Michal Valko, Michelle  
360 Restrepo, Mihir Patel, Mik Vyatskov, Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang,  
361 Miquel Jubert Hermoso, Mo Metanat, Mohammad Rastegari, Munish Bansal, Nandhini Santhanam,  
362 Natascha Parks, Natasha White, Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier,  
363 Nikhil Mehta, Nikolay Pavlovich Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia  
364 Hart, Omkar Salpekar, Ozlem Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro  
365 Rittner, Philip Bontrager, Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani,  
366 Pritish Yuvraj, Qian Liang, Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy,  
367 Raghu Nayani, Rahul Mitra, Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin  
368 Battey, Rocky Wang, Russ Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu,  
369 Samyak Datta, Sara Chugh, Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh  
370 Mahajan, Saurabh Verma, Seiji Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Shaun Lindsay,  
371 Sheng Feng, Shenghao Lin, Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang,  
372 Shuqiang Zhang, Sinong Wang, Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie  
373 Max, Stephen Chen, Steve Kehoe, Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta,  
374 Summer Deng, Sungmin Cho, Sunny Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman,  
375 Tal Remez, Tamar Glaser, Tamara Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun  
376 Zhang, Tim Matthews, Timothy Chou, Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria  
377 Montanez, Vijai Mohan, Vinay Satish Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru,  
378 Vlad Tiberiu Mihailescu, Vladimir Ivanov, Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz,  
379 Will Constable, Xiao Cheng Tang, Xiaojian Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv  
380 Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi, Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi,  
381 Youngjin Nam, Yu, Wang, Yu Zhao, Yuchen Hao, Yundi Qian, Yunlu Li, Yuzi He, Zach Rait,  
382 Zachary DeVito, Zef Rosenbriek, Zhaoduo Wen, Zhenyu Yang, Zhiwei Zhao, and Zhiyu Ma. The  
383 llama 3 herd of models, 2024. URL <http://arxiv.org/abs/2407.21783v3>.

384 Yanzhu Guo, Guokan Shang, and Chloé Clavel. Benchmarking linguistic diversity of large language  
385 models, 2024. URL <http://arxiv.org/abs/2412.10271v2>.

386 Robert Kirk, Ishita Mediratta, Christoforos Nalmpantis, Jelena Luketina, Eric Hambro, Edward  
387 Grefenstette, and Roberta Raileanu. Understanding the effects of rlhf on llm generalisation and  
388 diversity, 2023. URL <http://arxiv.org/abs/2310.06452v3>.

389 Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. A diversity-promoting  
390 objective function for neural conversation models. In Kevin Knight, Ani Nenkova, and Owen  
391 Rambow, editors, *Proceedings of the 2016 Conference of the North American Chapter of the*  
392 *Association for Computational Linguistics: Human Language Technologies*, pages 110–119, San  
393 Diego, California, June 2016. Association for Computational Linguistics. doi: 10.18653/v1/N16-  
394 1014. URL <https://aclanthology.org/N16-1014/>.

395 Meta AI. Llama 3.2: Revolutionizing edge AI and vision with open, customizable models. Meta  
396 AI Blog, September 2024. URL [https://ai.meta.com/blog/llama-3-2-connect-2024-](https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/)  
397 [vision-edge-mobile-devices/](https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/).

398 Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, Akshita Bhagia,  
399 Yuling Gu, Shengyi Huang, Matt Jordan, Nathan Lambert, Dustin Schwenk, Oyvind Tafjord, Taira  
400 Anderson, David Atkinson, Faeze Brahman, Christopher Clark, Pradeep Dasigi, Nouha Dziri,  
401 Allyson Ettinger, Michal Guerquin, David Heineman, Hamish Ivison, Pang Wei Koh, Jiacheng Liu,  
402 Saumya Malik, William Merrill, Lester James V. Miranda, Jacob Morrison, Tyler Murray, Crystal  
403 Nam, Jake Poznanski, Valentina Pyatkin, Aman Rangapur, Michael Schmitz, Sam Skjonsberg,  
404 David Wadden, Christopher Wilhelm, Michael Wilson, Luke Zettlemoyer, Ali Farhadi, Noah A.  
405 Smith, and Hannaneh Hajishirzi. 2 olmo 2 furious, 2024. URL [http://arxiv.org/abs/2501.](http://arxiv.org/abs/2501.00656v3)  
406 [00656v3](http://arxiv.org/abs/2501.00656v3).

407 Vishakh Padmakumar and He He. Does writing with language models reduce content diversity?,  
408 2023. URL <http://arxiv.org/abs/2309.05196v3>.

409 Max Peeperkorn, Tom Kouwenhoven, Dan Brown, and Anna Jordanous. Is temperature the creativity  
410 parameter of large language models?, 2024. URL <http://arxiv.org/abs/2405.00492v1>.

411 Tianyi Qiu, Ahmed Hani Ismail, Zhonghao He, and Shi Feng. Self-improvement as coherence  
412 optimization: A theoretical account, 2026. URL <http://arxiv.org/abs/2601.13566v1>.

413 Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever.  
414 Language models are unsupervised multitask learners. Technical report, OpenAI,  
415 2019. URL [https://cdn.openai.com/better-language-models/language\\_models\\_](https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf)  
416 [are\\_unsupervised\\_multitask\\_learners.pdf](https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf).

417 Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks,  
418 2019. URL <http://arxiv.org/abs/1908.10084v1>.

419 Guy Tevet and Jonathan Berant. Evaluating the evaluation of diversity in natural language generation.  
420 In Paola Merlo, Jorg Tiedemann, and Reut Tsarfay, editors, *Proceedings of the 16th Conference of*  
421 *the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 326–  
422 346, Online, April 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.eacl-  
423 main.25. URL <https://aclanthology.org/2021.eacl-main.25/>.

424 Zihan Wang, Kangrui Wang, Qineng Wang, Pingyue Zhang, Linjie Li, Zhengyuan Yang, Xing Jin,  
425 Kefan Yu, Minh Nhat Nguyen, Licheng Liu, Eli Gottlieb, Yiping Lu, Kyunghyun Cho, Jiajun Wu,  
426 Li Fei-Fei, Lijuan Wang, Yejin Choi, and Manling Li. Ragen: Understanding self-evolution in  
427 llm agents via multi-turn reinforcement learning, 2025. URL [http://arxiv.org/abs/2504.](http://arxiv.org/abs/2504.20073v2)  
428 [20073v2](http://arxiv.org/abs/2504.20073v2).

429 Jiaxin Wen, Zachary Ankner, Arushi Somani, Peter Hase, Samuel Marks, Jacob Goldman-Wetzler,  
430 Linda Petrini, Henry Sleight, Collin Burns, He He, Shi Feng, Ethan Perez, and Jan Leike. Unsu-  
431 pervised elicitation of language models, 2025. URL <http://arxiv.org/abs/2506.10139v2>.

432 An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li,  
433 Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin  
434 Yang, Jiayi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang,  
435 Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi  
436 Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan,  
437 Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report, 2024. URL  
438 <http://arxiv.org/abs/2412.15115v2>.

439 An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang  
440 Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu,  
441 Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin  
442 Yang, Jiayi Yang, Jing Zhou, Jingren Zhou, Junyang Lin, Kai Dang, Keqin Bao, Kexin Yang,  
443 Le Yu, Lianhao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui  
444 Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang  
445 Ren, Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yinger  
446 Zhang, Yu Wan, Yuqiong Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan  
447 Qiu. Qwen3 technical report, 2025. URL <http://arxiv.org/abs/2505.09388v1>.

448 Jiayi Zhang, Simon Yu, Derek Chong, Anthony Sicilia, Michael R. Tomz, Christopher D. Manning,  
449 and Weiyang Shi. Verbalized sampling: How to mitigate mode collapse and unlock llm diversity,  
450 2025a. URL <http://arxiv.org/abs/2510.01171v3>.

451 Tianhui Zhang, Bei Peng, and Danushka Bollegala. Evaluating the evaluation of diversity in com-  
452 mon-sense generation, 2025b. URL <http://arxiv.org/abs/2506.00514v1>.

453 Yiming Zhang, Harshita Diddee, Susan Holm, Hanchen Liu, Xinyue Liu, Vinay Samuel, Barry Wang,  
454 and Daphne Ippolito. Noveltybench: Evaluating language models for humanlike diversity, 2025c.  
455 URL <http://arxiv.org/abs/2504.05228v4>.

456 Yizhe Zhang, Michel Galley, Jianfeng Gao, Zhe Gan, Xiujun Li, Chris Brockett, and Bill Dolan. Gen-  
457 erating informative and diverse conversational responses via adversarial information maximization,  
458 2018. URL <http://arxiv.org/abs/1809.05972v5>.

459 Yaoming Zhu, Sidi Lu, Lei Zheng, Jiaxian Guo, Weinan Zhang, Jun Wang, and Yong Yu. Texygen: A  
460 benchmarking platform for text generation models, 2018. URL [http://arxiv.org/abs/1802.](http://arxiv.org/abs/1802.01886v1)  
461 01886v1.

## A Related Work

**Sampling diversity metrics.** Work on decoding-time diversity typically operates at the surface level:  $n$ -gram overlap [Li et al., 2016] and self-BLEU [Zhu et al., 2018]. Our approach operates at the distributional level and can distinguish policies that produce lexically varied but semantically redundant outputs.

**Diversity evaluation benchmarks.** Tevet and Berant [2021] introduced a systematic framework for evaluating diversity metrics, with their McDiv benchmark providing human-labeled response sets at two diversity levels. We use McDiv for validation (Section 4), though we identify a construction confound in how its low-diversity sets are produced (Appendix F). More recent work has revisited the meta-evaluation problem. NoveltyBench [Zhang et al., 2025c] defines diversity through pairwise functional equivalence rather than binary labels, training a DeBERTa classifier to group generations into equivalence classes and computing a  $Distinct_k$  score (the number of meaningfully different outputs in  $k$  samples), conceptually close to our mode count. However, their ground truth is itself a trained classifier (79% accuracy), making it unclear whether correlation with their labels validates a metric or merely measures agreement between two model-based scores. Zhang et al. [2025b] conduct a meta-evaluation of diversity metrics for constrained commonsense generation, using GPT-4o as an annotator in place of crowd workers. Guo et al. [2024] benchmark linguistic diversity of LLMs, building on Tevet and Berant’s framework with a broader set of generation tasks.

**Perplexity.** The coherence term  $C = 1/\text{PPL}$  (Section 3.2) connects our framework to the standard language modeling evaluation metric. The primary diversity score  $D_{Ca_n} = C \times a_n$  operates in bits/byte, weighting the residual surprise floor by output plausibility.

**Coherence as a signal.** Our coherence term  $C$  uses the base model’s predictive distribution as an unsupervised quality signal. This connects to a broader line of work using model-internal coherence without external supervision: Qiu et al. [2026] show theoretically that feedback-free self-improvement methods work by optimizing coherence (compressibility of context-to-behavior mappings), and Wen et al. [2025] use internal coherence maximization to elicit capabilities from language models. Our framework leverages a related insight: the base model’s ability to compress responses in context provides a meaningful signal about their diversity structure.

**Embedding-based diversity.** Du and Black [2019] introduced the embedding-based approach, clustering sentence embeddings of generated responses with  $k$ -means and reporting cluster inertia as the diversity score; subsequent work has explored alternative geometric statistics over embedded responses. These approaches are complementary to ours: they capture semantic distance in a learned representation space, while our metric captures statistical independence under a generative model. The approaches may disagree when  $\theta$ ’s in-context learning captures structure invisible to the embedding model, or vice versa.

**Output collapse in RL training.** Wang et al. [2025] document an “Echo Trap” pattern during multi-turn agent RL: early-stage agents produce varied symbolic reasoning, then collapse to deterministic templates as training proceeds. They detect this collapse using within-prompt reward standard deviation as an early-warning signal, a proxy that relies on the reward distinguishing degenerate outputs from diverse ones. The  $a_k$  framework provides a text-level alternative: collapse should appear as a drop in  $a_n$  on the agent’s own outputs, independent of reward structure. A downside of using the  $a_k$  framework is that it is unconnected from a downstream task. Responses can both be “diverse” as measured by  $a_k$ -based metrics and fail uniformly.

**Excess entropy.** The  $a_k$  curve also admits an excess-entropy summary related to the computational mechanics literature [Crutchfield and Feldman, 2003]; see Appendix E. We found it empirically inferior to  $D_{Ca_n}$  on Tevet and Berant’s diversity-eval benchmark [Tevet and Berant, 2021].

## B Practical Considerations

### B.1 Formatting the Conditioning Context

Computing  $\theta(r_k \mid r_{<k}, p)$  requires feeding  $\theta$  a context containing the prompt and previous responses. The formatting choice affects results, and we use two formats depending on whether the responses are best read as parallel answers to the prompt or as continuations of it. Both are implemented as `format_conditioning_context` in `src/icl_diversity/core.py` and selected via a `format_mode` argument throughout the pipeline.

**Instruct format.** The default, used for the synthetic scenarios (Appendix C.2), the mode-count experiments (Appendix C.3), and the OLMo-2-7B post-training case study (Section 5):

```
[prompt p]
Response A: [r_1]
Response B: [r_2]
Response C: [r_3]
```

The prompt appears once and each response is introduced by a labelled header (“Response A:”, “Response B:”, ..., with labels rolling over to “AA”, “AB” past “Z”), encouraging the base model’s in-context learning to treat the responses as parallel answers to the same prompt.

**Completion format.** Used for the Tevet and Berant evaluation (Section 4), where each response set is a continuation of a shared narrative prompt rather than an instruction-style answer:

```
1. [prompt p] [r_1]
2. [prompt p] [r_2]
3. [prompt p] [r_3]
```

The prompt is repeated immediately before each response so  $\theta$  scores  $r_k$  in the same prompt-as-context that produced it. Per-token log-probabilities are extracted only over the response tokens; the bits accumulated over the repeated prompt prefix are excluded from  $a_k$ , keeping the bits-numerator and the per-byte denominator both attributable to  $r_k$  alone.

**Base vs. instruction-tuned models.** We use a base (non-instruction-tuned) model as  $\theta$  for two reasons. First, an instruction-tuned model has had its output distribution shaped by RLHF or similar procedures, so using one as  $\theta$  reintroduces the kind of distributional bias we seek to avoid. Second, an instruction-tuned  $\theta$  may assign systematically different probabilities to text that follows or violates its alignment training, introducing a confound between coherence-as-fluency and coherence-as-alignment that the metric should not conflate. Modern base models already exhibit strong in-context learning from pretraining alone, so this choice does not sacrifice ICL capability. Optimizing the prompt format to maximally elicit in-context learning from  $\theta$  is an interesting direction but out of scope for this paper; we use the neutral “Response A/B/C” template throughout.

### B.2 Computational Cost

Because  $\theta$  is a causal language model, the full  $a_k$  curve for one response ordering can be computed from a **single forward pass**. The prompt and all responses are concatenated into one sequence (Figure 2). The forward pass produces the log-probability of every token conditioned on all preceding tokens.

The tokens belonging to  $r_k$  are conditioned on exactly  $p, r_1, \dots, r_{k-1}$  (plus formatting), which is the conditioning required for  $a_k$ . Partitioning the output log-probabilities by response boundary and summing within each partition yields all  $n$  values of  $a_k$  simultaneously for that ordering, with FLOPs scaling as  $O((|p| + n\bar{L}_{\text{tok}})^2)$  from causal attention, where  $\bar{L}_{\text{tok}}$  is the average response length

[prompt p] Response A: [r\_1] Response B: [r\_2] ... Response N: [r\_n]

Figure 2: Single-pass input layout for computing the full  $a_k$  curve. The prompt is followed by all  $n$  responses with “Response A/B/C/...” labels; one forward pass over this sequence yields every  $a_k$  value simultaneously, since causal attention conditions the tokens of  $r_k$  on exactly  $p, r_1, \dots, r_{k-1}$  (plus formatting).

in tokens. Because we report  $\bar{a}_k$  averaged over  $|\Sigma|$  random permutations of the response ordering (Section B.3), the total  $a_k$ -curve cost is  $|\Sigma|$  such passes; the permutations are independent and we batch them on a single GPU.

The coherence term  $C$  additionally requires the *unconditional* per-byte cross-entropies  $h_\theta(r_i | p)$  for each response. These cannot be extracted from the concatenated pass, since in that pass  $r_i$  for  $i > 1$  is conditioned on all prior responses. Computing them requires  $n$  independent forward passes, each over the short context  $(p, r_i)$ . These passes are embarrassingly parallel and batchable, and they do not depend on the ordering, so  $C$  is computed once per response set regardless of  $|\Sigma|$ .

In summary, scoring one (prompt, response set) tuple takes:

- $a_k$  curve:  $|\Sigma|$  forward passes (long context, one per permutation).
- $C$ :  $n$  forward passes (short contexts, batchable; not multiplied by  $|\Sigma|$ ).

With  $n = 10$  responses and  $|\Sigma| = 50$  permutations (our default for larger-scale experiments), this is  $50 + 10 = 60$  forward passes per (prompt, response set) tuple; the long-context  $|\Sigma|$  passes dominate wall-clock time.

### B.3 Dependence on Sample Ordering

The  $a_k$  values depend on the ordering of  $\{r_i\}$ . Individual responses differ in how surprising they are to  $\theta$  given the responses that precede them, so  $a_k$  depends on which response sits at position  $k$  and which others precede it, making a single ordering’s curve jagged. Averaging over random permutations removes this dependence: each position averages over many choices of response and preceding context, so the curve reflects only how  $\theta$ ’s predictions improve with more context. We therefore average over  $|\Sigma|$  random permutations; Section C.5 compares low- and high-permutation runs on the validation scenarios and finds that the low setting misranks adjacent scenarios. We have not swept intermediate values and cannot pinpoint a cheap-and-reliable minimum, so we default to  $|\Sigma| = 100$  for scenario-level experiments and  $|\Sigma| = 50$  elsewhere.

Per-response normalization by byte count must be performed *before* averaging across permutations, since the byte count depends on which response lands at each position. Concretely, the per-byte curve is a mean of per-permutation rates,

$$\bar{a}_k = \frac{1}{|\Sigma|} \sum_{\sigma \in \Sigma} \frac{a_k^\sigma}{\|r_{\sigma(k)}\|}, \quad (6)$$

not a ratio of permutation-averaged bits to permutation-averaged byte counts (which would degenerate, after enough permutations, into the total-bits curve rescaled by the mean response length).

### B.4 Choice of $n$

The diversity score depends on  $n$  as a measurement parameter, not through any estimate of an asymptotic floor. Given enough in-context examples, even genuinely diverse policies eventually become predictable to  $\theta$ , so  $a_k$  keeps decreasing rather than converging to a meaningful irreducible value: there is no  $a_\infty$  for  $a_n$  to approximate. Too small an  $n$  means  $\theta$  has not yet exploited the inter-response structure, and  $D_{Ca_n}$  overstates diversity. Too large an  $n$  means  $\theta$  has accumulated enough in-context examples to reduce its surprise even within genuine diversity; for any policy with finitely many distinct modes,  $a_n$  continues to decrease toward zero as  $\theta$  learns to predict within-mode variation from prior examples. For comparisons to be meaningful,  $n$  must be held fixed across all policies under evaluation. Our experiments use  $n = 5$  on Tevet (Section 4),  $n = 10$  on OLMo-2-7B and scenario validation, and  $n = 20$  on the mode-count experiment (Appendix C.3). This is one reason to report the full  $a_k$  curve rather than  $D_{Ca_n}$  alone.  $D$  is a lossy summary, but the curve shows directly how response surprise evolves at every context size.



Table 5: Scenario validation metrics (mean across 5 prompts, 100 permutations,  $n = 10$  responses, per-byte). We report several candidate scalars: unconditional surprise  $a_1$ , the curve’s last point  $a_n$ , coherence  $C$ , our recommended score  $D_{C a_n} = C \times a_n$  (**bold**), the excess entropy  $E$  (Appendix E),  $C \times E$ , and the coherence spread  $\sigma_\ell$  (the standard deviation of the per-byte cross-entropies  $\{h_\theta(r_i | p)\}_{i=1}^n$ ).  $D_{C a_n}$  correctly ranks multi-mode above one-mode and suppresses incoherent scenarios via  $C$ . Both  $E$  and  $C \times E$  go negative for the mixed scenario on GPT-2. Mixed empirically scores the highest  $D_{C a_n}$  on both models, exceeding multi-mode coherent rather than landing at the “mid” position Table 1 predicts; this is not the intended ranking. The coherence spread  $\sigma_\ell$  flags this case in its intended diagnostic role (the largest  $\sigma_\ell$  across scenarios on both models), reflecting the within-set heterogeneity between coherent and incoherent responses; reweighting variants such as  $C^\alpha \times a_n$  are discussed in Section C.6.

| Scenario       | GPT-2 (124M) |       |       |             |       |              |               |      | Qwen2.5-3B |       |             |             |       |              |               |  |
|----------------|--------------|-------|-------|-------------|-------|--------------|---------------|------|------------|-------|-------------|-------------|-------|--------------|---------------|--|
|                | $a_1$        | $a_n$ | $C$   | $D_{C a_n}$ | $E$   | $C \times E$ | $\sigma_\ell$ |      | $a_1$      | $a_n$ | $C$         | $D_{C a_n}$ | $E$   | $C \times E$ | $\sigma_\ell$ |  |
| Pure noise     | 7.52         | 6.94  | 0.005 | <b>0.04</b> | 0.68  | 0.004        | 0.20          | 7.18 | 6.68       | 0.007 | <b>0.05</b> | 0.75        | 0.005 | 0.11         |               |  |
| Multi-incoher. | 3.26         | 1.85  | 0.10  | <b>0.19</b> | 5.94  | 0.61         | 0.48          | 2.36 | 1.40       | 0.21  | <b>0.29</b> | 3.12        | 0.64  | 0.69         |               |  |
| Multi-mode     | 1.06         | 0.53  | 0.48  | <b>0.26</b> | 1.61  | 0.76         | 0.08          | 0.78 | 0.31       | 0.59  | <b>0.19</b> | 1.17        | 0.68  | 0.11         |               |  |
| One-mode       | 1.06         | 0.22  | 0.48  | <b>0.10</b> | 1.31  | 0.63         | 0.06          | 0.80 | 0.16       | 0.58  | <b>0.09</b> | 0.92        | 0.53  | 0.08         |               |  |
| Mixed          | 1.90         | 2.02  | 0.26  | <b>0.52</b> | −1.29 | −0.31        | 1.17          | 1.13 | 0.91       | 0.46  | <b>0.41</b> | 0.52        | 0.23  | 0.73         |               |  |

## C Experiments

### C.1 Experimental Setup

We evaluate the metric primarily using Qwen2.5-3B (3B parameters, 32K-token context window) [Yang et al., 2024], with GPT-2 (124M parameters, 1024-token context window) [Radford et al., 2019] as a smaller-model comparison point for scenario validation. Qwen2.5-3B is used for mode-count scaling, cross-mode learning, and external validation (Sections C.3, C.4, 4). The cross-model scaling study (Section C.6) additionally uses the full Llama 3 family for comparison.

All  $a_k$  curves are computed using single-pass forward computation (Section B.2). Unless otherwise noted, we use 100 permutations for scenario validation and 50 permutations for larger-scale experiments. Responses are formatted as described in Section B.2. All code and data, including the Olmo-2 RLHF generations under `results/rlhf_experiment/`, are publicly available.<sup>6</sup>

### C.2 Scenario Validation

We construct five synthetic scenarios with known diversity structure:

1. **Pure noise:** Random ASCII characters (letters, digits, punctuation, spaces) with no learnable structure.
2. **Multi-incoherent:** Responses from 5 distinct modes, each internally incoherent (scrambled words within a template).
3. **Multi-mode:** Responses from 5 distinct modes, each internally coherent (e.g., recipe, poem, code).
4. **One-mode:** All responses from a single coherent mode (paraphrases of the same content).
5. **Mixed:** A mixture of coherent and incoherent responses.

Each scenario contains 10 responses per prompt, 5 prompts per scenario, with 100 permutations and seed 42.

**Results.** Table 5 summarizes the metrics. Figure 3 shows the  $a_k$  curves for all scenarios.

The coherence term  $C$  correctly separates coherent from incoherent scenarios:  $C \approx 0.5$ – $0.6$  for multi-mode and one-mode,  $C \approx 0.1$ – $0.3$  for multi-incoherent, and  $C \approx 0.01$  for pure noise. The coherence spread  $\sigma_\ell$  discriminates mixed from pure scenarios:  $\sigma_\ell > 1.0$  for mixed (GPT-2), reflecting the within-set heterogeneity between coherent and incoherent responses.

<sup>6</sup><https://anonymous.4open.science/r/icl-diversity-67E6/>




Figure 3: Progressive conditional surprise curves ( $a_k$ , per-byte) for all five scenarios, comparing GPT-2 (top row) and Qwen2.5-3B (bottom row). In each panel, faint colored curves are individual permutations (100 per prompt), medium colored curves are per-prompt averages, and the bold black curve is the mean across prompts. Pure noise curves are flat (no learnable structure); multi-mode curves show progressive decline; one-mode curves decline steeply then flatten.

626  $C$  consistently improves with model strength (Qwen2.5-3B assigns higher coherence to coherent text  
627 than GPT-2).  $D_{Ca_n}$  correctly ranks multi-mode above one-mode on both models, and suppresses  
628 pure noise and multi-incoherent via the  $C$  factor.

### 629 C.3 Mode Count Scaling

630 To test whether the metric reflects the number of distinct modes, we construct response sets with  
631  $m \in \{1, \dots, 10\}$  modes drawn from a pool of 50 format-based generators (haiku, code, recipe, legal  
632 disclaimer, etc.), all responding to the same prompt (“Write a short piece about rain”). For Qwen2.5-  
633 3B:  $n = 20$  responses, 1000 random draws of mode assignments, averaged over permutations.




Figure 4: Mode count scaling on Qwen2.5-3B ( $n = 20$ , 1000 draws). The  $a_k$  curves fan out with increasing  $m$ : higher floors, slower convergence. All curves are exponential (no sigmoidal plateau), even at  $m = 10$ , due to cross-mode learning (Section C.4). Shaded bands around each curve are  $\pm 1$  standard error of the mean across the 1000 random draws of mode assignments (i.e., the across-draw standard deviation divided by  $\sqrt{1000}$ ); they are narrow because averaging over 1000 draws sharply reduces the across-draw spread.

**Key findings.** Figure 4 shows the  $a_k$  curves.  $a_n$  increases monotonically with  $m$  (9.3 bits at  $m = 1$  to 77.8 bits at  $m = 10$ ; see Table 6), reflecting that with more modes there are fewer same-mode repetitions within  $n$  responses, so each response remains more surprising even after  $\theta$  has seen the others. This is exactly the behavior  $D_{C_{a_n}} = C \times a_n$  needs:  $a_n$  tracks mode count directly. All curves are purely exponential (no sigmoidal plateau), even at  $m = 10$ ; Section C.4 investigates why. For a detailed comparison of the excess-entropy-based scores on this data, including the sigmoid fit parameters and the non-monotonicity of  $\hat{E}_n$ , see Appendix E.

#### C.4 Cross-Mode Learning and Curve Shape

The authors predicted that at high  $m$ , the  $a_k$  curve should be sigmoidal: an initial plateau (early responses come from different modes and are uninformative about each other), followed by decline once mode repetitions accumulate. GPT-2 shows hints of this pattern at some  $m$  values (e.g.  $m = 10$ ). Qwen2.5-3B does not: it shows immediate exponential decay at all  $m$ . This section investigates the discrepancy through systematic pairwise analysis.

**Pairwise cross-mode surprise matrix.** For 15 modes, we compute a  $15 \times 15$  matrix  $M_{ij}$ : the surprise reduction (in bits) when a response from mode  $j$  is in the context and mode  $i$  is the target. Diagonal entries measure same-mode learning; off-diagonal entries measure cross-mode information transfer. Each entry averages over 5 context–target samples (using different samples for context and target to avoid self-prediction inflation).

**Results.** (Cell standard deviations are the mean within-cell sampling std across the 15 diagonal or 210 off-diagonal matrix cells, each estimated from 5 context-target sample pairs.)

- **Qwen2.5-3B:** diagonal mean =  $60.5 \pm 14.4$  bits, off-diagonal mean =  $+1.9 \pm 2.5$  bits, 64% of off-diagonal entries positive.
- **GPT-2:** diagonal mean =  $83.0 \pm 20.0$  bits, off-diagonal mean =  $-3.7$  bits, only 30% of off-diagonal entries positive.

The sign of the off-diagonal determines the curve shape (Figure 5):

- **Positive off-diagonal (Qwen):** Every response in context lowers expected surprise on every other response. The  $a_k$  curve drops from position 1, producing exponential decay.
- **Negative off-diagonal (GPT-2):** Cross-mode responses actively raise surprise on subsequent responses. At  $m = 10$ , the expected first-step net drop under the additive independent-mode model is only +4.9 bits (the 8.3-bit diagonal gain from the  $1/m$  same-mode chance is largely offset by the  $-3.4$ -bit cross-mode damage from the  $(m-1)/m$  different-mode chance), producing a flat plateau until same-mode repetitions accumulate, the predicted sigmoid.

GPT-2 has a larger diagonal and a negative off-diagonal: it gains more from same-mode context but is actively penalised by cross-mode context. The two effects are linked, since probability mass that conditioning concentrates on same-mode continuations is mass taken away from other modes; a sharper conditional on the seen mode is necessarily worse on the unseen ones. Qwen does the opposite: a smaller diagonal but a positive off-diagonal, distributing its update across modes.

**Pairwise asymmetry.** The mean absolute asymmetry  $|M_{ij} - M_{ji}|$  is 5.6 bits ( $3.0 \times$  the off-diagonal mean), falsifying the hypothesis that pairwise surprise reduction is symmetric (Figure 6). Information-theoretically  $I(X; Y) = I(Y; X)$ , so this asymmetry cannot reflect asymmetric information structure between modes; it reflects  $\theta$ ’s imperfection as an in-context reasoner. The same content is more useful as context than as target depending on its surface form, mirroring the *reversal curse* observed in autoregressive LLMs trained on “A is B” but failing on “B is A” [Berglund et al., 2023].

**Connections to information-theoretic properties.** Two information-theoretic properties are relevant to interpreting the pairwise matrix:

1. **Non-negative conditioning** ( $H(X) \geq H(X | Y)$ ): Under the true distribution, conditioning can never increase entropy on average. For cross-entropies this is not guaranteed (the KL term can increase), but a well-calibrated  $\theta$  should show mostly non-negative off-diagonal entries.

684 2. **Symmetry of mutual information** ( $I(X; Y) = I(Y; X)$ ): Under the true joint, row means  
685 (how informative mode  $i$  is as context) should correlate with column means (how much  
686 mode  $i$  benefits from context).

687 Figure 7 compares the two models. Qwen shows a tighter row–column relationship (slope = 0.84,  
688  $R^2 = 0.34$ ) with only 1 of 15 modes having a negative column mean. GPT-2 shows widespread  
689 violations of non-negative conditioning (6 modes with negative column means) and a barely-existent  
690 row–column relationship (slope = 0.35,  $R^2 = 0.08$ ). Across the four Llama models [Grattafiori et al.,  
691 2024, Meta AI, 2024] we additionally tested (Section C.6, Figure 9), the fraction of positive off-  
692 diagonal entries rises monotonically with size (25% at 1B to 63% at 70B), and the off-diagonal mean  
693 follows the same trend (−4.9 bits at 1B to +2.1 bits at 70B): the *magnitude* of cross-mode information  
694 transfer increases with size. The row–column relationship does not:  $R^2$  peaks at Llama-3B (0.67)  
695 and degrades at the larger Llamas (0.09 at 70B), so the *symmetry* of the pairwise matrix does not  
696 correspondingly improve.

697 **Token-level attribution.** Figure 8 shows per-token surprise reduction (in bits) for one pair from  
698 each of three strata: same-mode (letter → letter), cross-mode where conditioning helps  
699 (math\_stats → json\_data), and cross-mode where conditioning hurts (scientific\_fact →  
700 haiku). Pairs were selected by stratified-median total  $|\Delta|$ .

701 **Discussion.** The paper’s sigmoid prediction assumes modes are approximately independent. For  
702 capable models, this fails: cross-mode learning creates positive information transfer between distinct  
703 modes. The Llama 1B/3B/8B/70B series confirms this scales with model size (Section C.6): the  
704 total cross-mode information  $\theta$  extracts grows monotonically with size, and with it non-negative  
705 conditioning is better satisfied. The matrix nevertheless remains asymmetric (row–column  $R^2$  peaks  
706 at Llama-3B and degrades at the larger Llamas), which under  $I(X; Y) = I(Y; X)$  signals that even  
707 Llama-70B is an imperfect in-context reasoner.

## 708 C.5 Practical Findings

709 **Permutation sensitivity.** The per-byte  $a_k$  curve depends on the response ordering. Comparing  
710 the  $D_{Ca_n}$  ranking of the 5 validation scenarios from Section C.2 between 3- and 100-permutation  
711 runs: on GPT-2, 2 of 5 scenarios shift rank between the two settings; on Qwen2.5-3B, 2 of 5. Where  
712 scenarios do shift rank, it is between adjacent scenarios with similar  $D_{Ca_n}$  values rather than dramatic  
713 reorderings. We use 100 permutations throughout for scenario-level experiments; we have not swept  
714 intermediate values, so we cannot pinpoint a minimum that is both cheap and reliable. The coherence  
715 term  $C$  is permutation-invariant by construction and unaffected.

716 **Boundary handling.** When Qwen merges a response’s final . with the following \n\n into a single  
717 .\n\n token, our boundary detector (which uses a character-span overlap rule) attributes that token to  
718 the response, not the separator. The response’s cross-entropy therefore includes a small contribution  
719 from predicting the upcoming separator alongside its own characters; the byte-count denominator  
720 remains the response’s literal byte length. We have not measured the magnitude of this bias on our  
721 reported numbers. Boundary correctness is verified by tests/test\_response\_boundaries.py.

## 722 C.6 Discussion

723  **$\sigma_\ell$  is a diagnostic, not a diversity signal.**  $\sigma_\ell$  increased monotonically with  $m$  in our synthetic  
724 experiments because our modes have very different formats (code vs. haiku vs. recipe → different  
725 per-byte cross-entropies). But  $\sigma_\ell$  measures coherence *heterogeneity*, not diversity: modes with  
726 similar fluency would have low  $\sigma_\ell$  despite high diversity, and a single mode with high quality variance  
727 could have high  $\sigma_\ell$  with zero diversity. Its role is diagnostic (detecting mixed coherence), not as a  
728 standalone diversity score.

729 **Cross-mode information scales with model quality.** We test how cross-mode information transfer  
730 varies with  $\theta$ ’s scale by running the pairwise matrix experiment on 4 dense Llama models (1B,  
731 3B, 8B, 70B) using the same 15 modes and 5 samples per mode. Figure 9 shows the results. The  
732 off-diagonal mean increases monotonically with model size: −4.9 (1B), −1.4 (3B), −1.2 (8B), +2.1  
733 (70B) bits, transitioning from negative (cross-mode context hurts) to positive (cross-mode context

helps). The fraction of positive off-diagonal entries follows the same monotone trend: 25%, 41%, 44%, 63%. This hypothesis was pre-registered before running any Llama models, and the probability of accidental monotonicity across 4 models under the null is  $1/4! \approx 4.2\%$ .

**The framework admits other metrics.** The  $a_k$  curve is the primary object of our framework, and  $D_{Ca_n} = C \times a_n$  is one natural summary of it among many. Reporting the curve directly preserves information that any single scalar (including ours) collapses. Alternative scalars could be derived from the same curve to capture different desiderata: weighted sums of  $a_k - a_\infty$  (emphasizing early or late positions), the slope or curvature at a specific  $k$  (measuring how quickly  $\theta$  learns), curve-shape descriptors, or coherence-weighted variants that combine  $C$  with quantities other than  $a_n$ . Different applications may warrant different choices; for instance, a policy-comparison setting might prefer a metric sensitive to *changes* in the curve shape, while a ranking setting needs a single well-behaved scalar. We have not explored this space systematically; we focus on  $D_{Ca_n}$  because it empirically works well across regimes and has a clean information-theoretic interpretation, but we expect the framework to support a family of related metrics with complementary properties. One untested family is coherence-reweighting: a practitioner with a clear preference for sharper Mixed penalisation could replace  $C \times a_n$  with  $C^\alpha \times a_n$  for  $\alpha > 1$  to drive low-coherence sets toward zero faster. We have not evaluated such variants on the Tevet or OLMo experiments and so make no empirical claim about them.

**Bits/byte vs. total bits, and the length-matching consequence.** A separate choice within the framework is whether to report total bits (length-dependent but tokenizer-agnostic) or per-byte bits/byte (also tokenizer-agnostic, with length scaled out). We default to bits/byte for the headline diversity score, but per-byte quantities are not invariant to response length: in a causal LM, longer responses give  $\theta$  more within-response context to predict later tokens, lowering per-byte surprise. For the OLMo-2-7B experiment (Section 5) raw response length is not monotone across stages, so we score the length-matched subset as the headline analysis: every response in a (stage, prompt) tuple is truncated to a common per-prompt UTF-8 byte length (the minimum across all 40 responses) before scoring. 111/300 prompts had a per-prompt common length below 50 bytes – driven by the base model on AlpacaEval (no instruction-following prior) and by SFT on NB-curated short-answer prompts – and were dropped because length-matching them would compare essentially-empty fragments. The headline numbers therefore condition on the 150/200 AlpacaEval and 39/100 NB-curated prompts where length-matching is well-defined; the dropped prompts are themselves a stage-specific behavioural signal we are not measuring here. We also verified that the monotone  $D$  drop holds on the un-truncated 200 + 100 prompt sets, with all three pre-registered H1 contrasts remaining Bonferroni-significant at  $p < 10^{-13}$  on both prompt sets. Full numbers and reproduction commands are in `investigations/length_matched_rlh/VERDICT.md`.

## D Future Work

Several directions remain unexplored.

**Prompt format optimization.** We use a fixed prompt template throughout and do not attempt to optimize it for ICL elicitation. Prompt engineering could improve the metric’s discriminative power. Dimension-specific framings (e.g., “Here are several stories in different genres”) are one such variant: priming  $\theta$  to expect variation along a dimension removes the small portion of the surprise attributable to the mere existence of that variation, but it does not isolate per-dimension diversity, since the rest of the cross-response surprise reduction remains.

**Ensembling base models.** A single base model  $\theta$  may produce non-monotone  $a_k$  curves when pushed out of distribution by long contexts. Ensembling multiple base models at the token level (averaging softmax probabilities) could stabilize the estimates while preserving the autoregressive structure. Our codebase supports token-level ensembling, but we have not yet validated it experimentally.

**Disagreement cases with other metrics.** We report aggregate correlations with embedding-based and reference metrics on Tevet’s tasks (Section 4) but do not look at *which* response sets the metrics rank differently. Identifying prompts or response sets where  $C \times a_n$  disagrees with SentBERT, distinct- $n$ , or McDiv is the most direct way to characterise what each metric is actually measuring; the

disagreements are where the choice of metric matters. We expect two regimes to drive disagreement: response sets that share arbitrary rare patterns a base model can pick up on but that fall outside what a fixed sentence embedder represents (where  $C \times a_n$  should detect repetition that embeddings miss), and response sets where one response is a mixture of two others (where progressive conditioning credits the mixture for being predictable from its components, while pairwise embedding-distance averages do not).

**Total-bits variants of  $D$ .** We report  $D_{Ca_n} = C \times a_n$  in bits/byte for length control, normalising by bytes rather than tokens so that scores remain tokenizer-agnostic. The implementation also exposes  $a_n$  in total bits, which is likewise tokenizer-agnostic but not length-controlled. A hybrid scalar  $C \times a_n^{\text{bits}}$  rewards length (longer responses have more positions at which to be surprising), which some practitioners may want; characterising when each variant is the appropriate measurement target is left to future work.

**Optimizing generators against  $D$ .** We use  $C \times a_n$  as an evaluation-time diagnostic, but it could in principle serve as a training or decoding objective. Future work could fine-tune or sample from a generator  $\pi$  with the goal of maximizing  $C \times a_n$  under a fixed base model  $\theta$  to push  $\pi$  toward response sets that are both coherent and diverse.

## E Excess Entropy and the $C \times E$ Score

This appendix develops the excess entropy  $E$ , an alternative summary of the  $a_k$  curve that is theoretically interesting (connecting to the excess entropy of computational mechanics [Crutchfield and Feldman, 2003] and to the total correlation) but empirically inferior to  $D_{Ca_n}$  on the external benchmarks of Section 4. We tried both the last-point estimator  $\hat{E}_n$  and a sigmoid-extrapolated projected floor for  $a_\infty$  (Section E.1); we stopped using  $E$  when it became clear it cannot distinguish diverse from non-diverse response sets in the few-draws regime where modes do not repeat within  $n$  samples. We include it here for theoretical completeness, as a null result, and to explain the empirical comparisons that motivate our preference for  $D_{Ca_n}$ .

### E.1 The Excess Entropy

The primary metric  $D_{Ca_n} = C \times a_n$  uses the curve’s endpoint  $a_n$ : as the surprise after conditioning on  $n - 1$  peer responses,  $a_n$  reflects both the curve’s level and the drop from  $a_1$  in a single scalar. Section 4 shows this combination is what discriminates diverse from non-diverse response sets in the few-draws regime, while  $a_1$  alone (the level) and  $E$  (the integrated drop, defined below) each underperform.

We thought that  $E$ ’s learnable structure would be interesting and track diversity. Following the concept of excess entropy from computational mechanics [Crutchfield and Feldman, 2003], define the **excess entropy**:

$$E = \sum_{k=1}^{\infty} (a_k - a_\infty) \quad (7)$$

Units: bits. This is the total learnable structural information in the progressive conditioning process, above the irreducible residual noise. It converges whenever  $\theta$ ’s surprise at new responses eventually stabilizes (since  $a_k \rightarrow a_\infty$  and the excess  $e_k = a_k - a_\infty$  decays to zero). The “structure” captured by  $E$  is not limited to discrete mode identity: it includes any inter-response regularity to which  $\theta$  assigns lower conditional surprise.

In practice,  $a_\infty$  is unknown. We estimate it as  $a_n$  (the last observed value) and compute:

$$\hat{E}_n = \sum_{k=1}^n (a_k - a_n) \quad (8)$$

This underestimates  $E$  (since  $a_n \geq a_\infty$ ), but the bias decreases with  $n$ .

826 **Parametric extrapolation of  $a_\infty$ .** Rather than using  $a_n$  directly, one can fit a parametric model  
 827 to the observed curve and extrapolate  $a_\infty$ . The  $a_k$  curve is theoretically expected to be sigmoidal  
 828 (with concave-up/exponential decay as the degenerate case when the inflection point  $k_0 < 1$ ). This  
 829 motivates fitting a four-parameter sigmoid:

$$a_k = a_\infty + \frac{\alpha}{1 + e^{\beta(k-k_0)}} \quad (9)$$

830 where  $a_\infty$  is the asymptotic floor,  $\alpha$  is the total drop from  $a_1$  to  $a_\infty$ ,  $\beta > 0$  controls the steepness  
 831 of the transition, and  $k_0$  is the inflection point. The fit yields  $a_\infty$  without requiring  $n$  to be large  
 832 enough for convergence. Section C.3 provides empirical evidence that the sigmoid-extrapolated  $E_{\text{fit}}$   
 833 recovers the expected monotonic relationship between mode count and excess entropy, while the raw  
 834  $\hat{E}_n$  estimator does not.

835 **Relationship to total correlation.** The excess entropy and total correlation are complementary  
 836 decompositions of the same underlying structure. Working in expectation over i.i.d. draws from  $\pi$ , let  
 837  $\bar{a}_1 = \mathbb{E}[-\log_2 \theta(r \mid p)]$  denote the expected unconditional cross-entropy,  $\bar{a}_k = \mathbb{E}[a_k]$  the expected  
 838 conditional cross-entropy at position  $k$ , and  $e_k = \bar{a}_k - a_\infty$  the excess at step  $k$ . The per-step mutual  
 839 information is  $I_k = \bar{a}_1 - \bar{a}_k = (\bar{a}_1 - a_\infty) - e_k$ . Summing:

$$\text{TC}_n = \sum_{k=1}^n I_k = n(\bar{a}_1 - a_\infty) - E_n \quad (10)$$

840 For large  $n$  where  $E_n \rightarrow E$ , the total correlation grows linearly with slope  $(\bar{a}_1 - a_\infty)$ :  $\text{TC}_n \approx$   
 841  $n(\bar{a}_1 - a_\infty) - E$ . Figure 10 illustrates. The three quantities have clean interpretations, all in bits:  
 842  $(\bar{a}_1 - a_\infty)$  is the per-response redundancy once  $\theta$  has fully learned the available structure;  $E$  is  
 843 the total finite structural information, consumed during  $\theta$ 's transient learning phase;  $\text{TC}_n$  is the  
 844 cumulative redundancy, which grows without bound.

## 845 E.2 Per-Byte Excess Entropy Rate

846 The excess entropy  $E = \sum_k (a_k - a_\infty)$  is in bits. For a length-normalized variant, we define the  
 847 **per-byte excess entropy rate** by normalizing each response's surprise by its byte count *before*  
 848 averaging across permutations. For a single permutation  $\sigma$ , the per-byte conditional surprise at  
 849 position  $k$  is  $a_k^\sigma / \|r_{\sigma(k)}\|$  (bits/byte). Averaging over permutations gives

$$\hat{a}_k^{\text{rate}} = \mathbb{E}_\sigma \left[ \frac{a_k^\sigma}{\|r_{\sigma(k)}\|} \right] \quad (11)$$

850 and the per-byte floor is estimated analogously from the last position:  $\hat{a}_\infty^{\text{rate}} = \hat{a}_n^{\text{rate}}$ . The per-byte  
 851 excess entropy rate is

$$E_{\text{rate}} = \sum_{k=1}^n (\hat{a}_k^{\text{rate}} - \hat{a}_\infty^{\text{rate}}). \quad (12)$$

852 This treats each response equally regardless of length. Note that  $E_{\text{rate}} \neq E/\bar{B}$  in general: because  
 853 total surprise  $a_k$  and byte count  $\|r_k\|$  are positively correlated,  $E/\bar{B}$  overweights long responses,  
 854 while the per-response normalization in  $E_{\text{rate}}$  avoids this.

## 855 E.3 The $C \times E$ Score

856 Combining  $E$  with the coherence term  $C$  (Section 3.2) yields scalar scores

$$C \times E \quad (\text{bits}) \quad \text{or} \quad C \times E_{\text{rate}} \quad (\text{bits/byte}). \quad (13)$$

857 The first measures total structural information, and the second normalizes per byte so long and short  
 858 responses are treated equally.

## 859 E.4 Why Not Weight Excess Entropy Inside the Sum?

860 A natural alternative to reporting  $E$  and  $C$  separately is to weight the terms within  $E$  by coherence,  
 861 computing  $E_w = \sum_k w_k \cdot (a_k - a_\infty)$  where  $w_k = 2^{-h_\theta(r_k|p)}$ . This would suppress contributions

from incoherent modes directly at the point of measurement. We considered and rejected this approach because it breaks the information-theoretic interpretation.  $E = \sum_k e_k$  has a clean meaning as the total structural information above the noise floor, connected to Crutchfield and Feldman’s excess entropy, and inserting weights makes  $E_w$  a hybrid quantity that is not the excess entropy of any well-defined process.

## E.5 Limitations of $C \times E$

$C \times E$  inherits two limitations from  $E$ . First,  $E$  measures *recurrence*: surprise reduction as multiple responses accumulate, rather than continuous spread. A policy that draws from a broad, roughly continuous distribution of coherent outputs has  $a_k \approx a_1$  for all  $k$ , yielding  $E \approx 0$  despite maximal diversity. Second,  $E$  is near-zero in the few-draws regime: with only 5–10 responses and no repeated modes,  $\theta$  finds little learnable structure regardless of diversity, so  $C \times E$  carries almost no signal. These limitations are the reason  $D_{C a_n}$  is preferred as the primary score:  $a_\infty$  is the *level* of the floor, which differs meaningfully between diverse and non-diverse response sets even when the curve stays nearly flat.

**Empirical inferiority.** On Tevet’s diversity-eval benchmark with only 5 responses per set (Section 4),  $C \times a_n$  achieves ROC AUC of 0.867 on McDiv\_nuggets prompt\_gen while the  $D_{\text{fit}} = C \times E_{\text{fit}}$  variant cannot be computed at all (the four-parameter sigmoid fails to converge on only five points) and the discretized  $D_{\text{disc}}$  variant is anti-correlated with diversity (AUC = 0.303, well below chance). Figure 11 shows the ROC curves.

**Why  $E$  fails in this regime.** With only 5 responses and no repeated modes,  $\theta$  finds little learnable structure regardless of diversity. The  $a_k$  curve stays approximately flat, yielding  $E \approx 0$  for both diverse and non-diverse sets. The  $E$ -based scores thus carry almost no signal. The asymptotic floor  $a_\infty$ , by contrast, differs substantially between the two conditions: diverse responses remain surprising after conditioning ( $a_\infty$  high), while paraphrases become predictable ( $a_\infty$  low). This is the fundamental reason  $D_{C a_n}$  is preferred.

## E.6 Mode Count Scaling: Excess-Entropy Metrics

In the synthetic mode count experiments (Section C.3), the sigmoid-extrapolated  $E_{\text{fit}}$  is monotonic in mode count for Qwen2.5-3B but has wide confidence intervals at high  $m$ , while the raw  $\hat{E}_n$  estimator is non-monotonic. Table 6 and Figure 12 show the details.

Table 6: Mode count scaling metrics for Qwen2.5-3B ( $n = 20$ , 1000 draws).  $E_{\text{fit}}$  is the sigmoid-extrapolated excess entropy;  $a_n$  is the mean floor at  $k = 20$ ;  $\sigma_\ell$  is per-byte coherence spread;  $k_0$  is the sigmoid inflection point. Confidence intervals are 95% bootstrap. The high- $m$  rows ( $m \geq 8$ ) have  $E_{\text{fit}}$  CIs spanning more than half the mean and should be read as preliminary; the trend direction is robust but the point values are imprecise.

| $m$ | $E_{\text{fit}}$ (bits) | $a_n$ (bits) | $\sigma_\ell$ (bits/byte) | $k_0$ |
|-----|-------------------------|--------------|---------------------------|-------|
| 1   | $296 \pm 11$            | 9.3          | 0.1                       | −10.0 |
| 2   | $568 \pm 24$            | 16.9         | 0.2                       | −10.0 |
| 3   | $704 \pm 40$            | 26.7         | 0.2                       | −10.0 |
| 4   | $799 \pm 67$            | 37.4         | 0.3                       | −10.0 |
| 5   | $855 \pm 113$           | 48.0         | 0.3                       | −10.0 |
| 6   | $940 \pm 127$           | 56.1         | 0.3                       | −10.0 |
| 7   | $1069 \pm 194$          | 61.5         | 0.3                       | −10.0 |
| 8   | $1138 \pm 253$          | 66.9         | 0.3                       | −10.0 |
| 9   | $1563 \pm 705$          | 72.8         | 0.3                       | −10.0 |
| 10  | $1807 \pm 1194$         | 77.8         | 0.3                       | −10.0 |

## F Dataset Construction Confounds in McDiv

When validating against the McDiv\_nuggets benchmark [Tevet and Berant, 2021], we observed that the mean  $\bar{a}_1$  curve (per-byte) for the “low diversity” group starts *above* the “high diversity” group in



the story\_gen task (Figure 14). Since  $a_1$  measures the surprise of the first response conditioned only on the prompt (before any other responses appear in context), this gap cannot reflect diversity. It is a confound in the dataset construction.

## F.1 Mechanism

The McDiv protocol (Tevet and Berant §6.4), from which McDiv\_nuggets is sampled (Tevet and Berant Appendix C.2 specifies McDiv\_nuggets as the 3K subset of McDiv on which distinct- $n$  correlates at zero), pairs a low-diversity set with each high-diversity set as follows. Workers first write five *different* continuations (the high-diversity set). The same worker is then asked to *self-select one* of their own five responses and paraphrase it five times, preserving content while varying form (the low-diversity set). Tevet and Berant do not characterize the distribution of which responses workers self-select.

Empirically (see Table 7 and Figure 14, and the illustrative examples from the McDiv\_nuggets story\_gen dataset in Section F.3 below), we observe that the self-selected endings for the low-diversity sets tend to be more specific or dramatic (e.g., “He scored winner” or “they quit”) than the typical high-diversity continuations (e.g., “Joel fired the cook when things went too far downhill”). This means that individual low-diversity responses are intrinsically more surprising to the base model, not because of diversity, but because of the specificity of the endings workers self-selected for paraphrasing.

## F.2 Evidence

Table 7 summarizes the gap. The per-byte  $a_1$  difference is substantial (a content effect) while the total-bits difference is small because high-diversity responses are on average several bytes longer.

Table 7: Unconditional surprise ( $a_1$ ) by diversity label, McDiv\_nuggets story\_gen.

|                     | High diversity | Low diversity | Gap    |
|---------------------|----------------|---------------|--------|
| $a_1$ (bits/byte)   | 0.977          | 1.154         | +0.177 |
| $a_1$ (total bits)  | 49.1           | 50.5          | +1.4   |
| Mean response bytes | 52.7           | 46.5          | −6.3   |
| $n$                 | 103            | 97            |        |

The gap persists after binning by response length (Table 8), confirming it is primarily a content effect rather than a length artifact: the high-minus-low per-byte gap remains positive across the bulk of the length range.

Table 8: Unconditional per-byte surprise by response length bin, story\_gen.

| Byte range | High div. mean | Low div. mean | Gap    |
|------------|----------------|---------------|--------|
| [20, 40)   | 1.172          | 1.340         | +0.167 |
| [40, 60)   | 0.878          | 1.019         | +0.141 |
| [60, 80)   | 0.853          | 1.029         | +0.176 |
| [80, 120)  | 0.905          | 0.908         | +0.003 |

## F.3 Illustrative Examples

Figures 15 and 16 show per-token surprise (in bits) for the first response of a high-diversity and low-diversity sample, respectively. Samples are drawn by ranking each label group by per-byte  $a_1$  (ascending for high-diversity, descending for low-diversity) and taking the sample at the  $\lfloor N/4 \rfloor$  position, chosen to be representative of the expected confound direction rather than an extreme outlier; the specific pinned IDs (sa\_00698, sa-b\_00279) are fixed in scripts/dataset\_confound\_analysis.py for reproducibility. Red bars are measured response tokens; grey bars are masked context tokens.

**High diversity** (Figure 15). Context: “Joel owned a restaurant. He hired a cook that didn’t care about his work. The cook didn’t clean after himself. The kitchen was a mess.” Response: “Joel fired

927 *the cook when things went too far downhill.* This is a natural, predictable continuation ( $a_1 = 0.783$   
928 bits/byte).

929 **Low diversity** (Figure 16). Context: *“Harry and his basketball team was losing the game. The*  
930 *coach called for a timeout. Harry boosted morale by talking to his team. The team caught up and the*  
931 *game was tied.”* Response: *“He scored winner.”* This is a specific, somewhat ungrammatical ending  
932 ( $a_1 = 1.777$  bits/byte). All five responses in this set are paraphrases of the same idea (“Harry scored  
933 the winning point”).

## 934 F.4 Implications

935 The construction confound (low-diversity sets are paraphrases of specific dramatic endings rather  
936 than a random subsample of low-diversity content) lifts the entire  $\bar{a}_k$  curve on low-diversity sets,  
937 since their content is intrinsically surprising under  $\theta$  at every  $k$ , and pushes  $C$  down because the  
938 same content is less coherent.  $a_1$ ’s upward shift opposes the diversity label;  $C$ ’s downward shift  
939 happens to align with it (a confound artifact, not coherence genuinely tracking diversity). At  $a_n$ ,  
940 the diversity signal we are measuring overrides the confound’s upward pull: paraphrases collapse to  
941 predictable once  $\theta$  has seen the others while genuinely diverse responses do not, so  $a_n$  ends up lower  
942 on low-diversity sets and  $C \times a_n$  tracks the label despite the confound (Section 4).

## 943 G Aggregation Across Prompts

944 This appendix concerns aggregation across prompts. Aggregation *within* a single prompt, across  
945 permutations of the response ordering, is described in Section B.3 and given by Eq. (6): the per-  
946 byte  $\bar{a}_k$  is a mean of per-permutation per-byte rates, not a ratio of permutation-averaged bits to  
947 permutation-averaged byte counts.<sup>7</sup>

948 The diversity score  $D_{Ca_n}$  is defined relative to a single prompt  $p$ . To summarize a policy’s diversity  
949 across a prompt suite  $P = \{p_1, \dots, p_M\}$ , one can average:

$$\bar{D}_{Ca_n}(\pi) = \frac{1}{M} \sum_{j=1}^M D_{Ca_n}(\pi, p_j) \quad (14)$$

950 To compare two policies  $\pi_1$  and  $\pi_2$ , the natural scalar is the difference:

$$\Delta D_{Ca_n} = \bar{D}_{Ca_n}(\pi_1) - \bar{D}_{Ca_n}(\pi_2) \quad (15)$$

951  $\Delta D_{Ca_n}$  answers “how much more plausibility-weighted residual diversity does  $\pi_1$  have than  $\pi_2$ , on  
952 average across prompts?” No division is involved, so  $\Delta D_{Ca_n}$  is stable even when either policy’s  
953 score is near zero.

## 954 H Scaling the Base Model: Qwen2.5-3B vs Qwen3-30B-A3B-Base

955 To test whether a stronger base model improves the metric on Tevet’s benchmarks, we re-ran the full  
956 evaluation using Qwen3-30B-A3B-Base [Yang et al., 2025] (a 30B-parameter mixture-of-experts  
957 model with  $\sim 3B$  active parameters per token) via the Tinker API,<sup>8</sup> with the same setup as in Section 4:  
958 completion-format prompting, 50 permutations, no fine-tuning. Table 9 reports per-task  $C \times a_n$   
959 (per-byte) for both base models.

960 **Result: scaling did not help on the binary tasks.** On the 12 binary classification tasks (McDiv,  
961 McDiv\_nuggets, ConTest), Qwen2.5-3B wins on 11 of 12, with ConTest prompt\_gen the only  
962 marginal win for Qwen3-30B-A3B (+0.005 AUC). The mean degradation from scaling is small but

<sup>7</sup>An earlier revision of our implementation computed the ratio-of-means form, in which the bits and byte counts at each position are each averaged across permutations before dividing. After enough permutations, the byte counts at each position approach the mean response length, so that estimator collapses into the total-bits curve rescaled by a single constant, losing the per-permutation per-byte signal that Eq. (6) preserves. Switching the implementation to the form in Eq. (6) shifts the headline numbers in this paper by 0–17% (in mean-of-ratios’ favor), with no qualitative change to any ranking we report.

<sup>8</sup><https://thinkingmachines.ai/tinker>

Table 9: Per-task comparison of  $C \times a_n$  (per-byte) for Qwen2.5-3B vs Qwen3-30B-A3B-Base on Tevet diversity-eval.  $\Delta\text{AUC}$  is positive when the larger model wins. ROC AUC is reported for the binary tasks (McDiv, McDiv\_nuggets, ConTest); only Spearman  $\rho$  is meaningful for DecTest (continuous temperature labels).


| Section       | Task                  | Qwen2.5-3B |       | Qwen3-30B-A3B |       | $\Delta\text{AUC}$ |
|---------------|-----------------------|------------|-------|---------------|-------|--------------------|
|               |                       | $\rho$     | AUC   | $\rho$        | AUC   |                    |
| McDiv         | prompt_gen (no_hds)   | 0.729      | 0.921 | 0.718         | 0.915 | -0.006             |
|               | resp_gen (no_hds)     | 0.500      | 0.788 | 0.483         | 0.779 | -0.009             |
|               | story_gen (no_hds)    | 0.523      | 0.802 | 0.493         | 0.785 | -0.017             |
| McDiv_nuggets | prompt_gen (no_hds)   | 0.636      | 0.867 | 0.615         | 0.855 | -0.012             |
|               | prompt_gen (with_hds) | 0.683      | 0.895 | 0.655         | 0.878 | -0.017             |
|               | resp_gen (no_hds)     | 0.345      | 0.699 | 0.325         | 0.687 | -0.012             |
|               | resp_gen (with_hds)   | 0.437      | 0.753 | 0.425         | 0.746 | -0.007             |
|               | story_gen (no_hds)    | 0.317      | 0.683 | 0.293         | 0.669 | -0.014             |
|               | story_gen (with_hds)  | 0.405      | 0.734 | 0.376         | 0.717 | -0.017             |
| ConTest       | prompt_gen (with_hds) | 0.584      | 0.837 | 0.592         | 0.842 | +0.005             |
|               | resp_gen (with_hds)   | 0.391      | 0.726 | 0.333         | 0.692 | -0.034             |
|               | story_gen (with_hds)  | 0.686      | 0.896 | 0.653         | 0.877 | -0.019             |
| DecTest       | prompt_gen (no_hds)   | 0.842      | —     | 0.877         | —     | —                  |
|               | prompt_gen (with_hds) | 0.845      | —     | 0.875         | —     | —                  |
|               | resp_gen (no_hds)     | 0.771      | —     | 0.804         | —     | —                  |
|               | resp_gen (with_hds)   | 0.760      | —     | 0.777         | —     | —                  |
|               | story_gen (no_hds)    | 0.763      | —     | 0.771         | —     | —                  |
|               | story_gen (with_hds)  | 0.785      | —     | 0.785         | —     | —                  |

consistent (mean  $\Delta\text{AUC} = -0.013$  across the 12 binary tasks). On the 6 DecTest tasks (continuous temperature labels), Qwen3-30B-A3B is mildly better on 5 of 6 (mean  $+0.021$  in Spearman  $\rho$ , with one tie).


**Interpretation.** The binary result is a clear negative: scaling from 3B to 30B active parameters does not improve discrimination on McDiv or ConTest. We do not have a confident explanation. The Qwen3-30B-A3B run was not repeated after a bug fix applied to the primary experiments, so the comparison should be treated as preliminary. We leave the question of whether larger base models improve the metric on binary tasks to future work.

## I Cross-Metric Agreement on the OLMo-2-7B RLHF Experiment

The RLHF case study (Section 5) reports the headline monotone  $D_{Ca_n}$  drop in Table 4. Figure 17 below visualises the underlying  $\bar{a}_k$  curves and per-prompt  $D_{Ca_n}$  distributions for the same length-matched subset, and Figure 18 plots  $D_{Ca_n}$  against the lexical (EAD) and semantic (SentBERT) baselines, confirming that the three metrics see the same diversity-loss signal.



(a) Qwen2.5-3B: positive off-diagonal.



(b) GPT-2: negative off-diagonal.

Figure 5: Pairwise cross-mode surprise reduction matrices. Each cell  $(i, j)$  shows how many bits of surprise reduction mode  $i$  (target) receives from seeing mode  $j$  (context). Qwen shows diagonal dominance with pervasive positive off-diagonal (+1.9 bits mean); GPT-2 shows diagonal dominance with negative off-diagonal (-3.7 bits mean).






Figure 6: Pairwise symmetry scatter for Qwen2.5-3B: each point is one  $(i, j)$  pair, plotting  $M_{ij}$  vs.  $M_{ji}$ . Under symmetric mutual information, points would lie on  $y = x$ . The best-fit slope is 0.46 ( $R^2 = 0.24$ ), and many pairs have opposite signs: mode  $i$  can reduce surprise for mode  $j$  while  $j$  *increases* surprise for  $i$ .



(a) Qwen2.5-3B: slope = 0.84,  $R^2 = 0.34$ .



(b) GPT-2: slope = 0.35,  $R^2 = 0.08$ .

Figure 7: Row mean (how informative as context) vs. column mean (how much benefit from context) for each mode. Qwen shows tighter correlation closer to the identity, consistent with approximate symmetry of mutual information. GPT-2 shows widespread violations of non-negative conditioning (modes in the negative-row-mean region).




Figure 8: Per-token surprise reduction (bits) for three representative context–target pairs (Qwen2.5-3B). Bars are mean reductions across 5 context samples per pair; error bars are  $\pm 1$  standard deviation. Green bars indicate the conditioning context lowered the per-token surprise; red bars indicate it raised it. The three pairs are stratified-median selections (by total  $|\Delta|$ ) from the same-mode, positive cross-mode, and negative cross-mode strata of a 9-pair pool. Note how in scientific fact to haiku, the early newline triggers the largest surprise. Scientific facts may discuss "still water", but they almost never have a line with three words in it.




Figure 9: Cross-mode information transfer scales with model size. **Left:** Mean off-diagonal surprise reduction (with bootstrap 95% CI) transitions from negative to positive as models grow, indicating that larger models extract more information from cross-mode context. **Right:** The fraction of mode pairs showing positive cross-mode transfer increases monotonically across the 4 Llama models (blue;  $p = 4.2\%$  under monotone null). GPT-2 and Qwen2.5-3B (gray) are included for context but differ in architecture.




Figure 10: Decomposition of the gap between unconditional surprise  $\bar{a}_1$  and the asymptotic floor  $a_\infty$  at each step  $k$ . The per-step gap  $\bar{a}_1 - a_\infty$  splits into mutual information  $I_k = \bar{a}_1 - \bar{a}_k$  (red, above the curve) and excess  $e_k = \bar{a}_k - a_\infty$  (blue, below the curve). Summing across  $k$ : the red area is the total correlation  $TC_n$ ; the blue area is the excess entropy  $E$ . As  $k$  grows,  $e_k \rightarrow 0$  and each step contributes the full  $\bar{a}_1 - a_\infty$  to  $TC_n$ , so  $TC_n$  grows linearly while  $E$  converges. All quantities are in bits. This figure is theoretical: the sigmoidal  $\bar{a}_k$  shape reflects our initial expectation, but we subsequently found that the empirical curves are exponential decay without an initial plateau (see Section C.4).




Figure 11: ROC curves for McDiv\_nuggets binary classification (Qwen2.5-3B, 50 permutations). Five per-byte metrics are overlaid per panel:  $C \times a_n$  (red, solid),  $a_n$  alone (orange, dashed),  $D_{\text{fit}} = C \times E_{\text{fit}}$  (blue, solid),  $D_{\text{disc}} = C \times \hat{E}_n$  (cyan, dashed), and  $a_1$  (gray, dotted). Line style is redundant with color so the encoding remains legible under common color-vision deficiencies.  $C \times a_n$  dominates the other four scores on every task; the  $E$ -based scores ( $D_{\text{fit}}, D_{\text{disc}}$ ) hug or fall below the chance diagonal in this 5-response regime.




Figure 12: Summary metrics vs. mode count  $m$  (Qwen2.5-3B,  $n = 20$ , 1000 draws).  $E_{\text{fit}}$  (sigmoid-extrapolated) increases monotonically, while the raw  $\hat{E}_n$  does not.



### Sigmoid Fit Parameters vs Mode Count — Qwen/Qwen2.5-3B




Figure 13: Sigmoid fit parameters vs. mode count  $m$  (Qwen2.5-3B). The inflection point  $k_0$  remains at the lower bound ( $-10$ ) across all  $m$ , indicating pure exponential decay without an initial plateau (see Section C.4).




Figure 14: Distribution of  $a_1$  (unconditional surprise of the first response) for high- vs. low-diversity story\_gen samples. Left: per-byte. Right: total bits. The per-byte gap (see Table 7) is a dataset construction artifact.




Figure 15: Per-token surprise for a high-diversity sample’s first response. The continuation (“Joel fired the cook...”) is predictable, yielding low per-token surprise across response tokens.






Figure 16: Per-token surprise for a low-diversity sample’s first response. The specific ending (“He scored winner”) is inherently more surprising to the base model, despite being labeled “low diversity.”




(a) AlpacaEval  $\bar{a}_k$  curves (length-matched).

(b) Per-prompt  $D_{Ca_n}$  by stage (length-matched).

Figure 17: Progressive conditional surprise curves and per-prompt diversity distributions across the four OLMo-2-7B stages on the length-matched AlpacaEval subset. Each later stage’s  $\bar{a}_k$  curve lies below the base curve at every  $k \geq 2$ ; the per-prompt  $D_{Ca_n}$  distribution shifts toward lower values as the pipeline advances.



(a) AlpacaEval (length-matched,  $n = 150$  prompts).



(b) NoveltyBench curated (length-matched,  $n = 39$  prompts).

Figure 18: Per-prompt  $D_{Ca_n}$  versus EAD (left subpanel) and SentBERT-similarity diversity (right subpanel), coloured by stage, on the length-matched subset of prompts. Length-matching truncates each (stage, prompt) tuple’s responses to a common per-prompt byte budget so the per-byte conditional surprise that defines  $a_n$  is not depressed by response length. Pearson  $r$  and Spearman  $\rho$  (two-sided  $p$ ) are reported in each panel. The rank correlations are positive across both prompt sets and both baselines:  $D_{Ca_n}$  tracks the same diversity-loss signal EAD and SentBERT detect, and the agreement is not an artefact of response length.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Section 1 state that we (i) define the per-byte  $a_k$  progressive-surprise curve and the scalar  $D_{Ca_n} = C \times a_n$ , (ii) evaluate it on Tevet & Berant’s McDiv human-grounded benchmark, and (iii) apply it to the OLMo-2-7B post-training pipeline; Section 4 and Section 5 report exactly these results, with limitations (Section 6) flagging where the metric trails embedding baselines.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 is dedicated to limitations: measurement is relative to the base model  $\theta$ , cross-mode behaviour is model-dependent, the metric form  $D = C \times a_n$  was selected with full visibility into the Tevet benchmark, performance trails the strongest embedding baseline, the prompt format is fixed, and the per-byte denominator forces length-matching that drops short-response prompts.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best

judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The contribution is empirical: an estimator and a benchmark study, not a theorem. The information-theoretic identities we draw on (excess entropy, entropy rate) are standard and cited; informal derivations connecting the  $a_k$  curve to those quantities are in Appendix E.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 3 fully specifies the metric (prompt template, permutation-averaged single-pass forward computation, per-byte cross-entropy with explicit byte-count denominators); Section 4 and Section 5 state the base model ( $\theta$ ), permutation count, response set sizes, and length-matching protocol used in each experiment; the anonymous code mirror at <https://anonymous.4open.science/r/icl-diversity-67E6/> contains the exact run scripts (scripts/run\_scenarios.py, scripts/analyze\_c\_ainf.py) that produced every reported number.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example

- (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
- (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Anonymous code and data mirror at <https://anonymous.4open.science/r/icl-diversity-67E6/> contains the full implementation (src/icl\_diversity/), all experiment scripts, the released OLMo-2-7B generations under results/r1hf\_experiment/, and command-line invocations in CLAUDE.md that reproduce every reported scalar and figure.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: We do not train models; the metric is a forward-pass-only estimator over a base  $\theta$ . Section 3 specifies the response-conditioning template, permutation count ( $|\Sigma| = 50$  at scale,  $|\Sigma| = 20$  for synthetic studies), choice of  $\theta$  (Qwen2.5-3B base for headline; ablations in Appendix H), and the length-matching protocol used in the OLMo-2-7B experiment.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification:  $a_k$  curves are reported with permutation-averaged means and standard-deviation bands across response orderings (Section 3). The OLMo-2-7B comparison reports per-prompt  $D_{C_{a_n}}$  means with standard deviations and pre-registered Bonferroni-corrected paired tests for the H1a–H1c stage contrasts, with the H1' comparison explicitly marked as exploratory (Section 5, Table 4).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The metric requires only forward passes through a base language model. Section B.2 reports the cost as  $|\Sigma| + n$  forward passes per (prompt, response set). Local experiments (Qwen2.5-3B base, GPT-2) ran on a workstation with  $2 \times$  Quadro RTX 8000 GPUs (48 GB VRAM each); each individual experiment completed in at most a few hours. The Qwen3-30B-A3B-Base ablation (Appendix H) was run through the Tinker API rather than on local hardware; total Tinker API spend across all 30B runs (including those reported in the paper) was under \$17. No model training was performed in this work.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: We use only publicly released models (GPT-2, Qwen2.5-3B, Qwen3-30B-A3B-Base, OLMo-2-7B) and the publicly released Tevet & Berant McDiv benchmark; no human subjects were enrolled, no data was scraped, and no PII is processed.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Impact Statement (immediately preceding the references) discusses both the positive impact (better measurement of mode collapse and creative-writing diversity supports evaluation and oversight of generative AI) and the potential negative impacts (use in pipelines that select for more capable AI; possible automation of human creative roles).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]



1239 Justification: We release neither a high-risk pretrained model nor scraped data. The released  
1240 artefact is the set of OLMo-2-7B generations on AlpacaEval and NoveltyBench-curated  
1241 prompts (a transcript of standard sampled outputs from a public model), plus the metric  
1242 implementation; none of these introduce misuse risks beyond what the underlying public  
1243 OLMo-2 release already poses.

1244 Guidelines:

- 1245 • The answer [N/A] means that the paper poses no such risks.
- 1246 • Released models that have a high risk for misuse or dual-use should be released with  
1247 necessary safeguards to allow for controlled use of the model, for example by requiring  
1248 that users adhere to usage guidelines or restrictions to access the model or implementing  
1249 safety filters.
- 1250 • Datasets that have been scraped from the Internet could pose safety risks. The authors  
1251 should describe how they avoided releasing unsafe images.
- 1252 • We recognize that providing effective safeguards is challenging, and many papers do  
1253 not require this, but we encourage authors to take this into account and make a best  
1254 faith effort.

## 1255 12. Licenses for existing assets

1256 Question: Are the creators or original owners of assets (e.g., code, data, models), used in  
1257 the paper, properly credited and are the license and terms of use explicitly mentioned and  
1258 properly respected?

1259 Answer: [Yes]

1260 Justification: All models and datasets used in the paper are cited by their canonical references.  
1261 Models: GPT-2 [Radford et al., 2019], Qwen2.5-3B [Yang et al., 2024], Qwen3-30B-A3B-  
1262 Base [Yang et al., 2025], OLMo-2-7B [OLMo et al., 2024], and the Llama 3 family at 8B/70B  
1263 [Grattafiori et al., 2024] together with the Llama 3.2 lightweight 1B/3B variants [Meta AI,  
1264 2024]; SentBERT-style sentence embeddings [Reimers and Gurevych, 2019] are used as  
1265 a baseline. Datasets: Tevet & Berant’s diversity-eval / McDiv [Tevet and Berant, 2021],  
1266 AlpacaFarm/AlpacaEval prompts [Dubois et al., 2023], and NoveltyBench-curated prompts  
1267 [Zhang et al., 2025c]. Most assets are released under permissive licenses (Apache-2.0 for the  
1268 Qwen and OLMo families, MIT for GPT-2 and the diversity-eval repo, research-use licenses  
1269 for the benchmark datasets); Llama 3 and Llama 3.2 are released under the bespoke Meta  
1270 Llama Community License, which is more restrictive than Apache-2.0/MIT but permits  
1271 research use. All assets are used within their stated terms.

1272 Guidelines:

- 1273 • The answer [N/A] means that the paper does not use existing assets.
- 1274 • The authors should cite the original paper that produced the code package or dataset.
- 1275 • The authors should state which version of the asset is used and, if possible, include a  
1276 URL.
- 1277 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1278 • For scraped data from a particular source (e.g., website), the copyright and terms of  
1279 service of that source should be provided.
- 1280 • If assets are released, the license, copyright information, and terms of use in the  
1281 package should be provided. For popular datasets, paperswithcode.com/datasets  
1282 has curated licenses for some datasets. Their licensing guide can help determine the  
1283 license of a dataset.
- 1284 • For existing datasets that are re-packaged, both the original license and the license of  
1285 the derived asset (if it has changed) should be provided.
- 1286 • If this information is not available online, the authors are encouraged to reach out to  
1287 the asset’s creators.

## 1288 13. New assets

1289 Question: Are new assets introduced in the paper well documented and is the documentation  
1290 provided alongside the assets?

1291 Answer: [Yes]

1292 Justification: New assets released: (i) the metric implementation (open-source code at  
1293 <https://anonymous.4open.science/r/icl-diversity-67E6/> with README, unit  
1294 tests, and a CLI documented in CLAUDE.md); (ii) the OLMo-2-7B generations dataset (200  
1295 AlpacaEval + 100 NoveltyBench-curated prompts  $\times$  4 stages  $\times$   $K = 10$  generations)  
1296 released under `results/r1hf_experiment/` with a sidecar README describing schema,  
1297 sampling parameters, and license. No new model weights are released.

1298 Guidelines:

- 1299 • The answer [N/A] means that the paper does not release new assets.
- 1300 • Researchers should communicate the details of the dataset/code/model as part of their  
1301 submissions via structured templates. This includes details about training, license,  
1302 limitations, etc.
- 1303 • The paper should discuss whether and how consent was obtained from people whose  
1304 asset is used.
- 1305 • At submission time, remember to anonymize your assets (if applicable). You can either  
1306 create an anonymized URL or include an anonymized zip file.

#### 1307 14. Crowdsourcing and research with human subjects

1308 Question: For crowdsourcing experiments and research with human subjects, does the paper  
1309 include the full text of instructions given to participants and screenshots, if applicable, as  
1310 well as details about compensation (if any)?

1311 Answer: [N/A]

1312 Justification: We did not run crowdsourcing or human-subjects experiments. The human  
1313 ratings used in the Tevet evaluation are from the previously published McDiv benchmark,  
1314 which is cited as the source.

1315 Guidelines:

- 1316 • The answer [N/A] means that the paper does not involve crowdsourcing nor research  
1317 with human subjects.
- 1318 • Including this information in the supplemental material is fine, but if the main contribu-  
1319 tion of the paper involves human subjects, then as much detail as possible should be  
1320 included in the main paper.
- 1321 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,  
1322 or other labor should be paid at least the minimum wage in the country of the data  
1323 collector.

#### 1324 15. Institutional review board (IRB) approvals or equivalent for research with human 1325 subjects

1326 Question: Does the paper describe potential risks incurred by study participants, whether  
1327 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)  
1328 approvals (or an equivalent approval/review based on the requirements of your country or  
1329 institution) were obtained?

1330 Answer: [N/A]

1331 Justification: No human-subjects research was performed in this work; we use only previ-  
1332 ously released human ratings from the Tevet & Berant McDiv benchmark.

1333 Guidelines:

- 1334 • The answer [N/A] means that the paper does not involve crowdsourcing nor research  
1335 with human subjects.
- 1336 • Depending on the country in which research is conducted, IRB approval (or equivalent)  
1337 may be required for any human subjects research. If you obtained IRB approval, you  
1338 should clearly state this in the paper.
- 1339 • We recognize that the procedures for this may vary significantly between institutions  
1340 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the  
1341 guidelines for their institution.
- 1342 • For initial submissions, do not include any information that would break anonymity (if  
1343 applicable), such as the institution conducting the review.

#### 1344 16. Declaration of LLM usage

1345 Question: Does the paper describe the usage of LLMs if it is an important, original, or  
1346 non-standard component of the core methods in this research? Note that if the LLM is used  
1347 only for writing, editing, or formatting purposes and does *not* impact the core methodology,  
1348 scientific rigor, or originality of the research, declaration is not required.

1349 Answer: [\[Yes\]](#)

1350 Justification: An LLM is the central object of the proposed method: a base language model  $\theta$   
1351 (e.g. Qwen2.5-3B) is used at inference time as the scoring model that produces the  $a_k$  curve  
1352 and the per-byte cross-entropy  $C$  (Section 3). The choice of  $\theta$  is documented per-experiment,  
1353 and ablations across model scales are reported in Appendix H.

1354 Guidelines:

- 1355 • The answer [\[N/A\]](#) means that the core method development in this research does not
- 1356 involve LLMs as any important, original, or non-standard components.
- 1357 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
- 1358 be described.